

NEURAL NETWORKS AND DEEP LEARNING

MICHELLE KUCHERA
DAVIDSON COLLEGE

MACHINE LEARNING WORKSHOP
TRIUMF
5 MAY 2022

LECTURE 1 TOPICS

- Computational graphs
- Gradient-descent optimization
- Logistic regression
- Regression neural networks

MICHELLE KUCHERA
DAVIDSON COLLEGE

CPS-FR

MIT

25 AUGUST 2022

GOALS

- Each of us learns something today
- Stop me with any questions

MICHELLE KUCHERA
DAVIDSON COLLEGE

CPS-FR

MIT

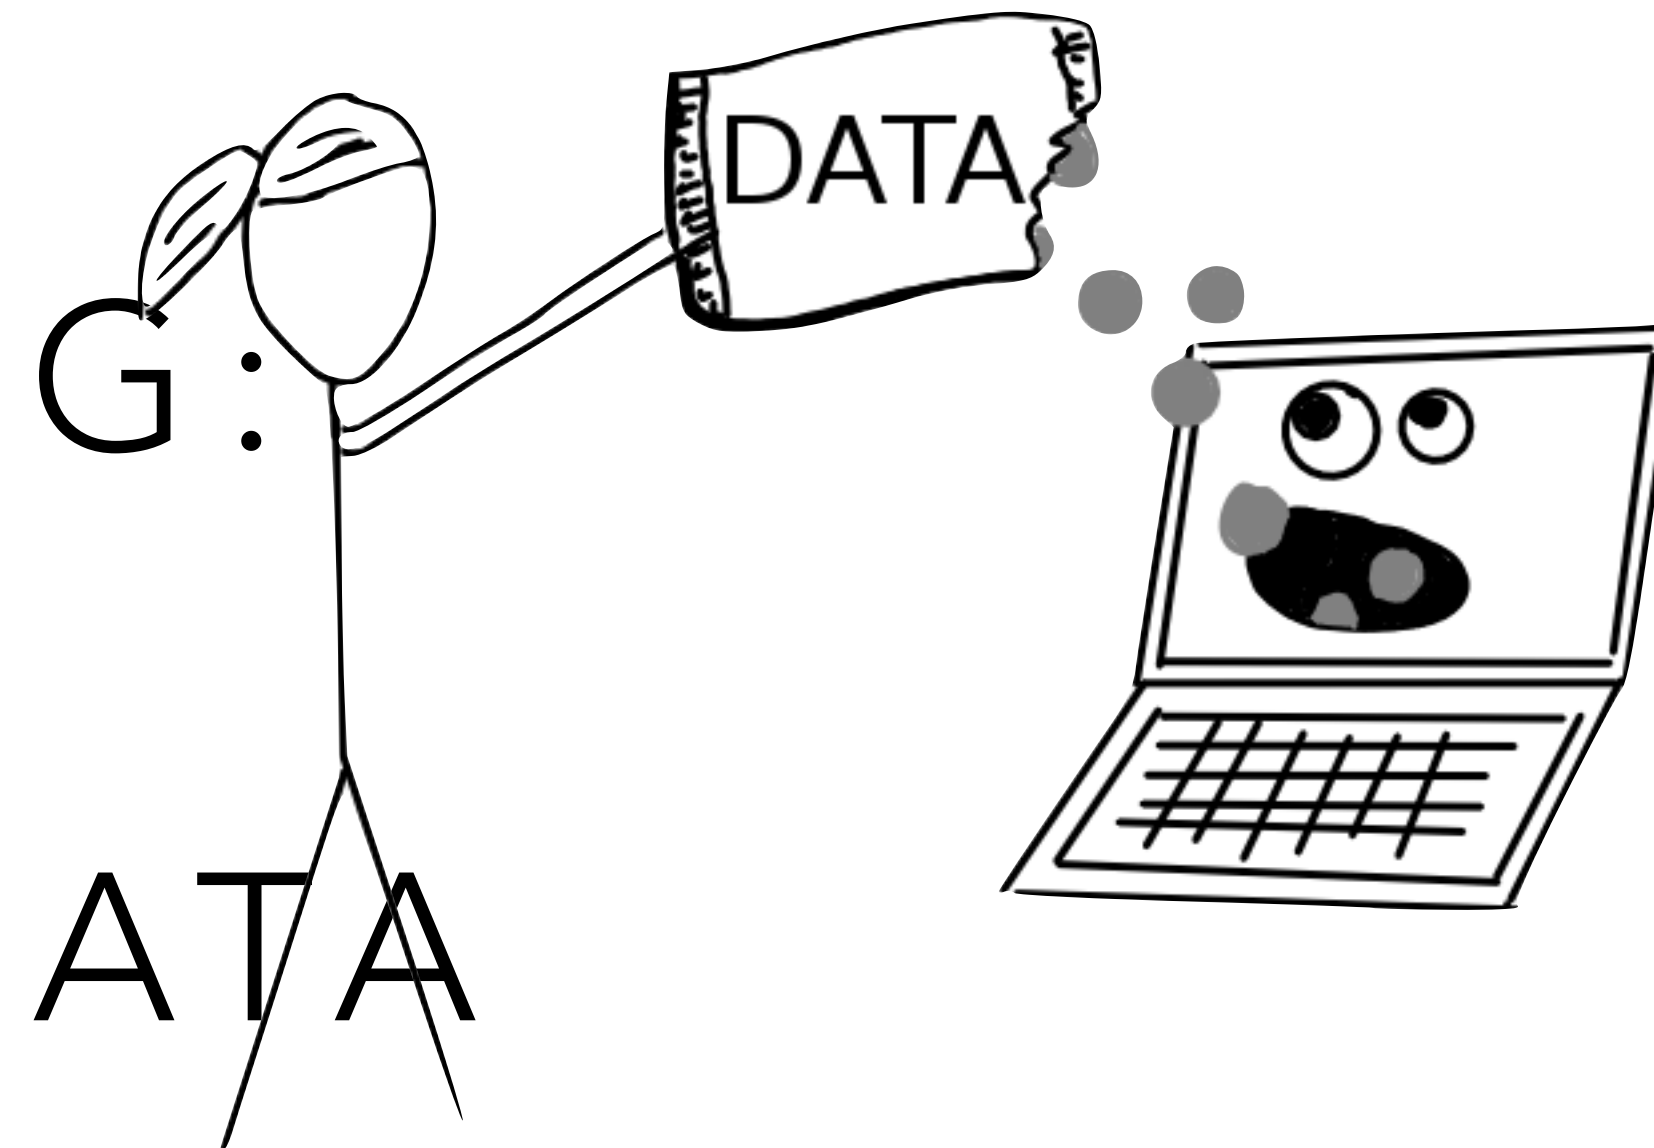
25 AUGUST 2022

COMMUNITY

- Each of you arrived here with your own backgrounds, specialty, and path in life
- Your experience and expertise are valuable here, no matter what it is
- If the activity is within your background, help others!
- If you are totally (or a little) lost, ask for help!
- It is our shared goal to have **each** of us leave with some new skill/knowledge/understanding

Without Machine Learning

With Machine Learning

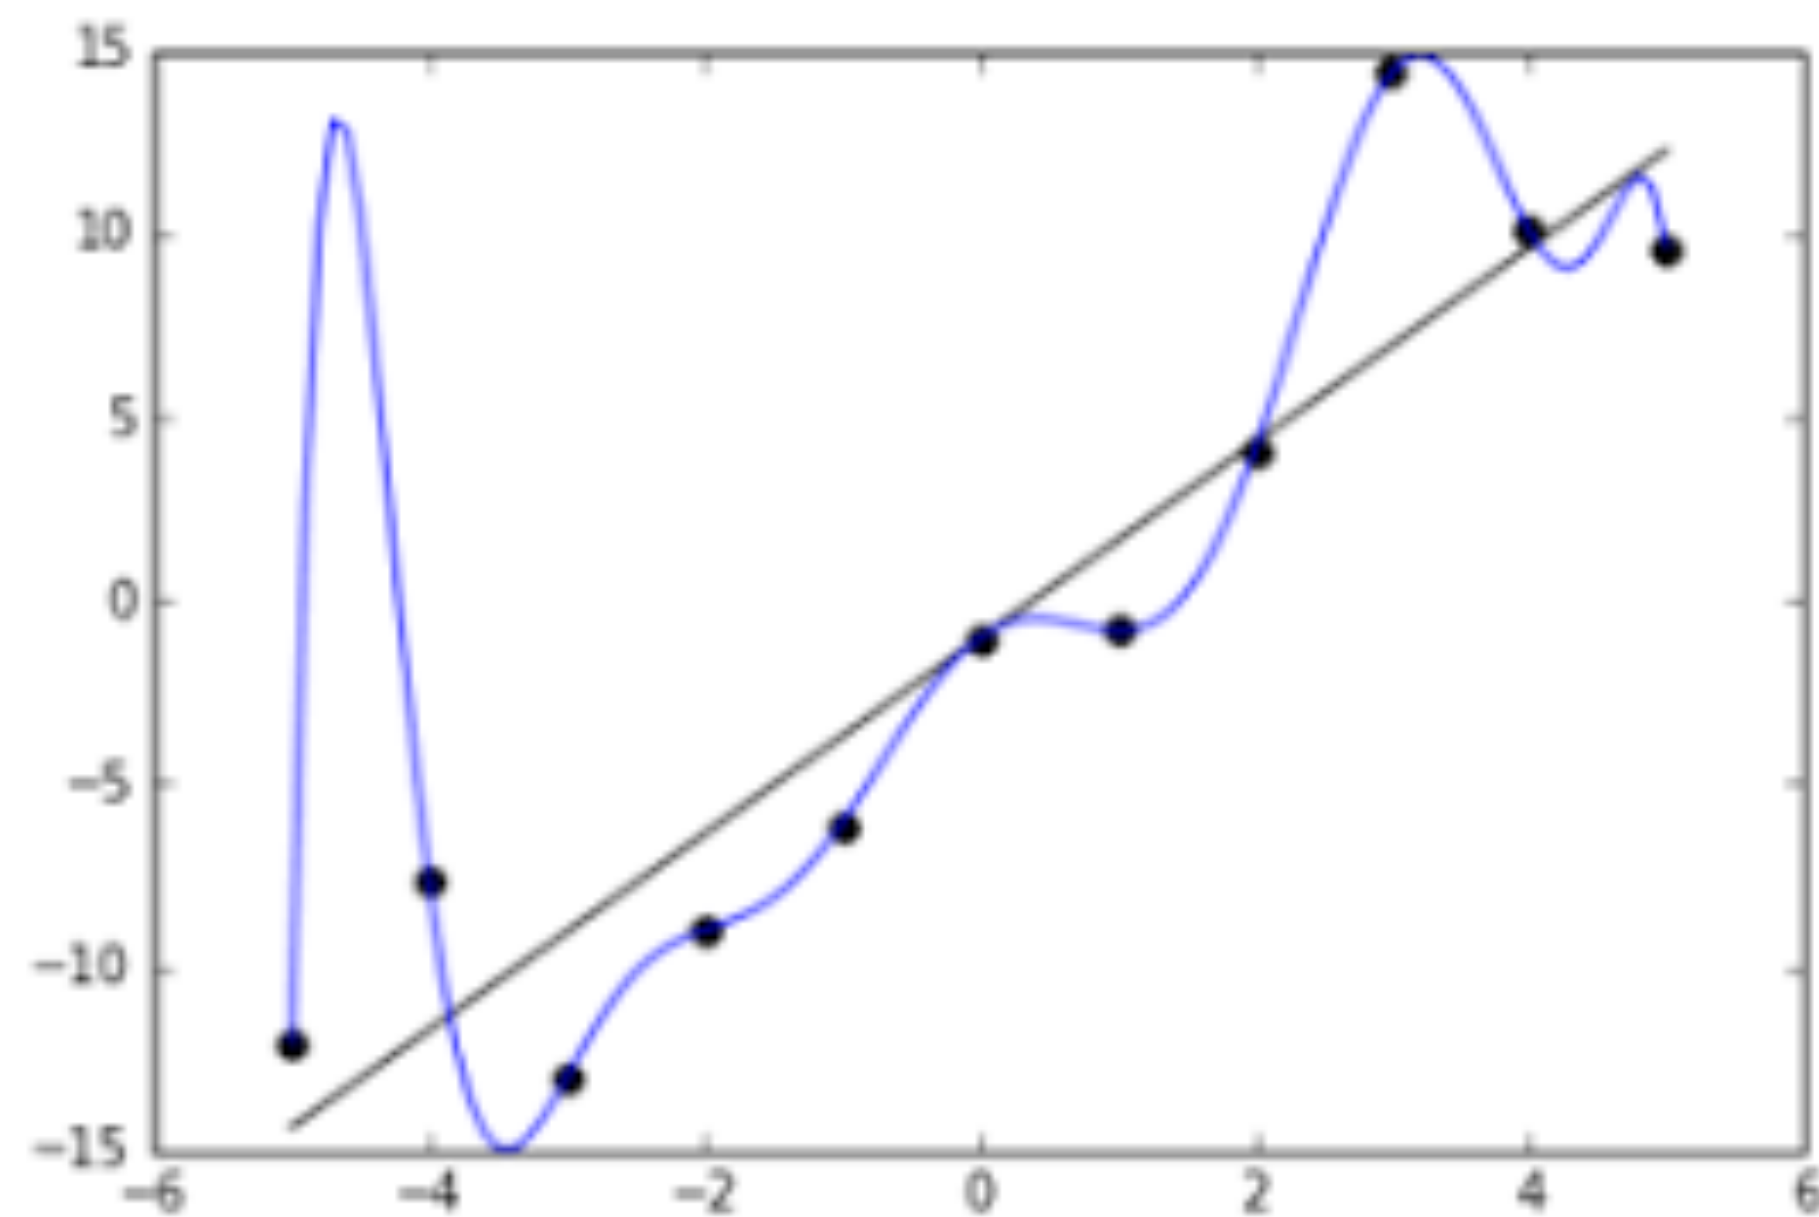


Without Machine Learning

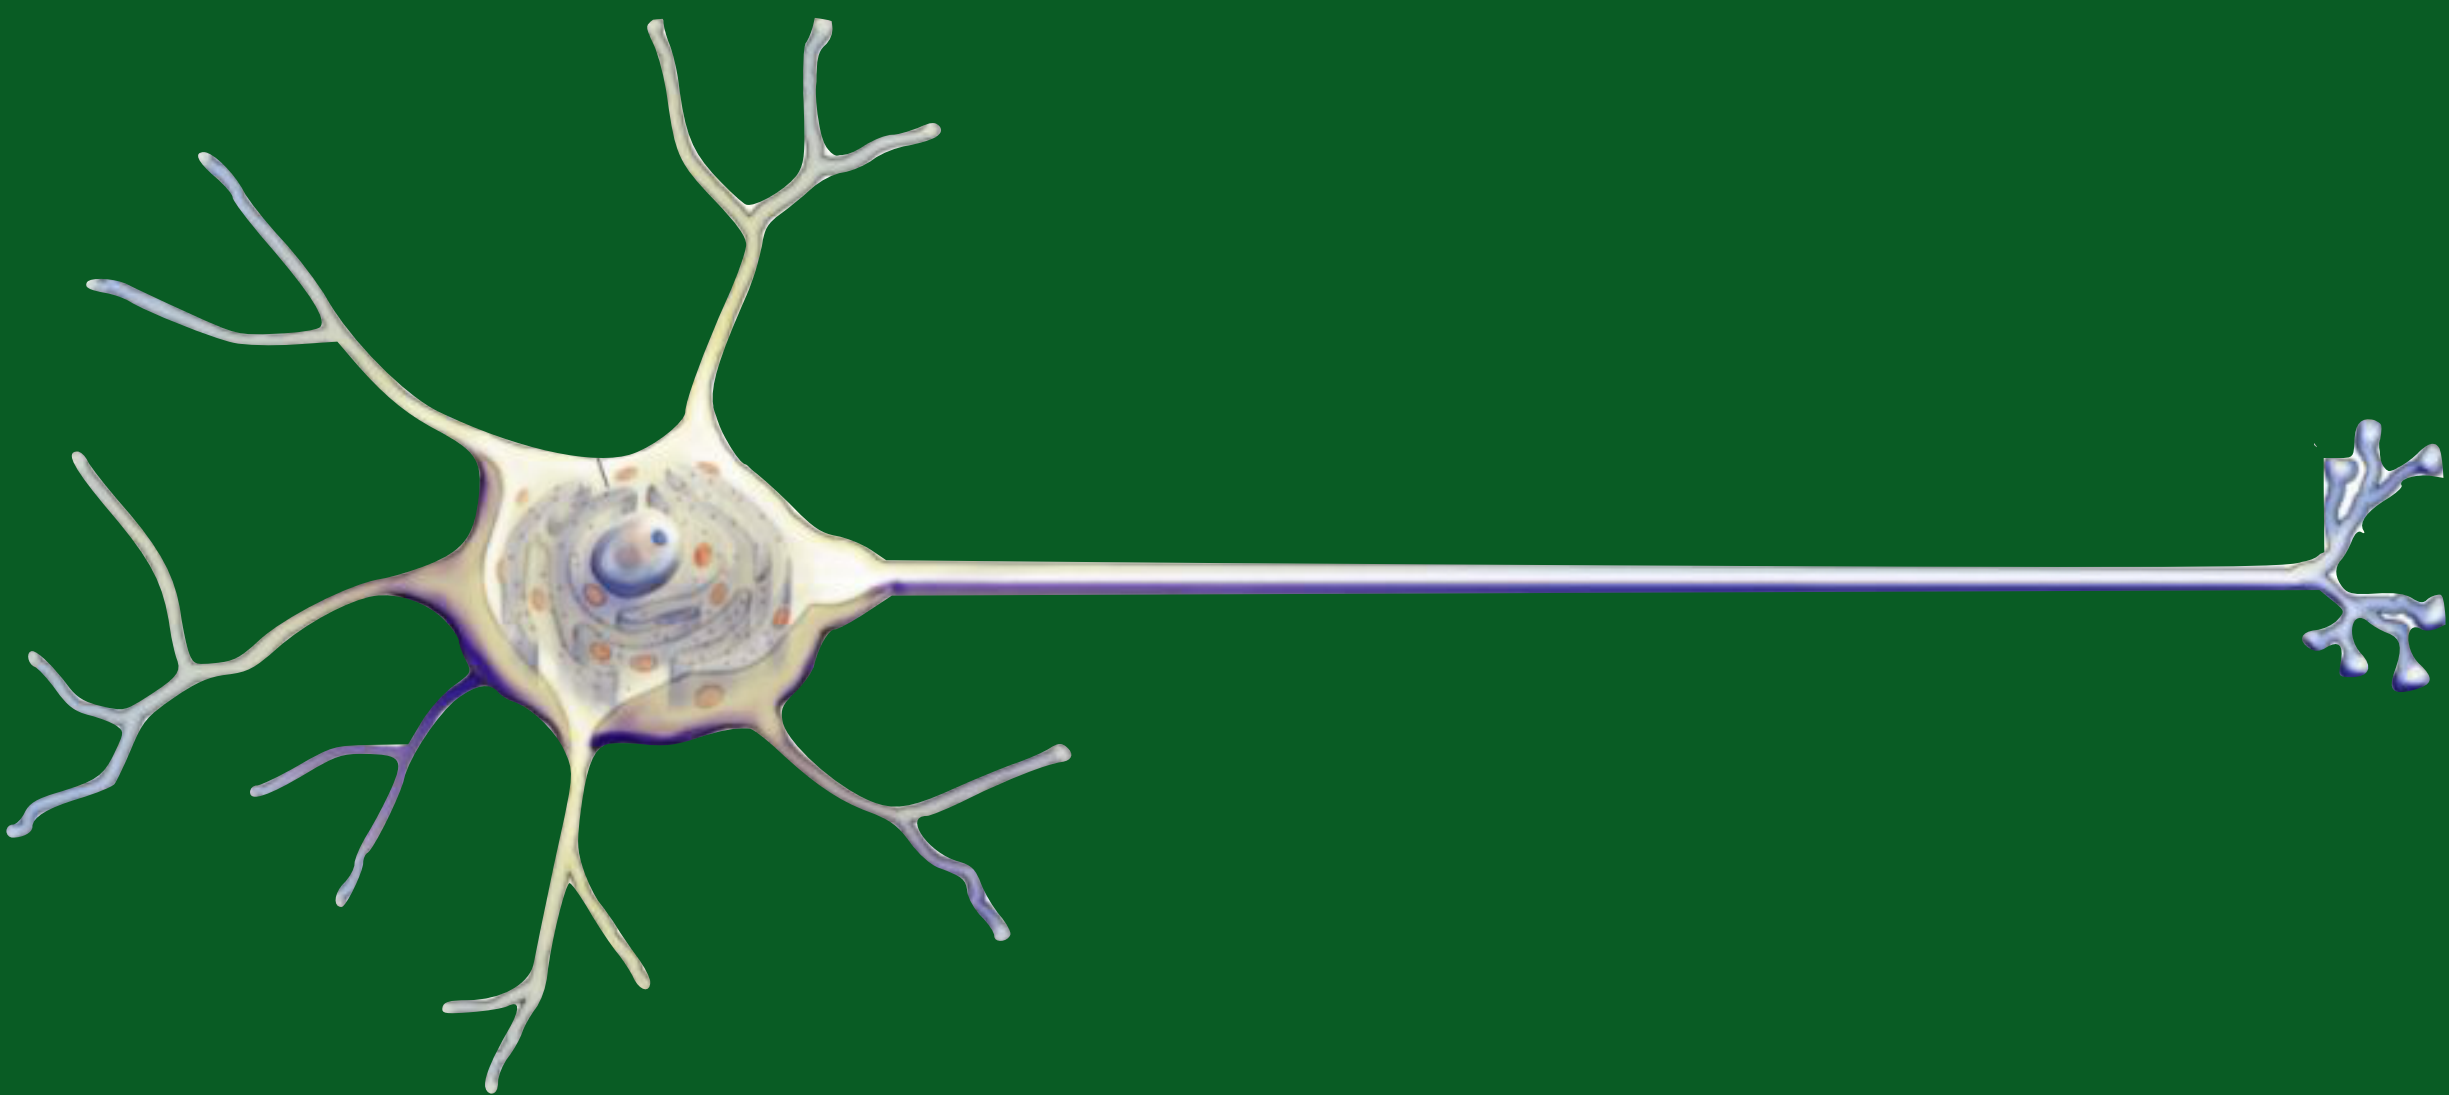
With Machine Learning

Learning from data is a **paradigm shift** in thinking about predictive models

* VERY SPECIFIC INSTRUCTIONS



NEURON



MATHEMATICS



Neural Networks
Volume 4, Issue 2, 1991, Pages 251-257



Approximation capabilities of multilayer feedforward networks

Kurt Hornik 

[Show more](#) 

 Share  Cite

[https://doi.org/10.1016/0893-6080\(91\)90009-T](https://doi.org/10.1016/0893-6080(91)90009-T)

[Get rights and content](#)

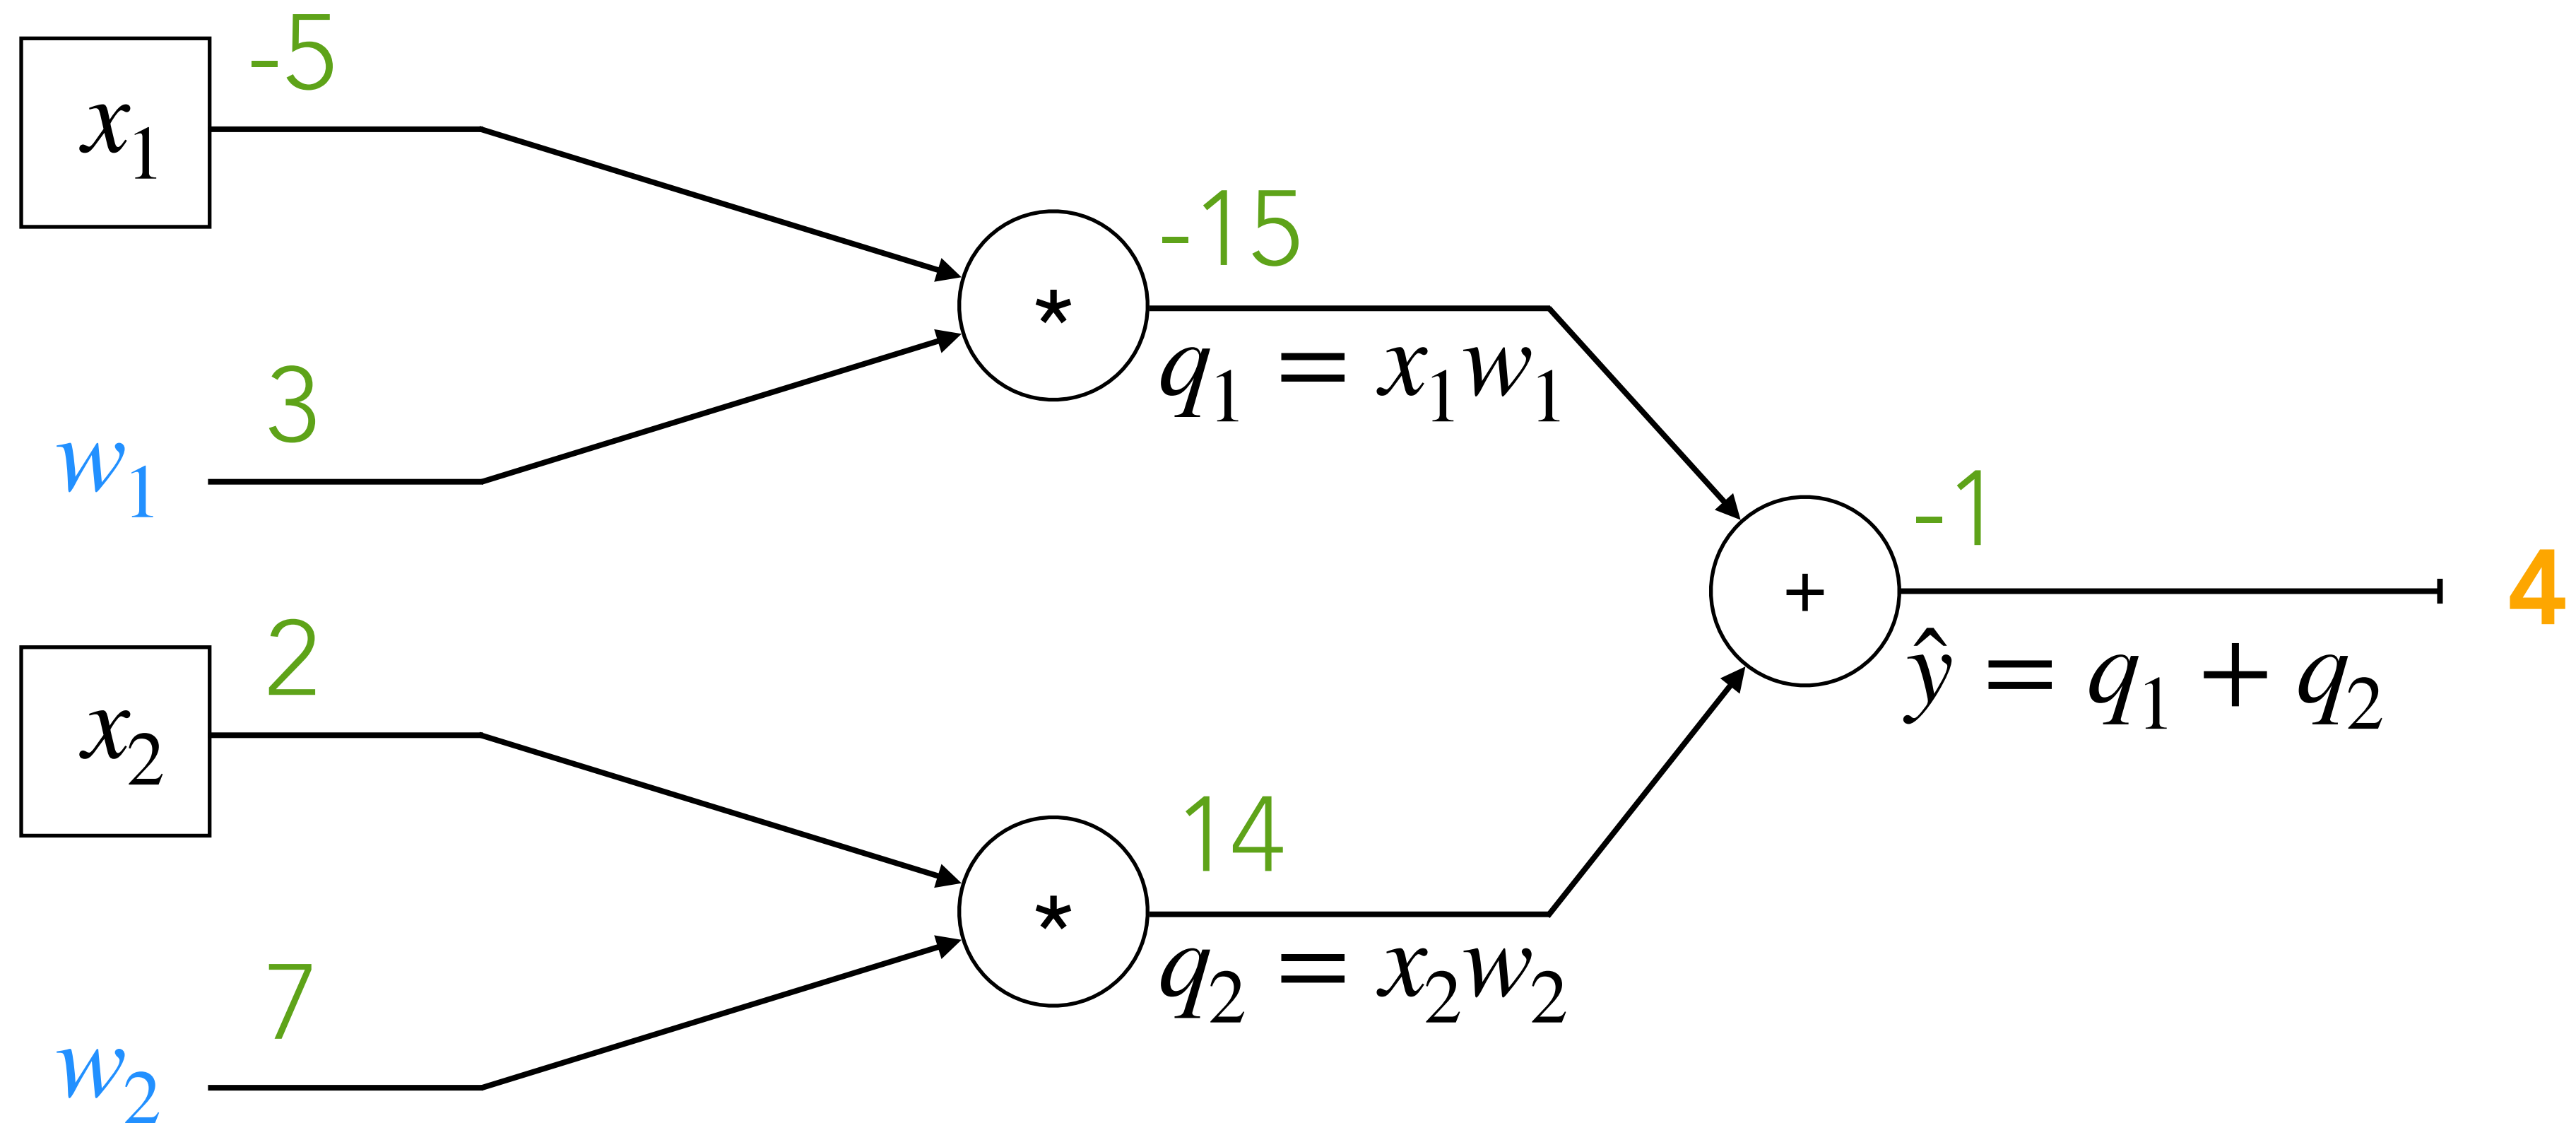
Abstract

We show that standard multilayer feedforward networks with as few as a single hidden layer and arbitrary bounded and nonconstant activation function are universal approximators with respect to $L^p(\mu)$ performance criteria, for arbitrary finite input environment measures μ , provided only that sufficiently many hidden units are available. If the activation function is continuous, bounded and nonconstant, then continuous mappings can be learned uniformly over compact input sets. We also give very general conditions ensuring that networks with sufficiently smooth activation functions are capable of arbitrarily accurate approximation to a function and its derivatives.

MATHEMATICS

COMPUTATIONAL GRAPH

$$\hat{y} = x_1 w_1 + x_2 w_2$$

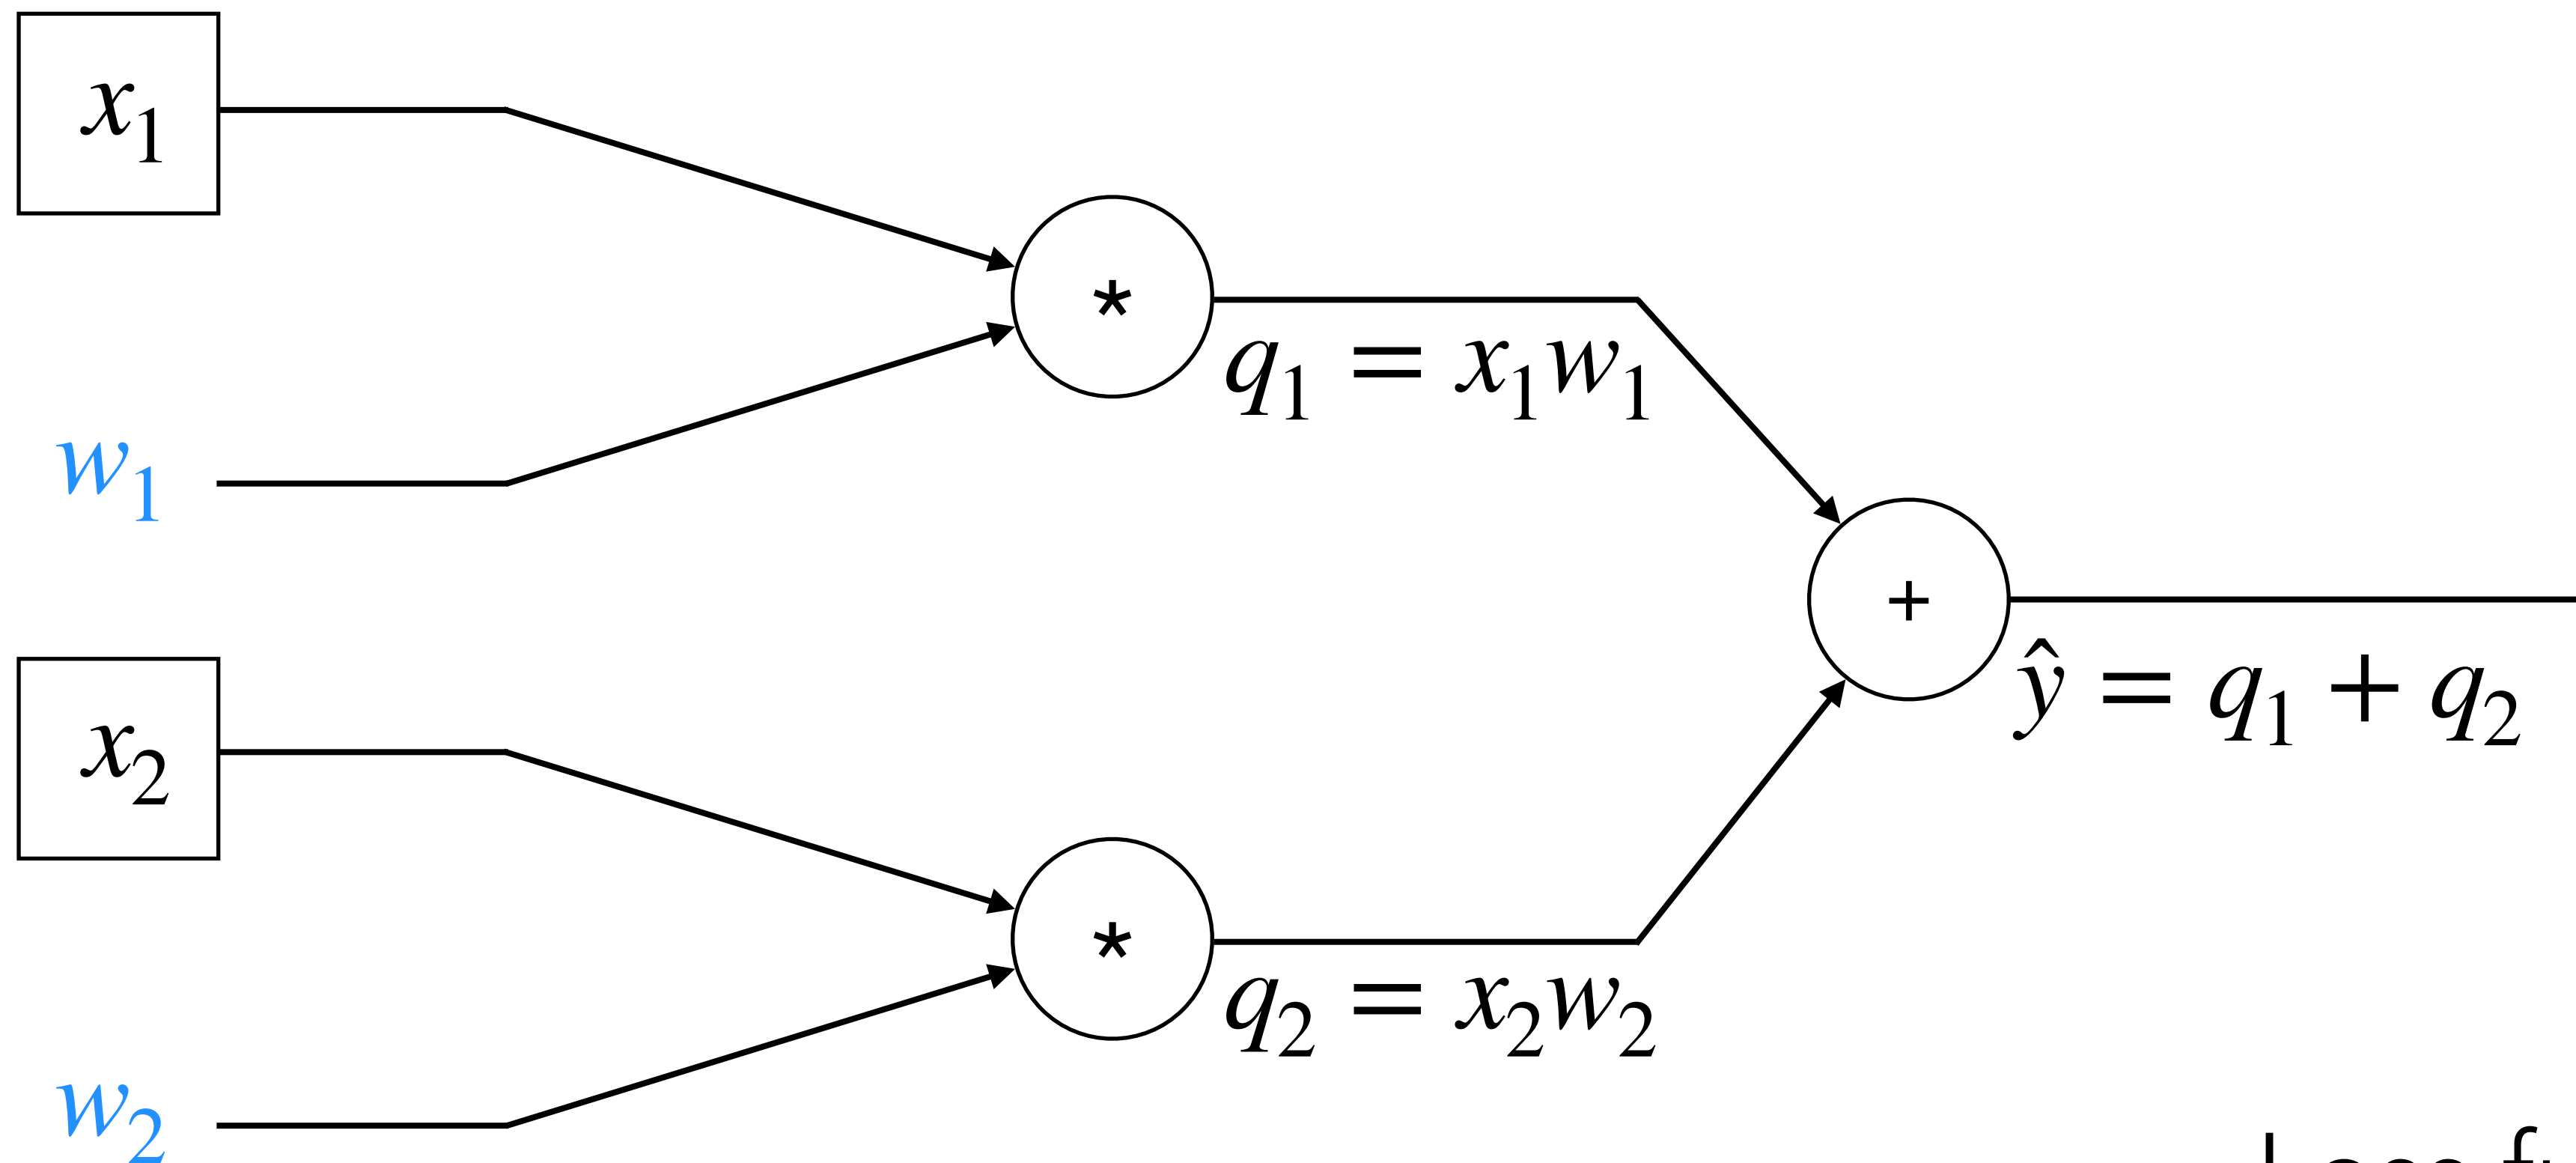


MACHINE LEARNING

SUPERVISED LEARNING

REGRESSION

$$\hat{y} = x_1 w_1 + x_2 w_2$$

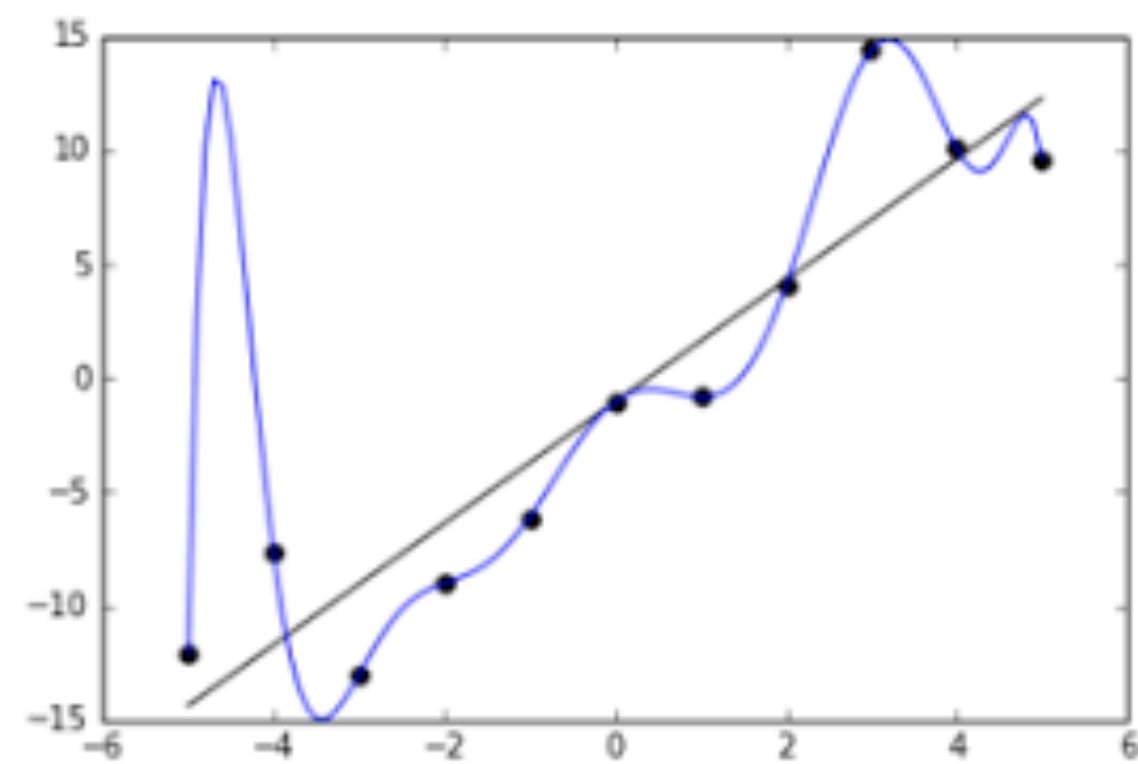
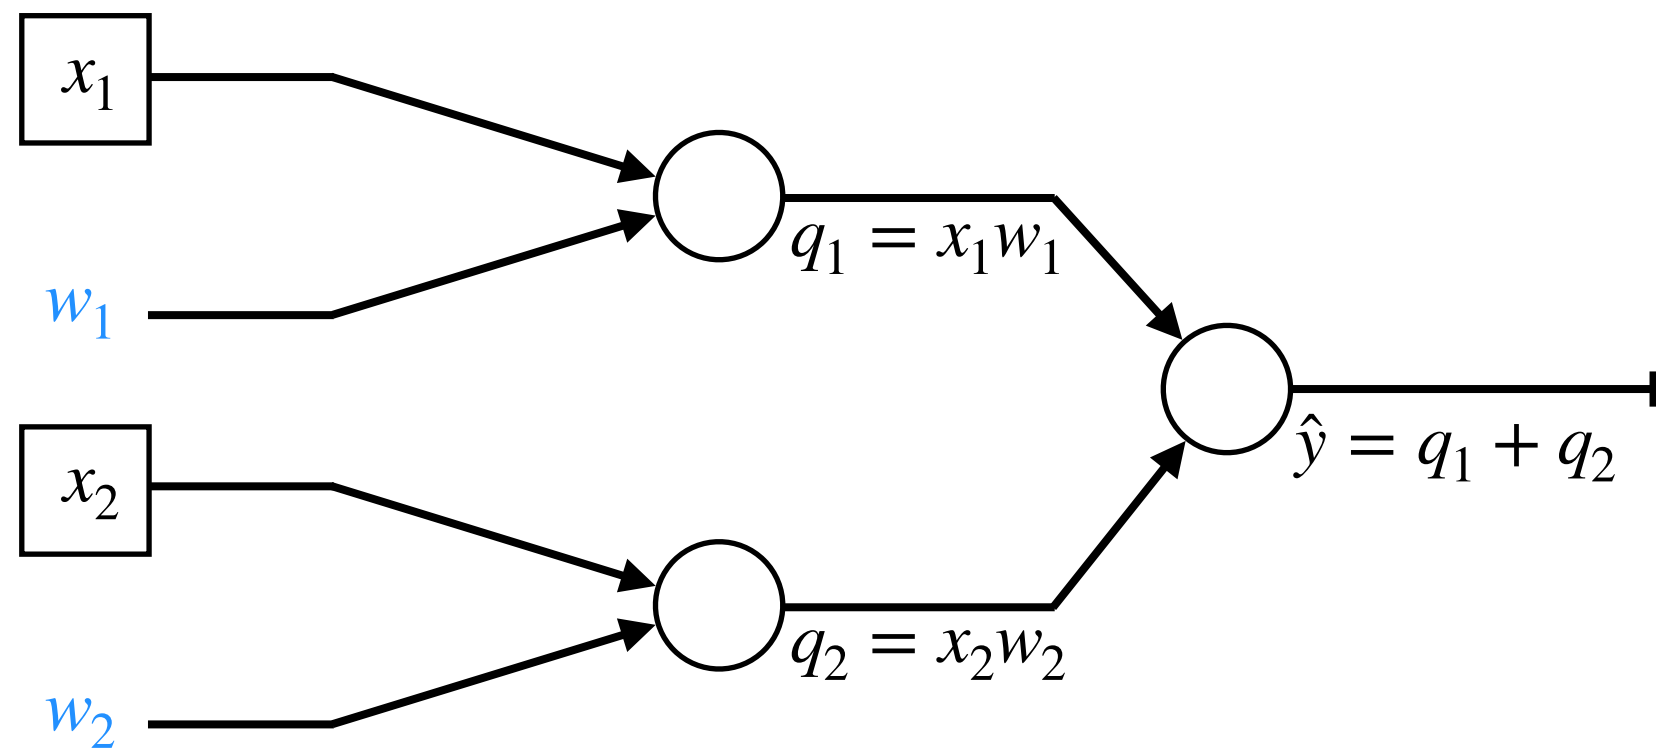


Loss function

$$J(w) = \hat{y} - y$$

SUPERVISED LEARNING

$$\hat{y} = x_1 w_1 + x_2 w_2$$



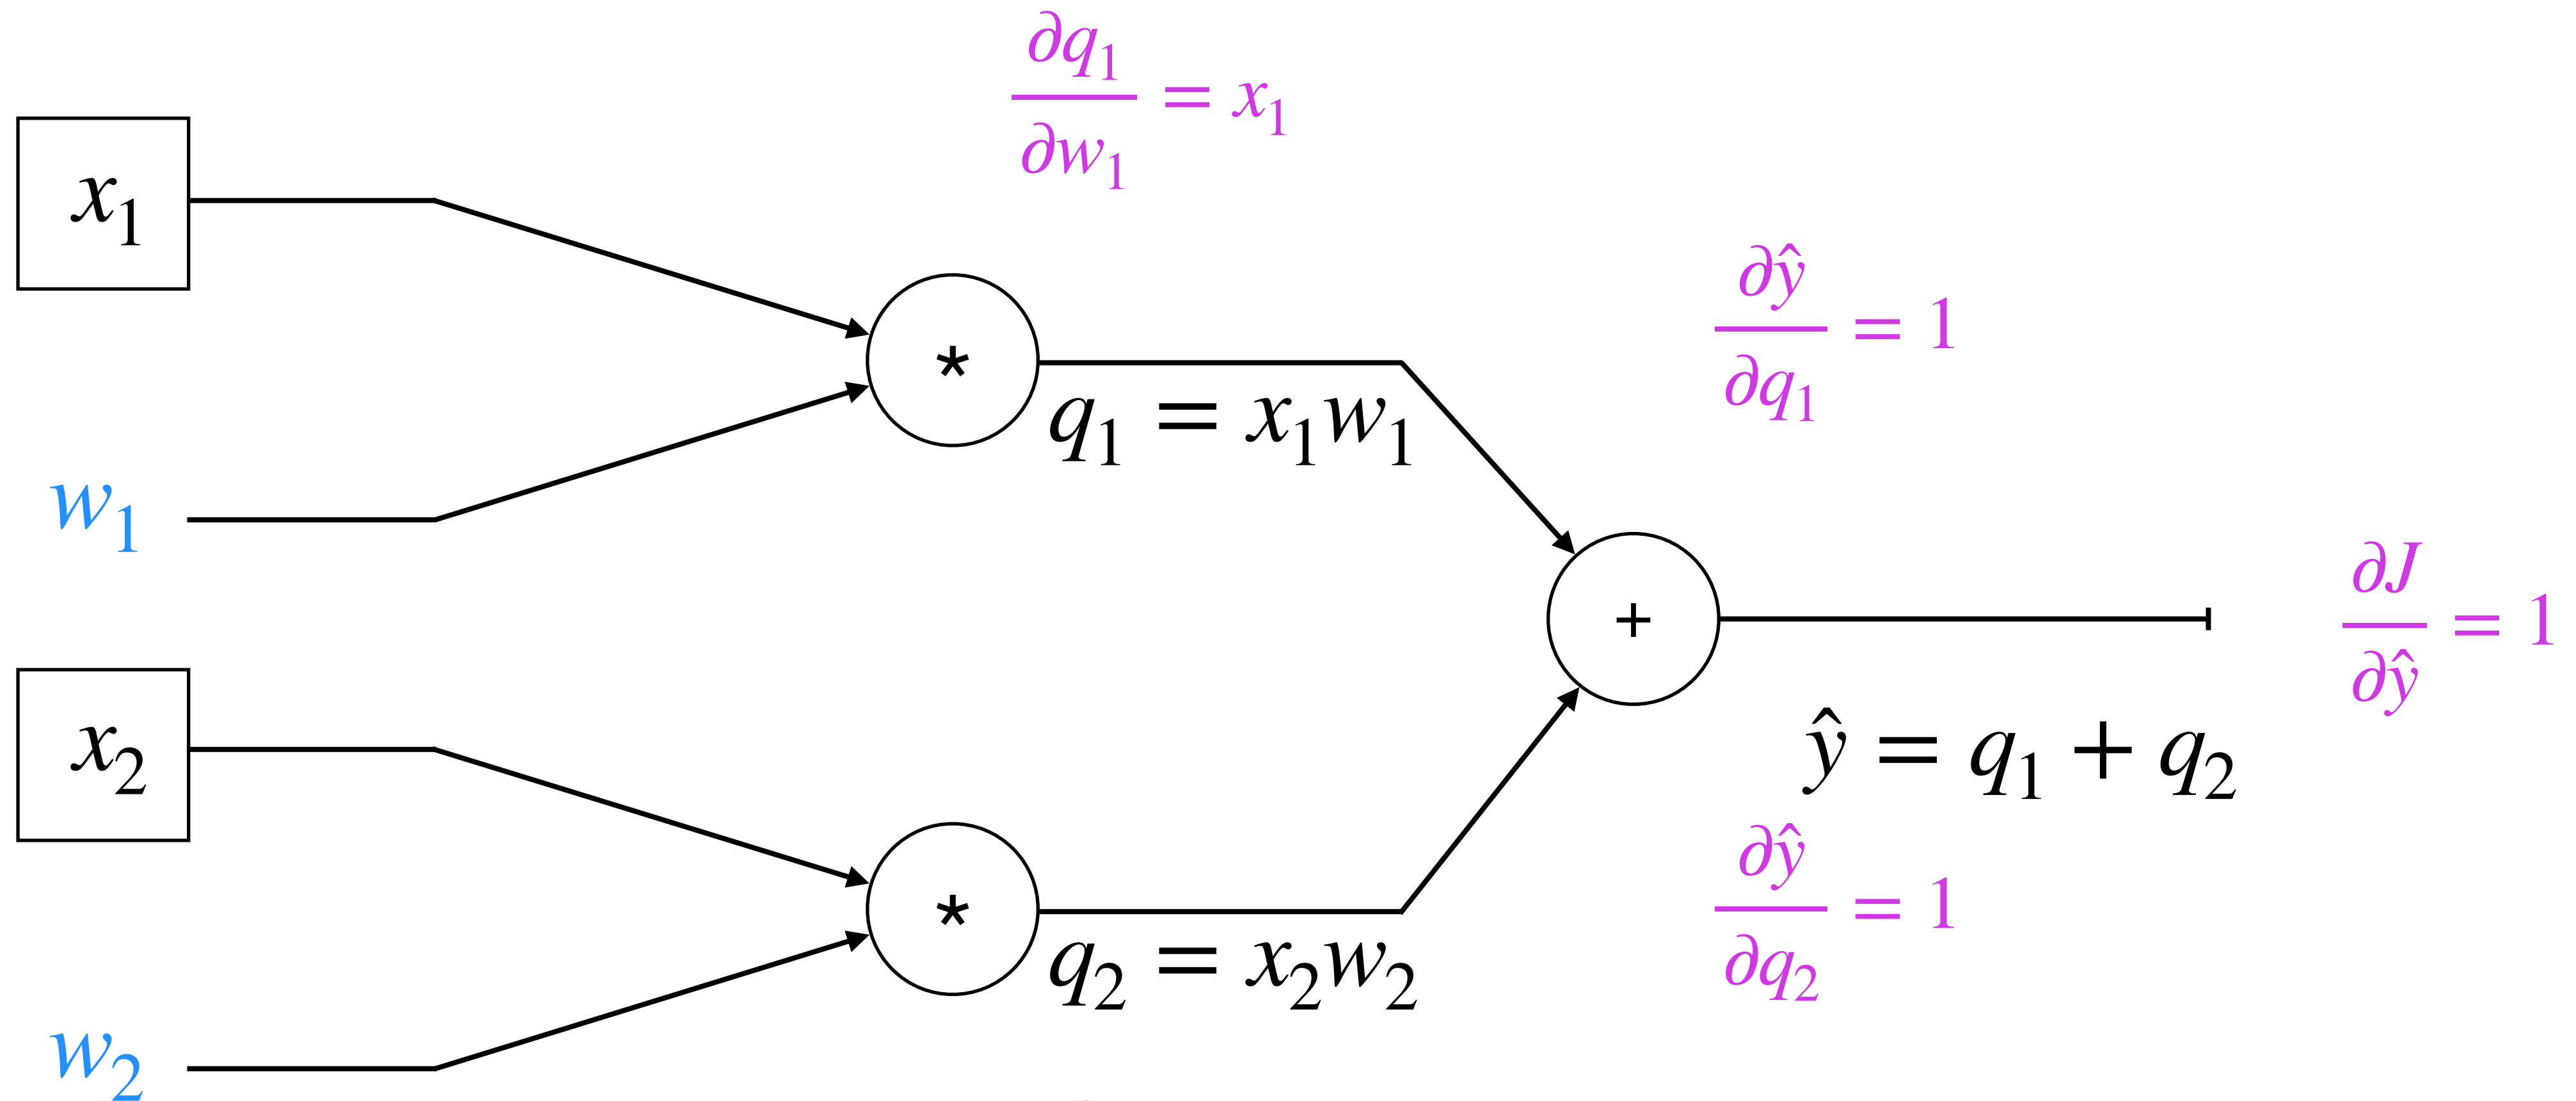
	Feature 1	Feature 2	Feature 3	Target
Example 1				
Example 2				
Example 3				
Example 4				

Loss function
MSE across N examples

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

BACKPROPAGATION

$$w_1 = w_1 - \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial f}{\partial q_1} \frac{\partial q_1}{\partial w_1}$$

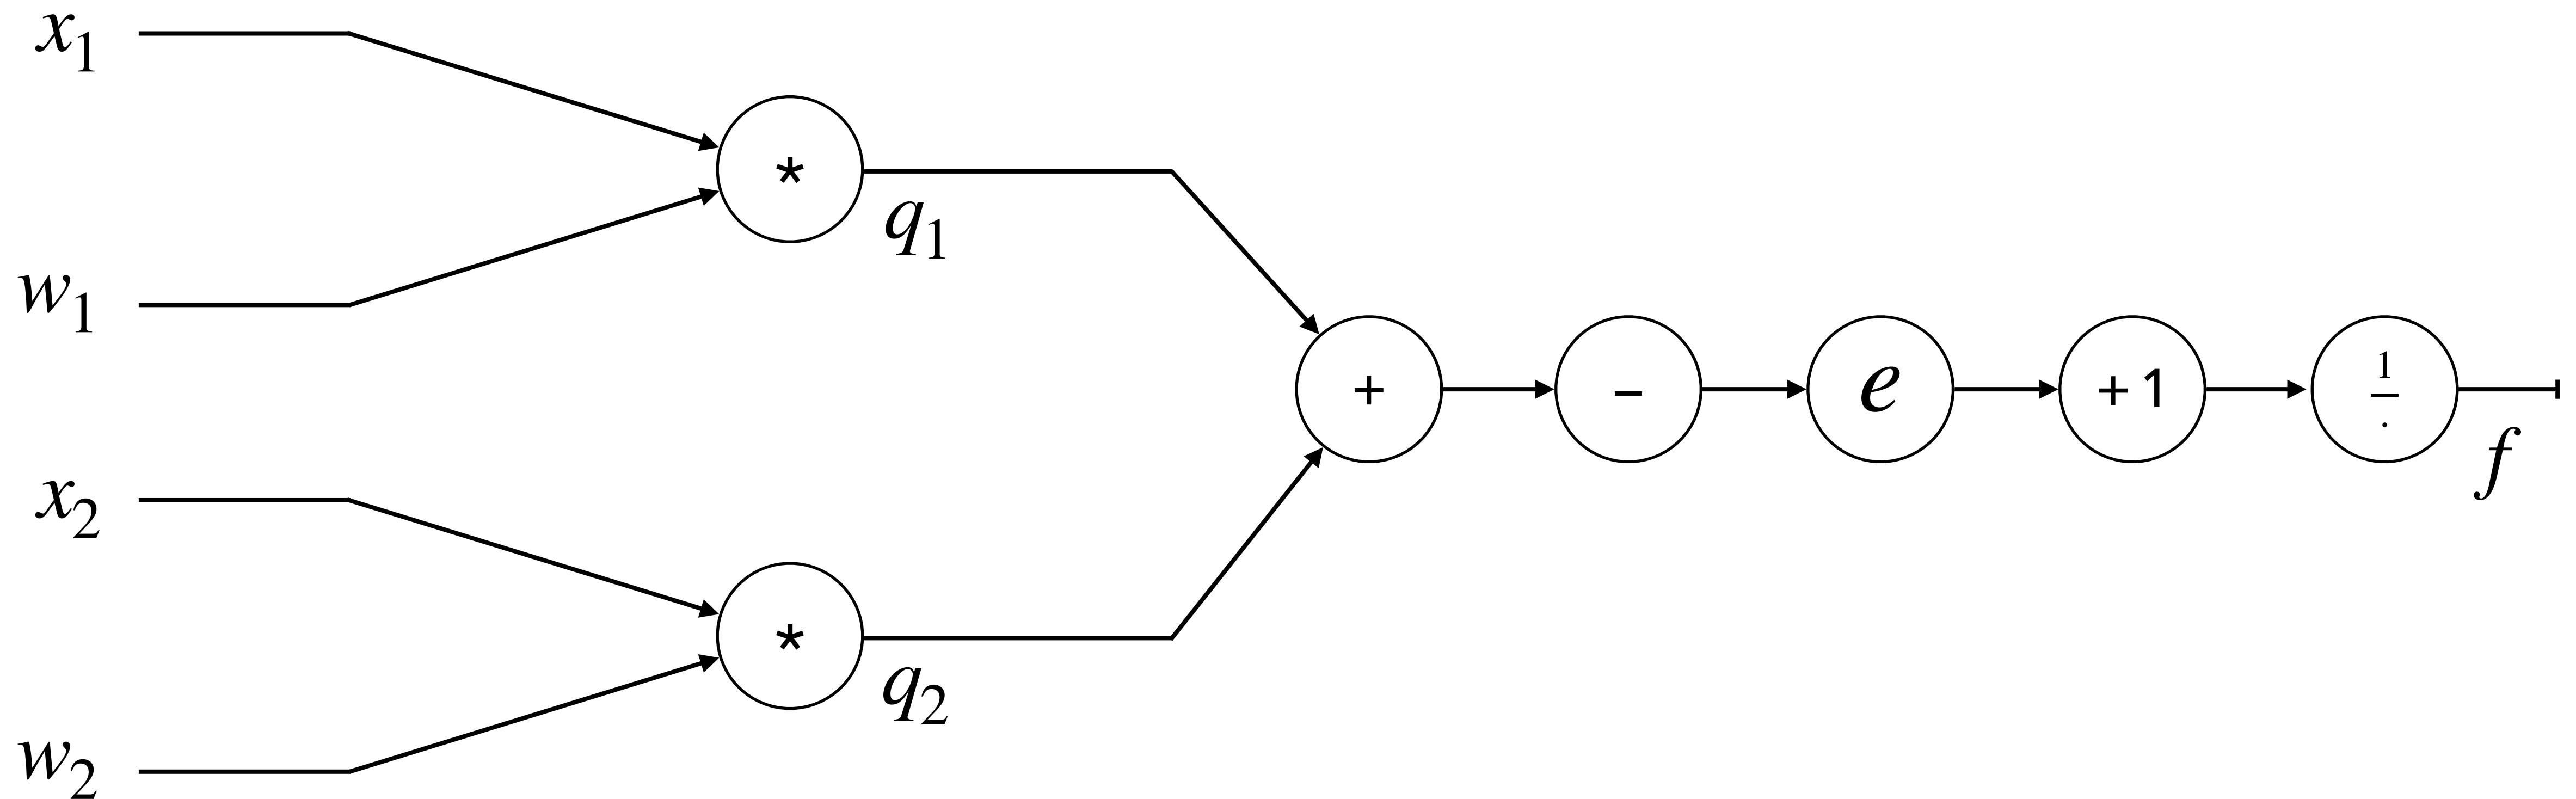


$$w_2 = w_2 - \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial f}{\partial q_2} \frac{\partial q_2}{\partial w_2}$$

Loss function

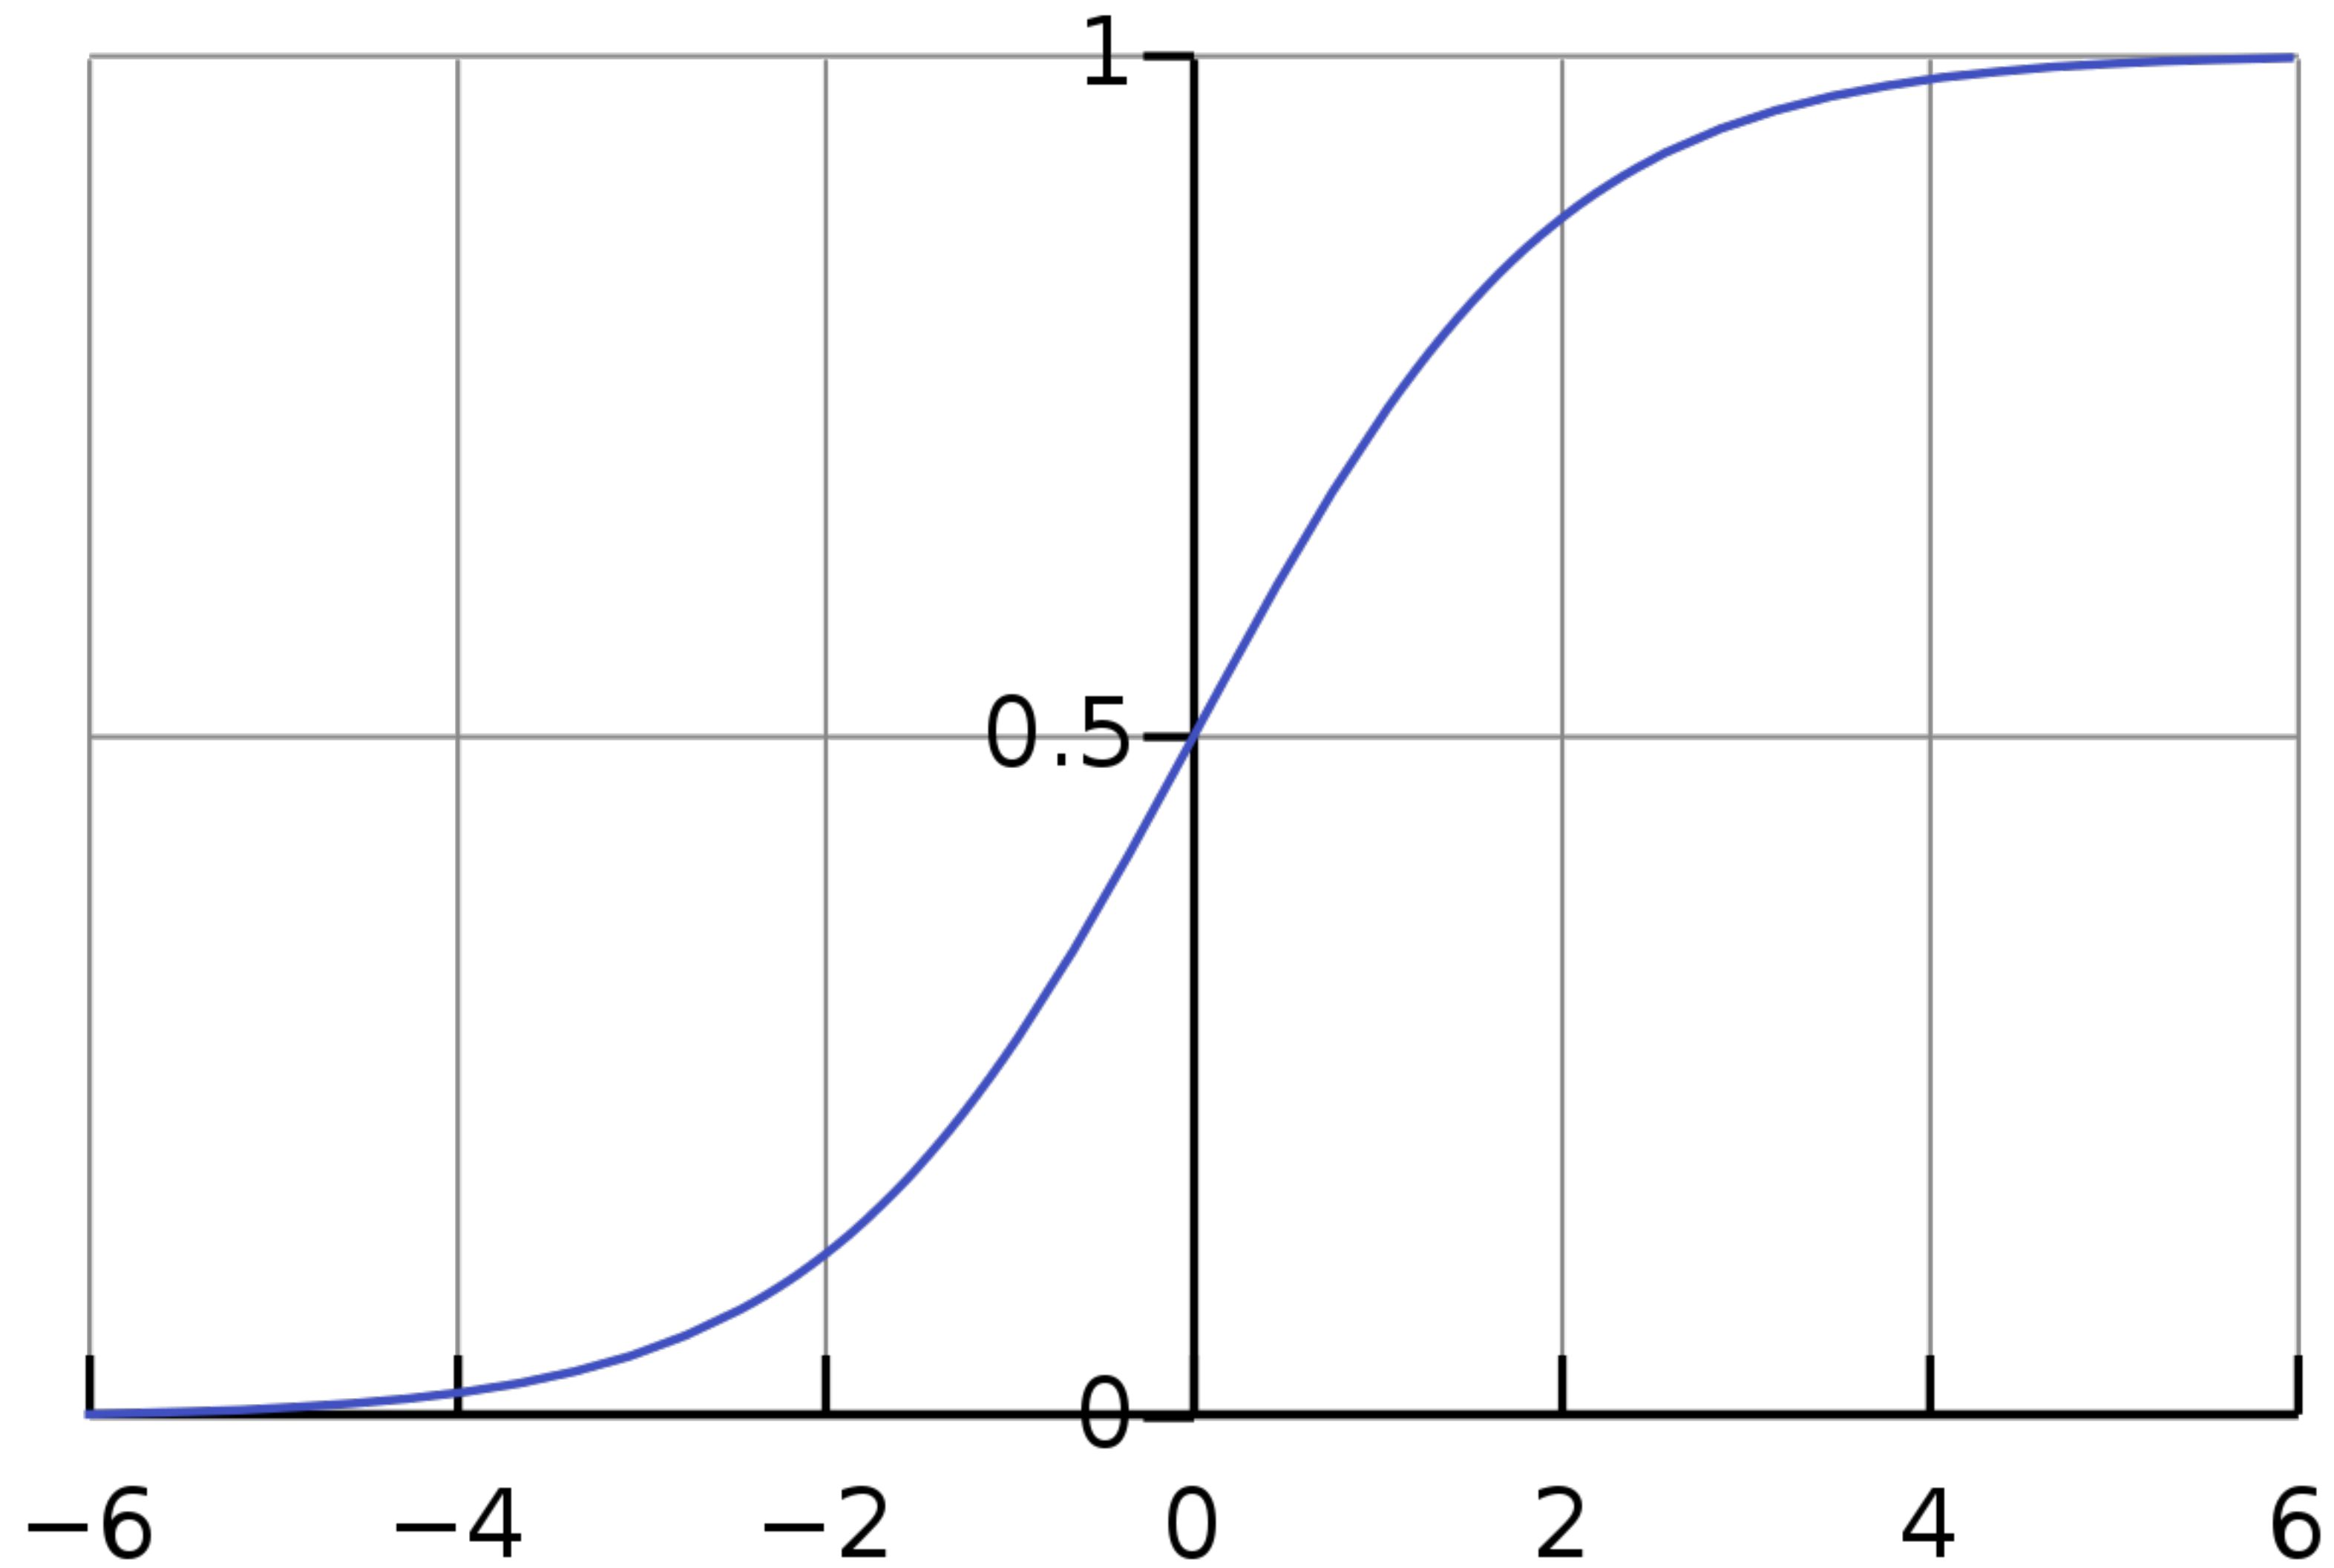
$$J(w) = \hat{y} - y$$

LOGISTIC REGRESSION

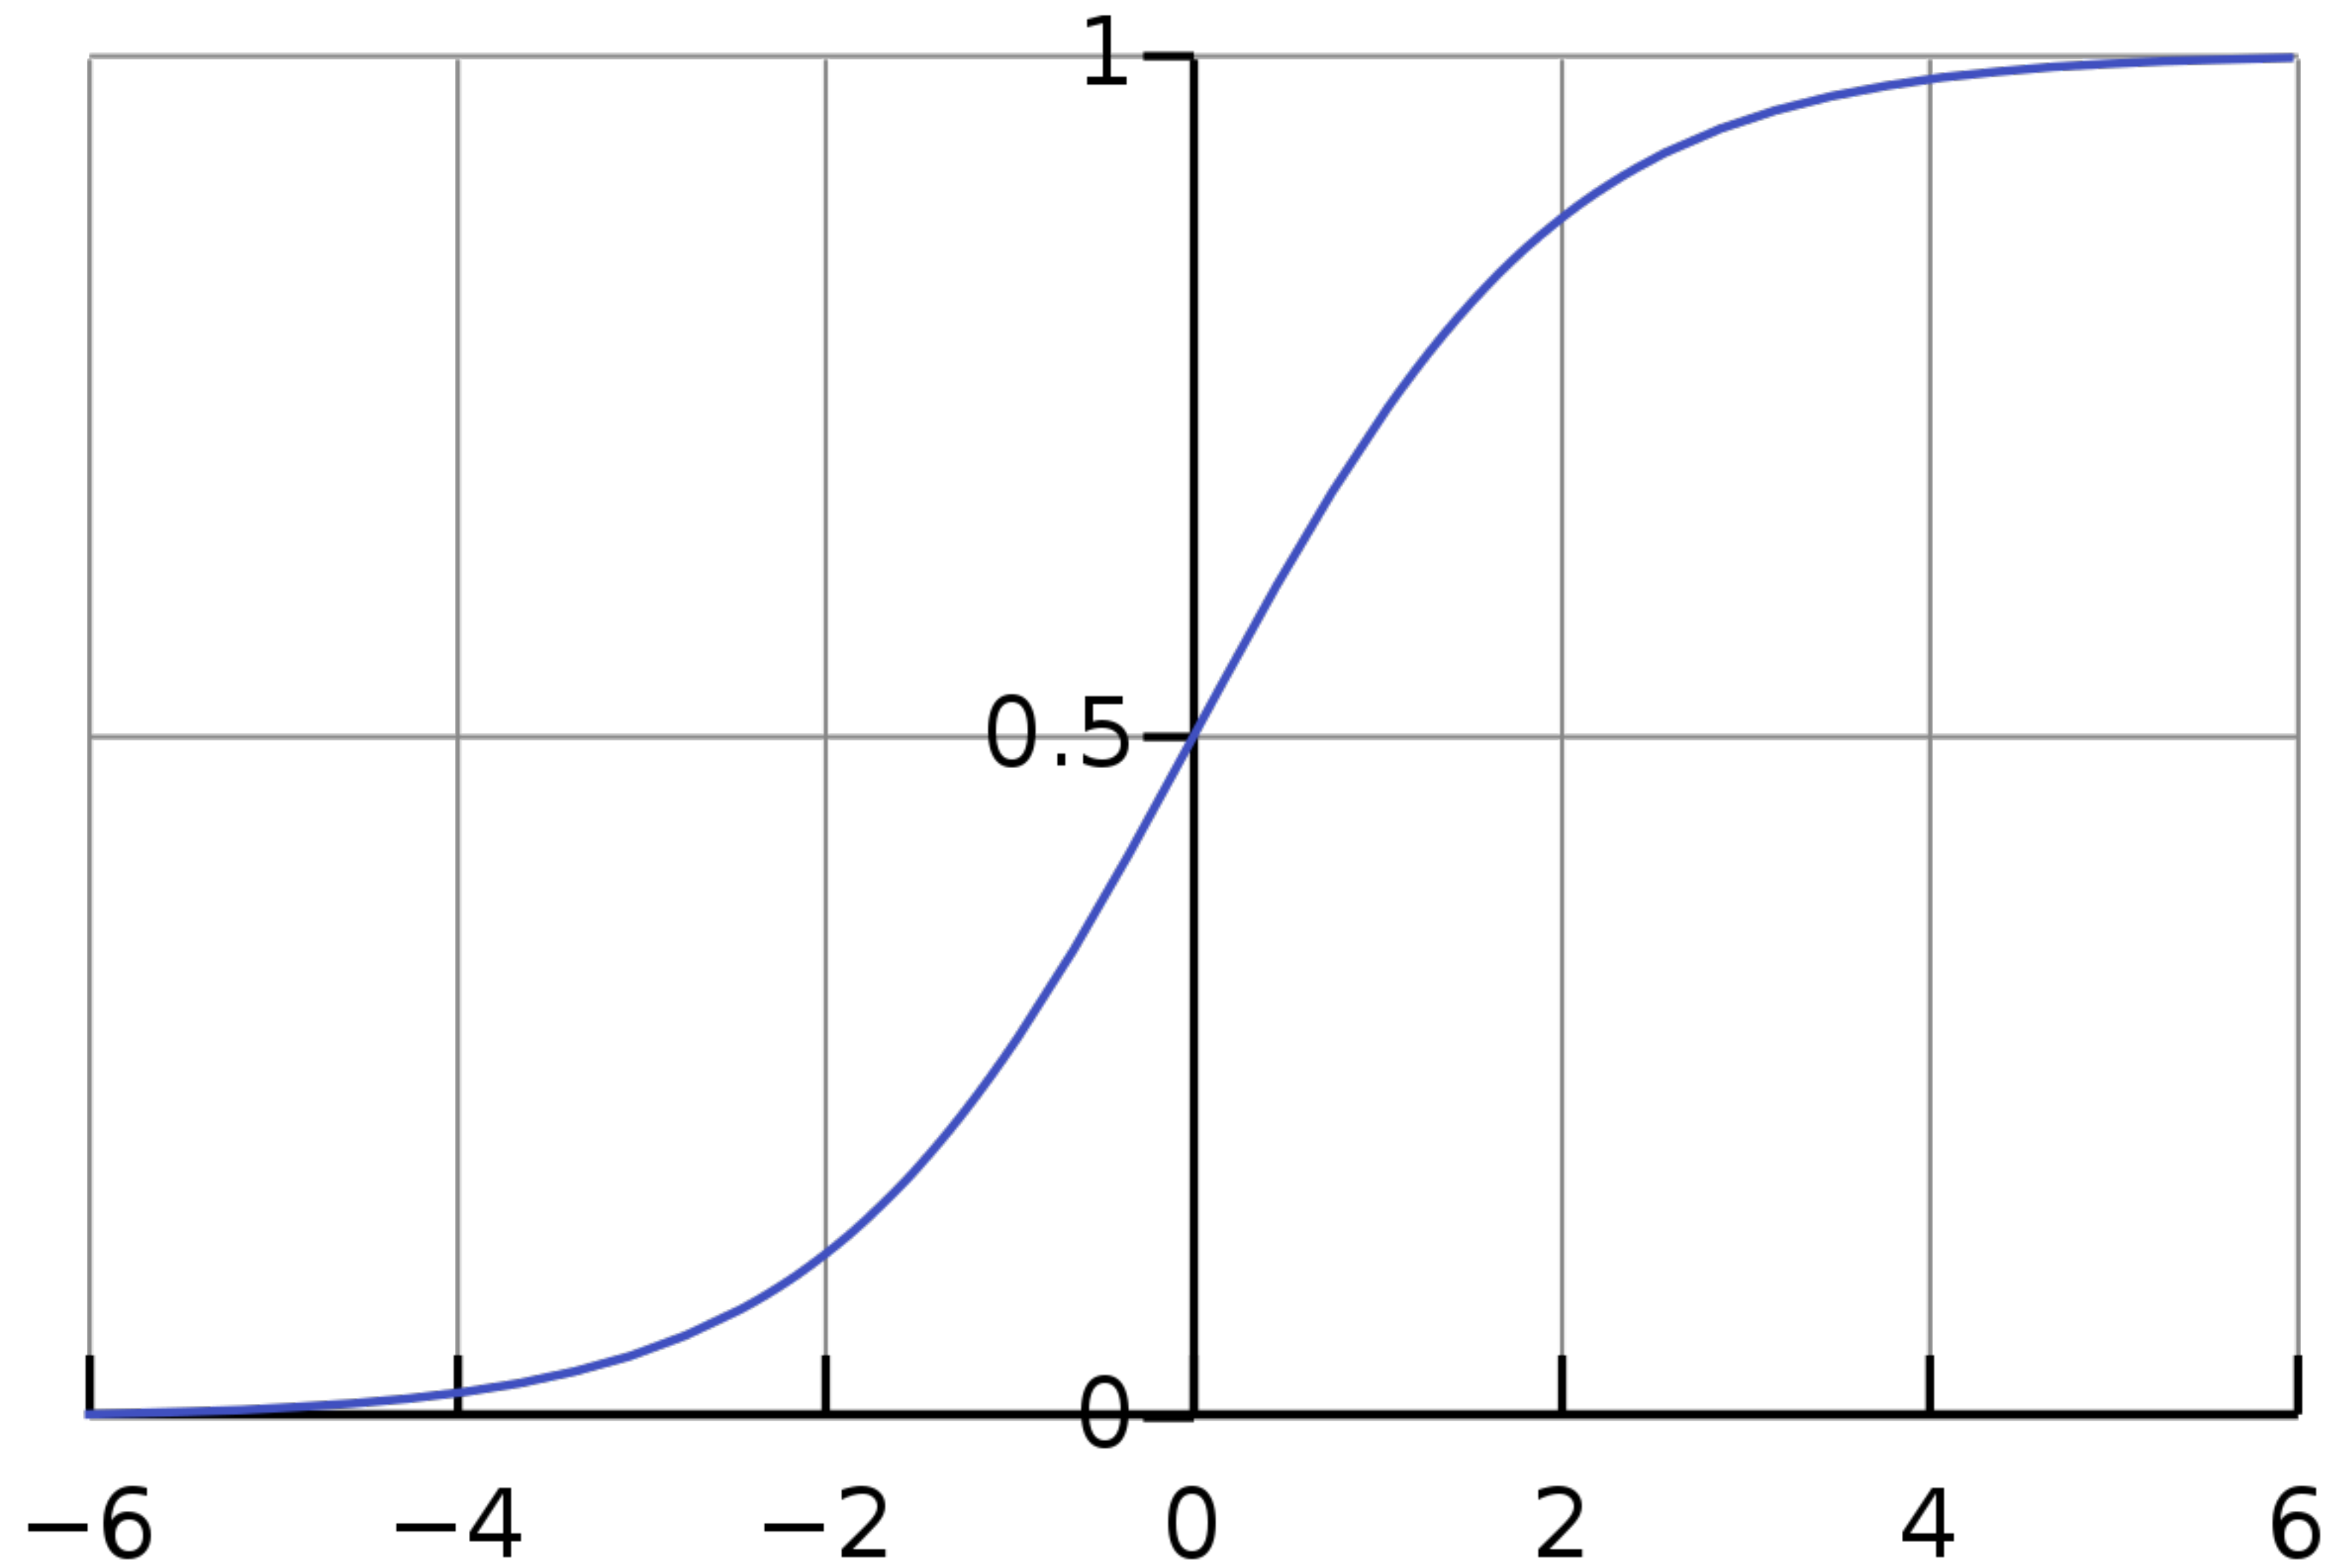


$$f = \frac{1}{1 + e^{-(x_1 w_1 + x_2 w_2)}}$$

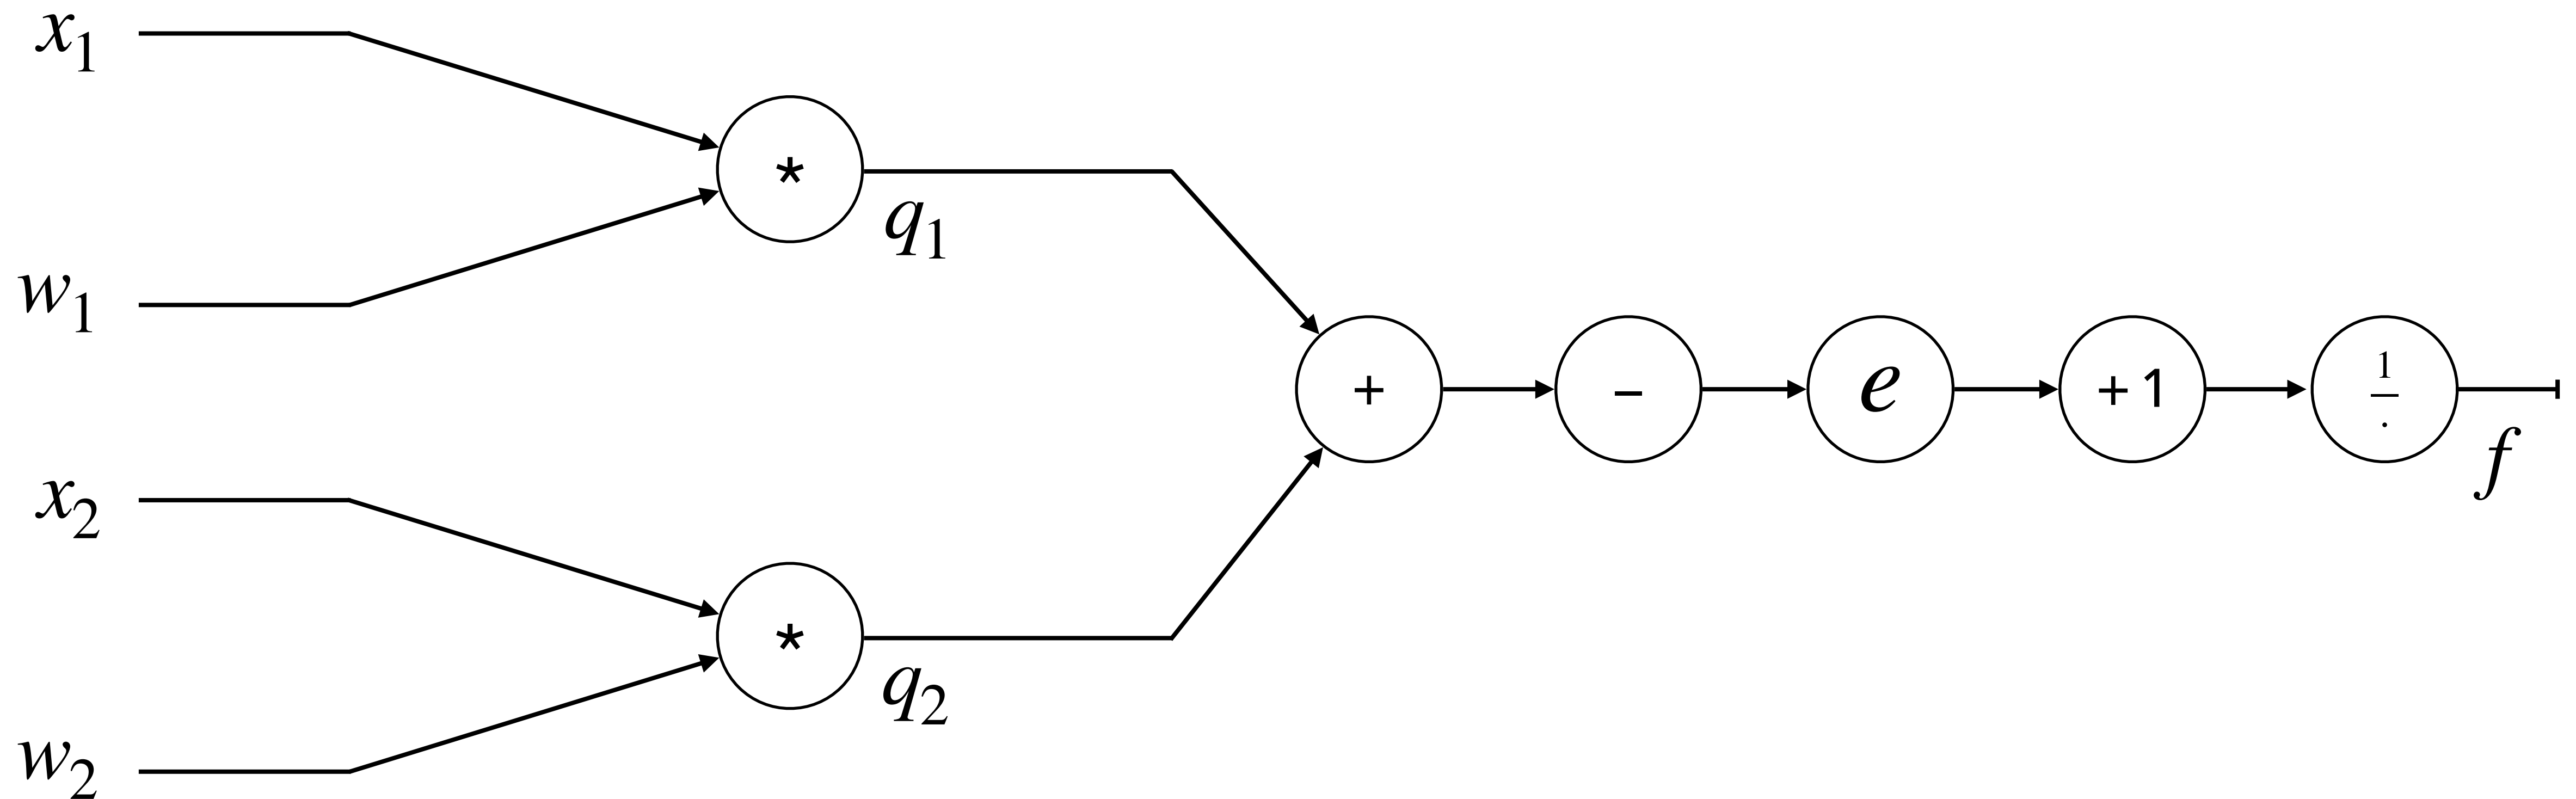
LOGISTIC REGRESSION



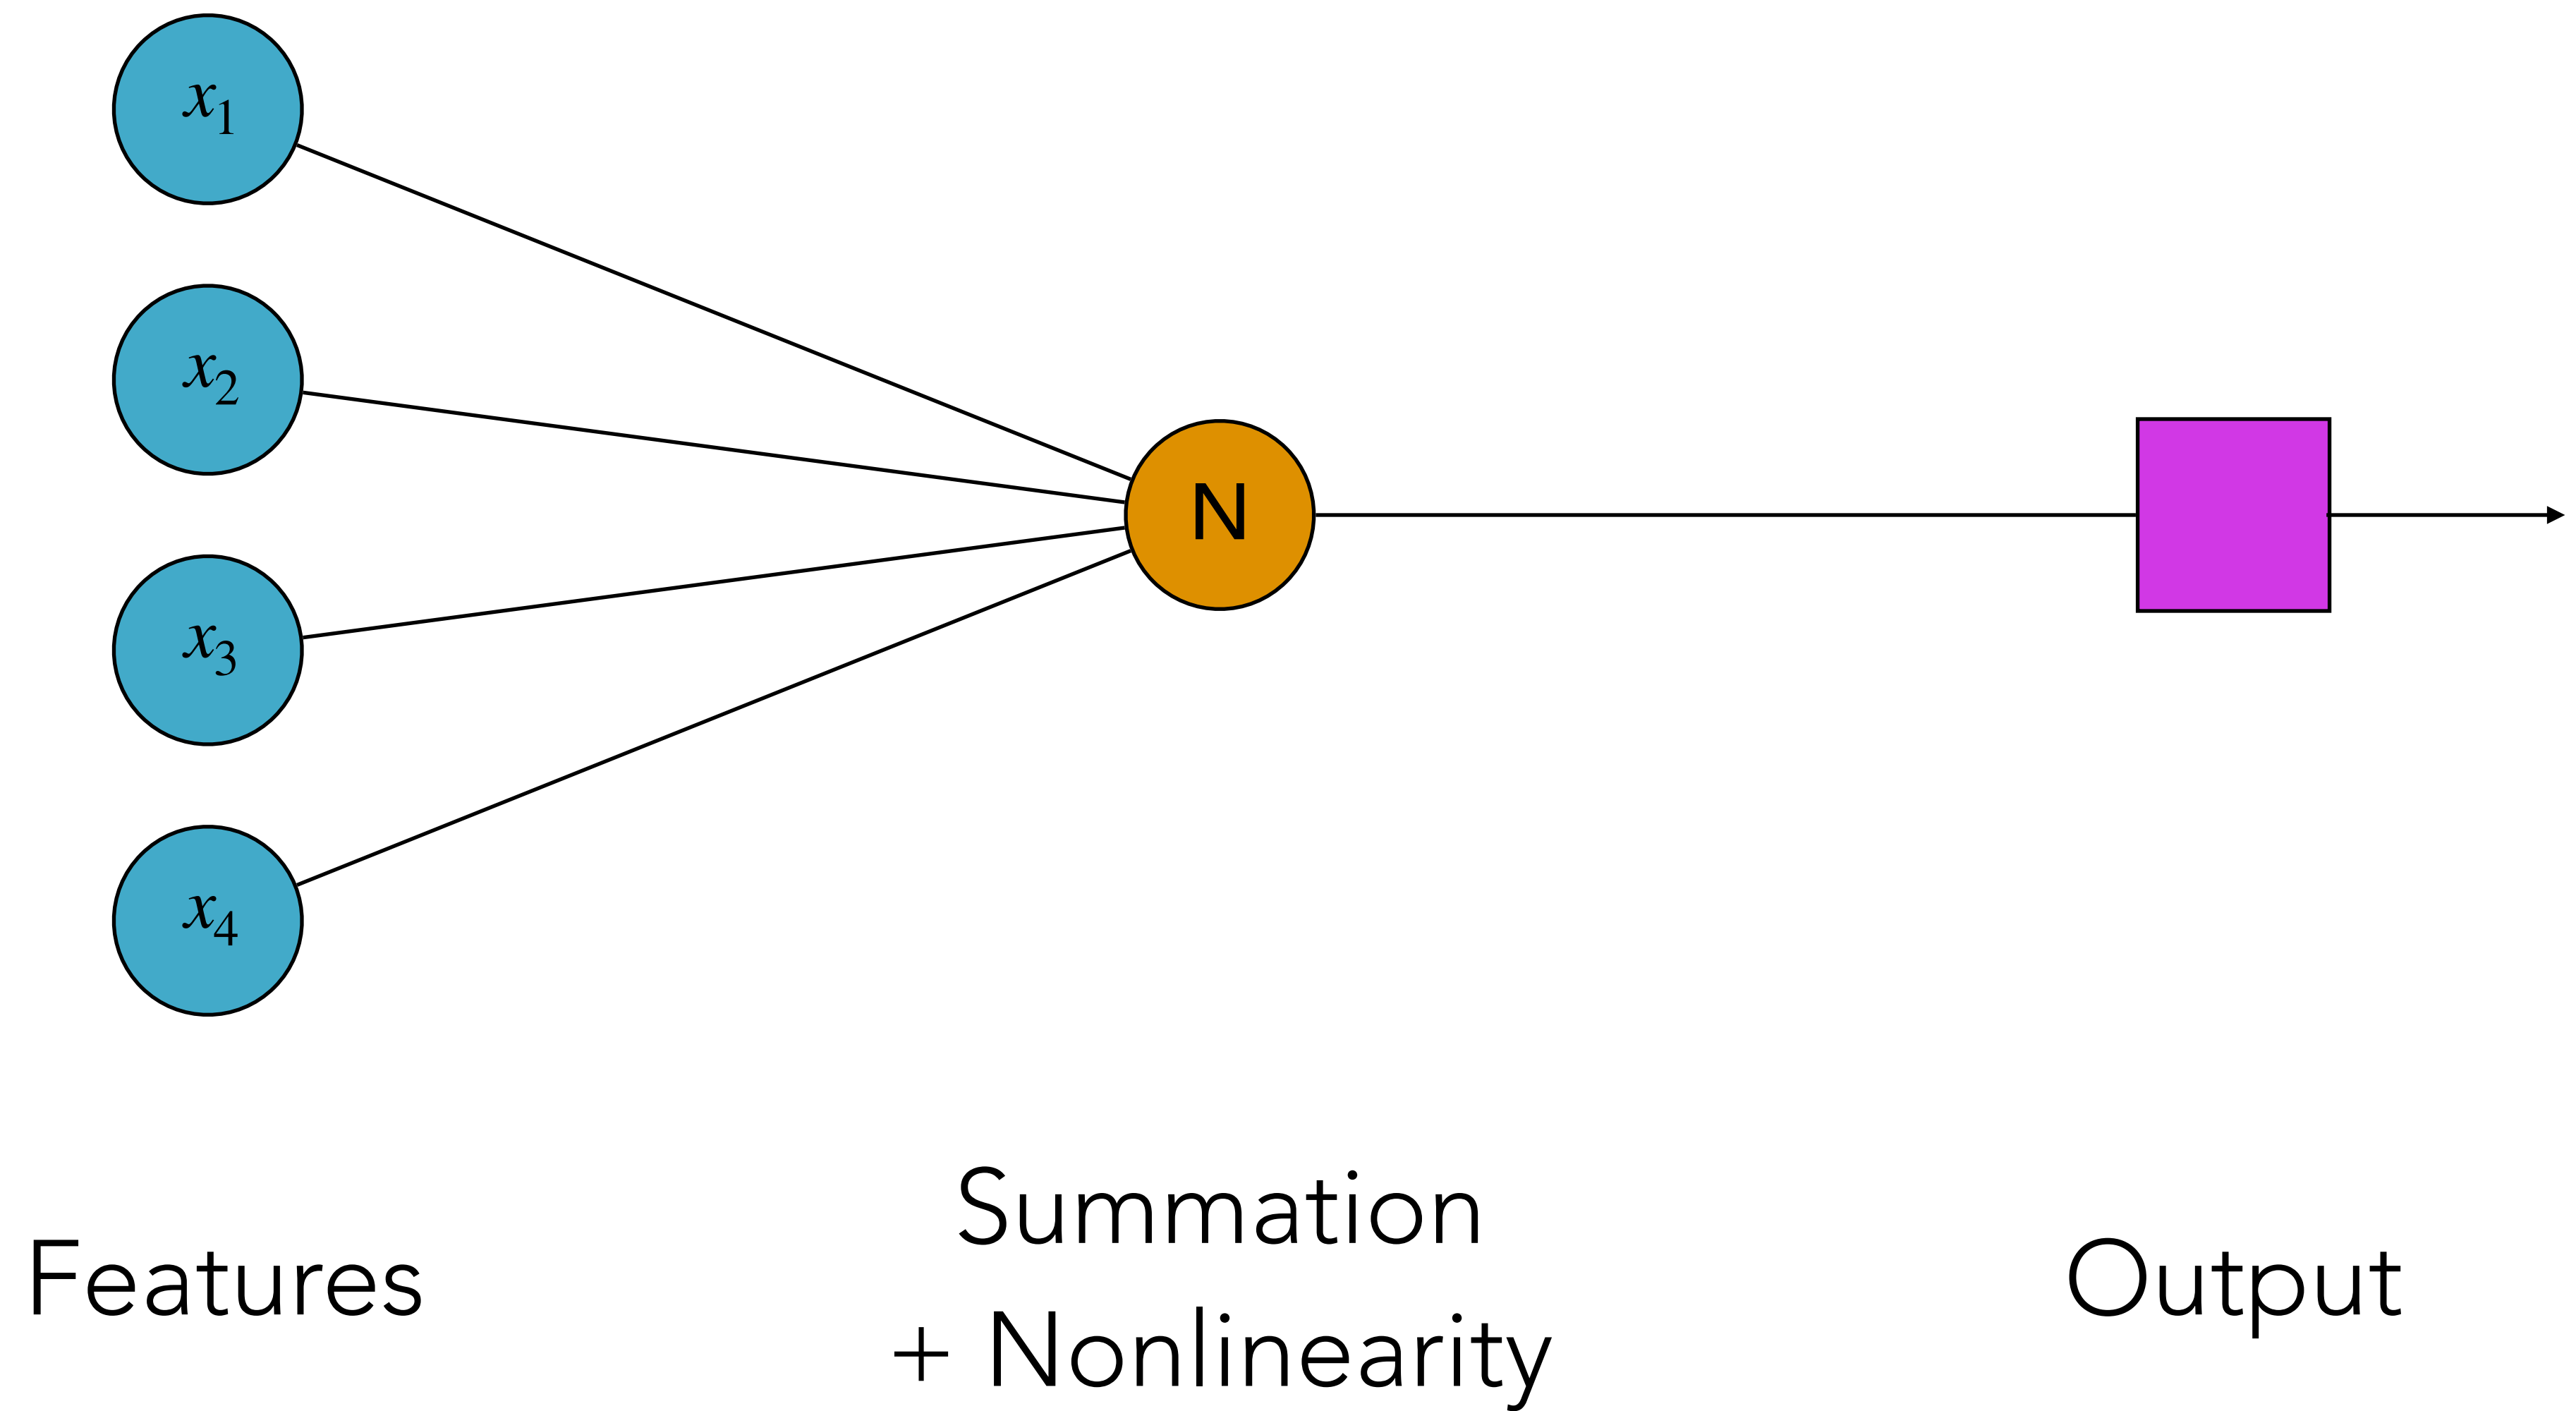
CLASSIFICATION



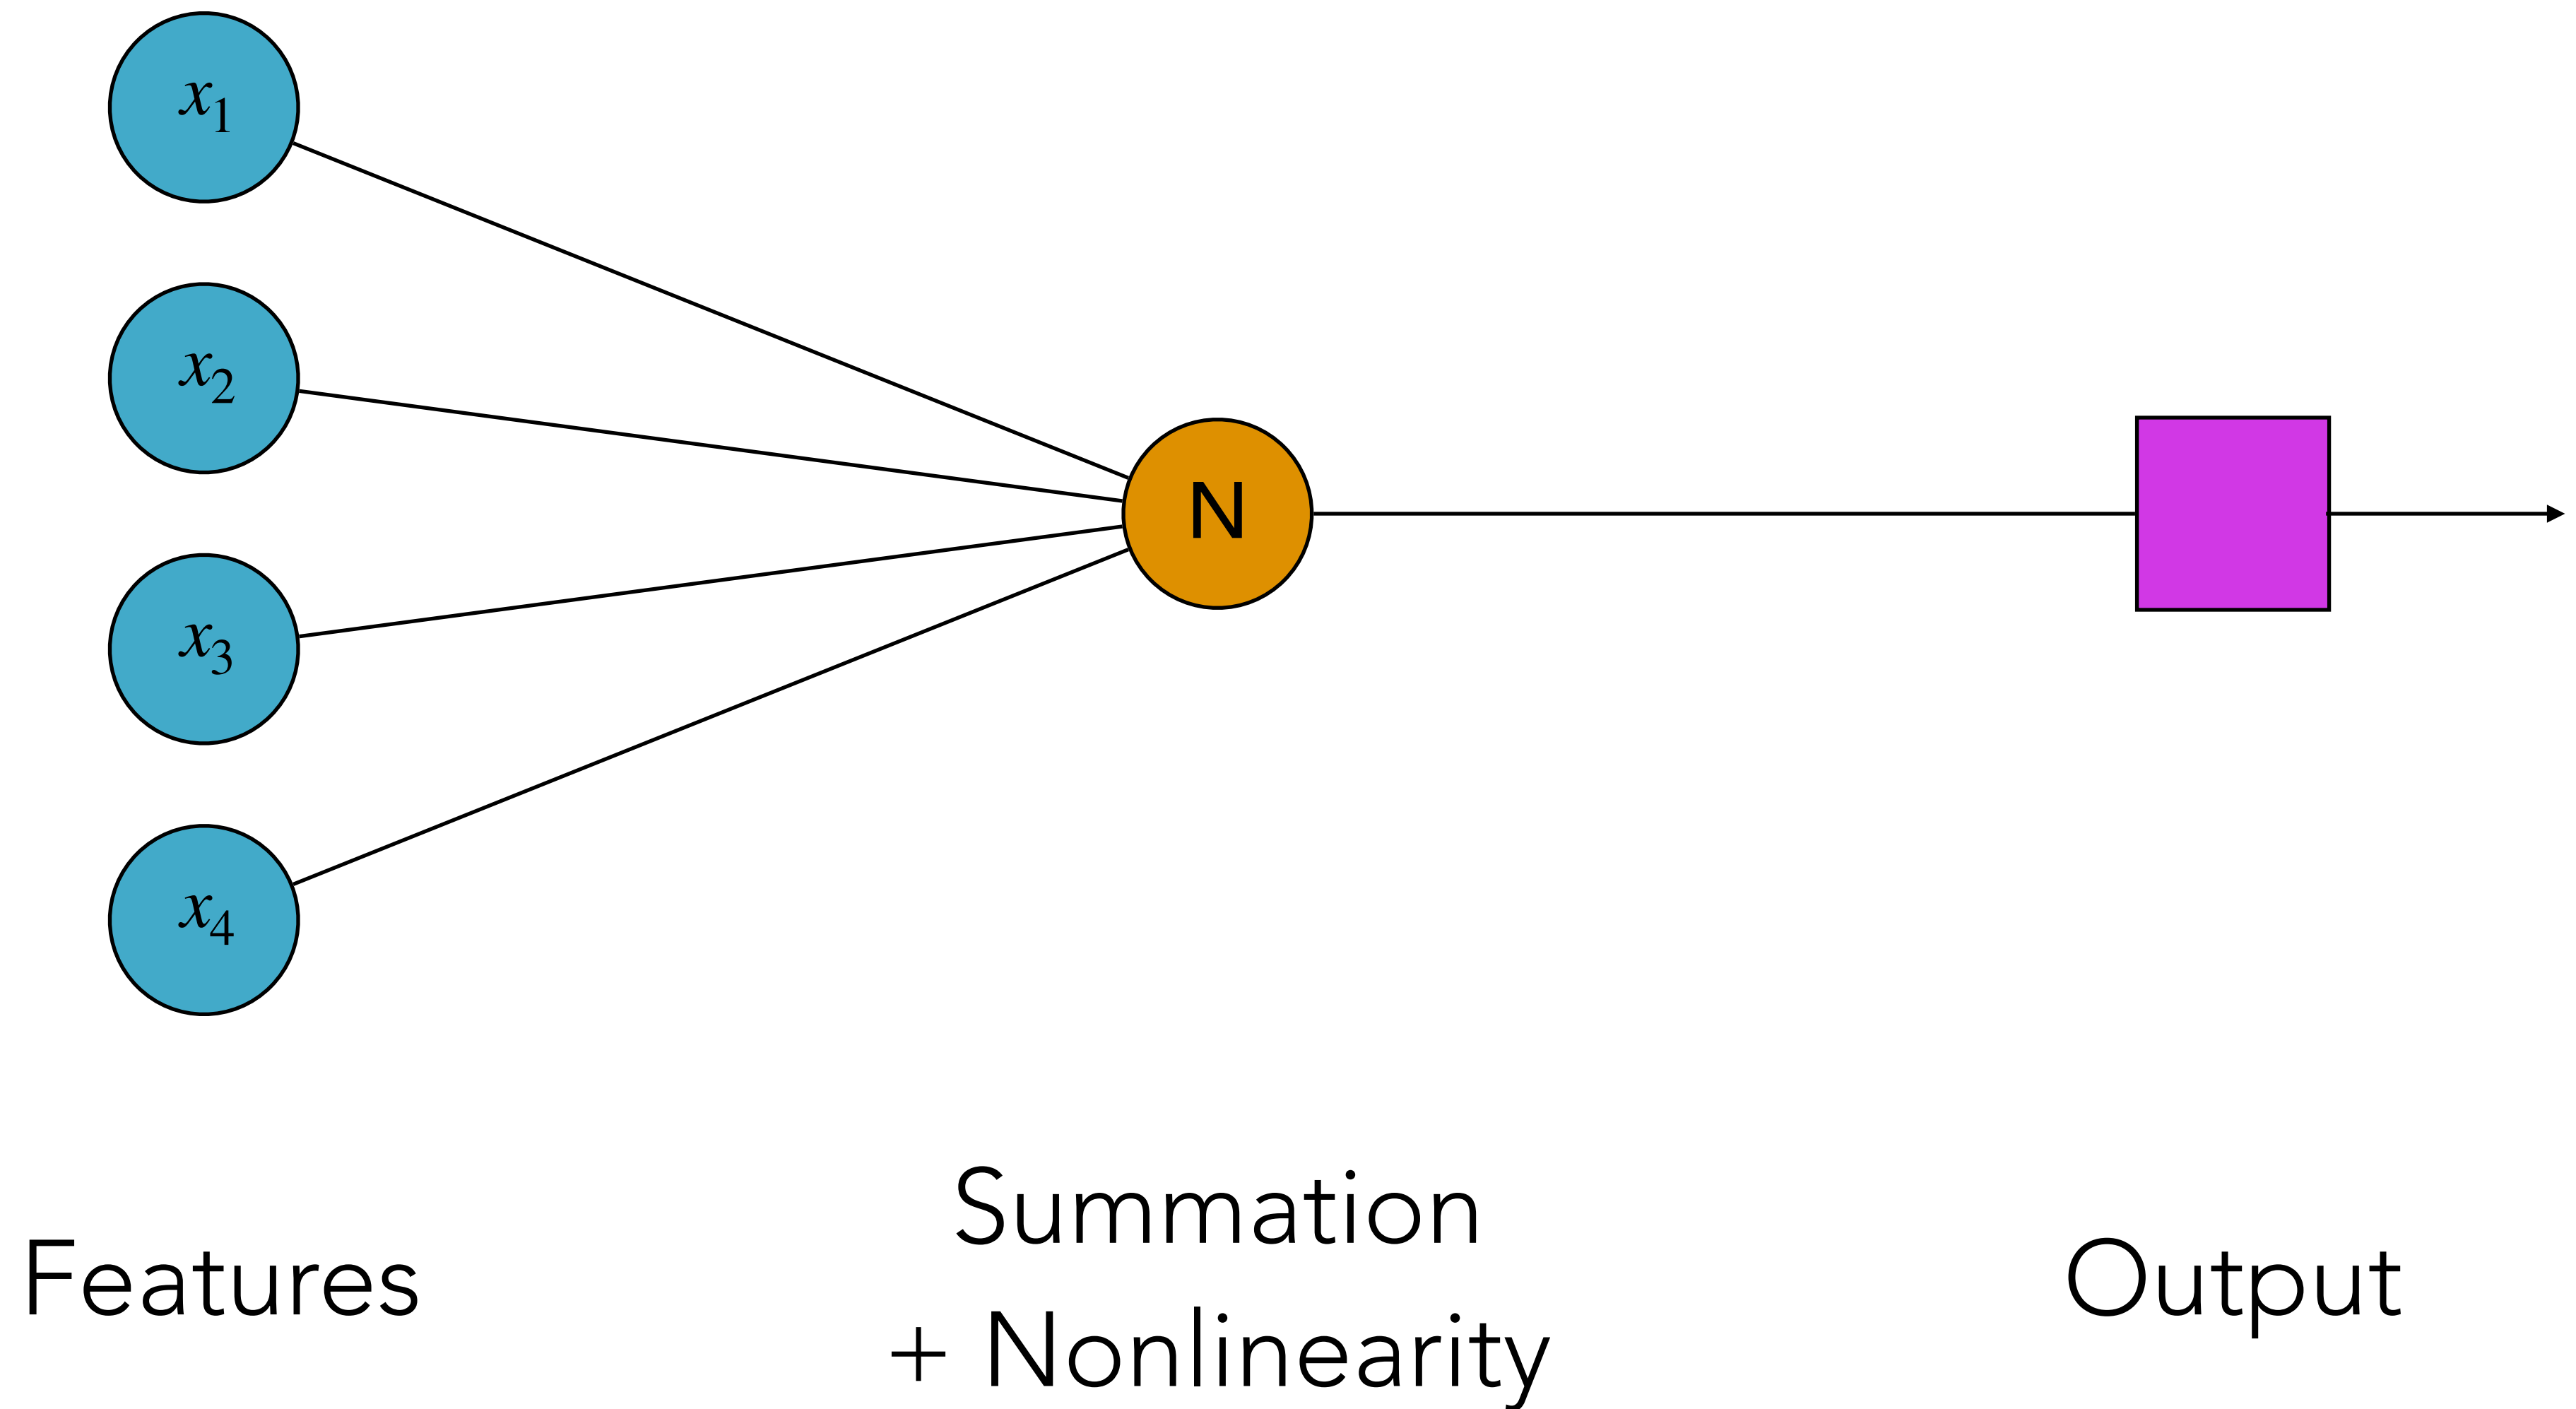
LOGISTIC REGRESSION

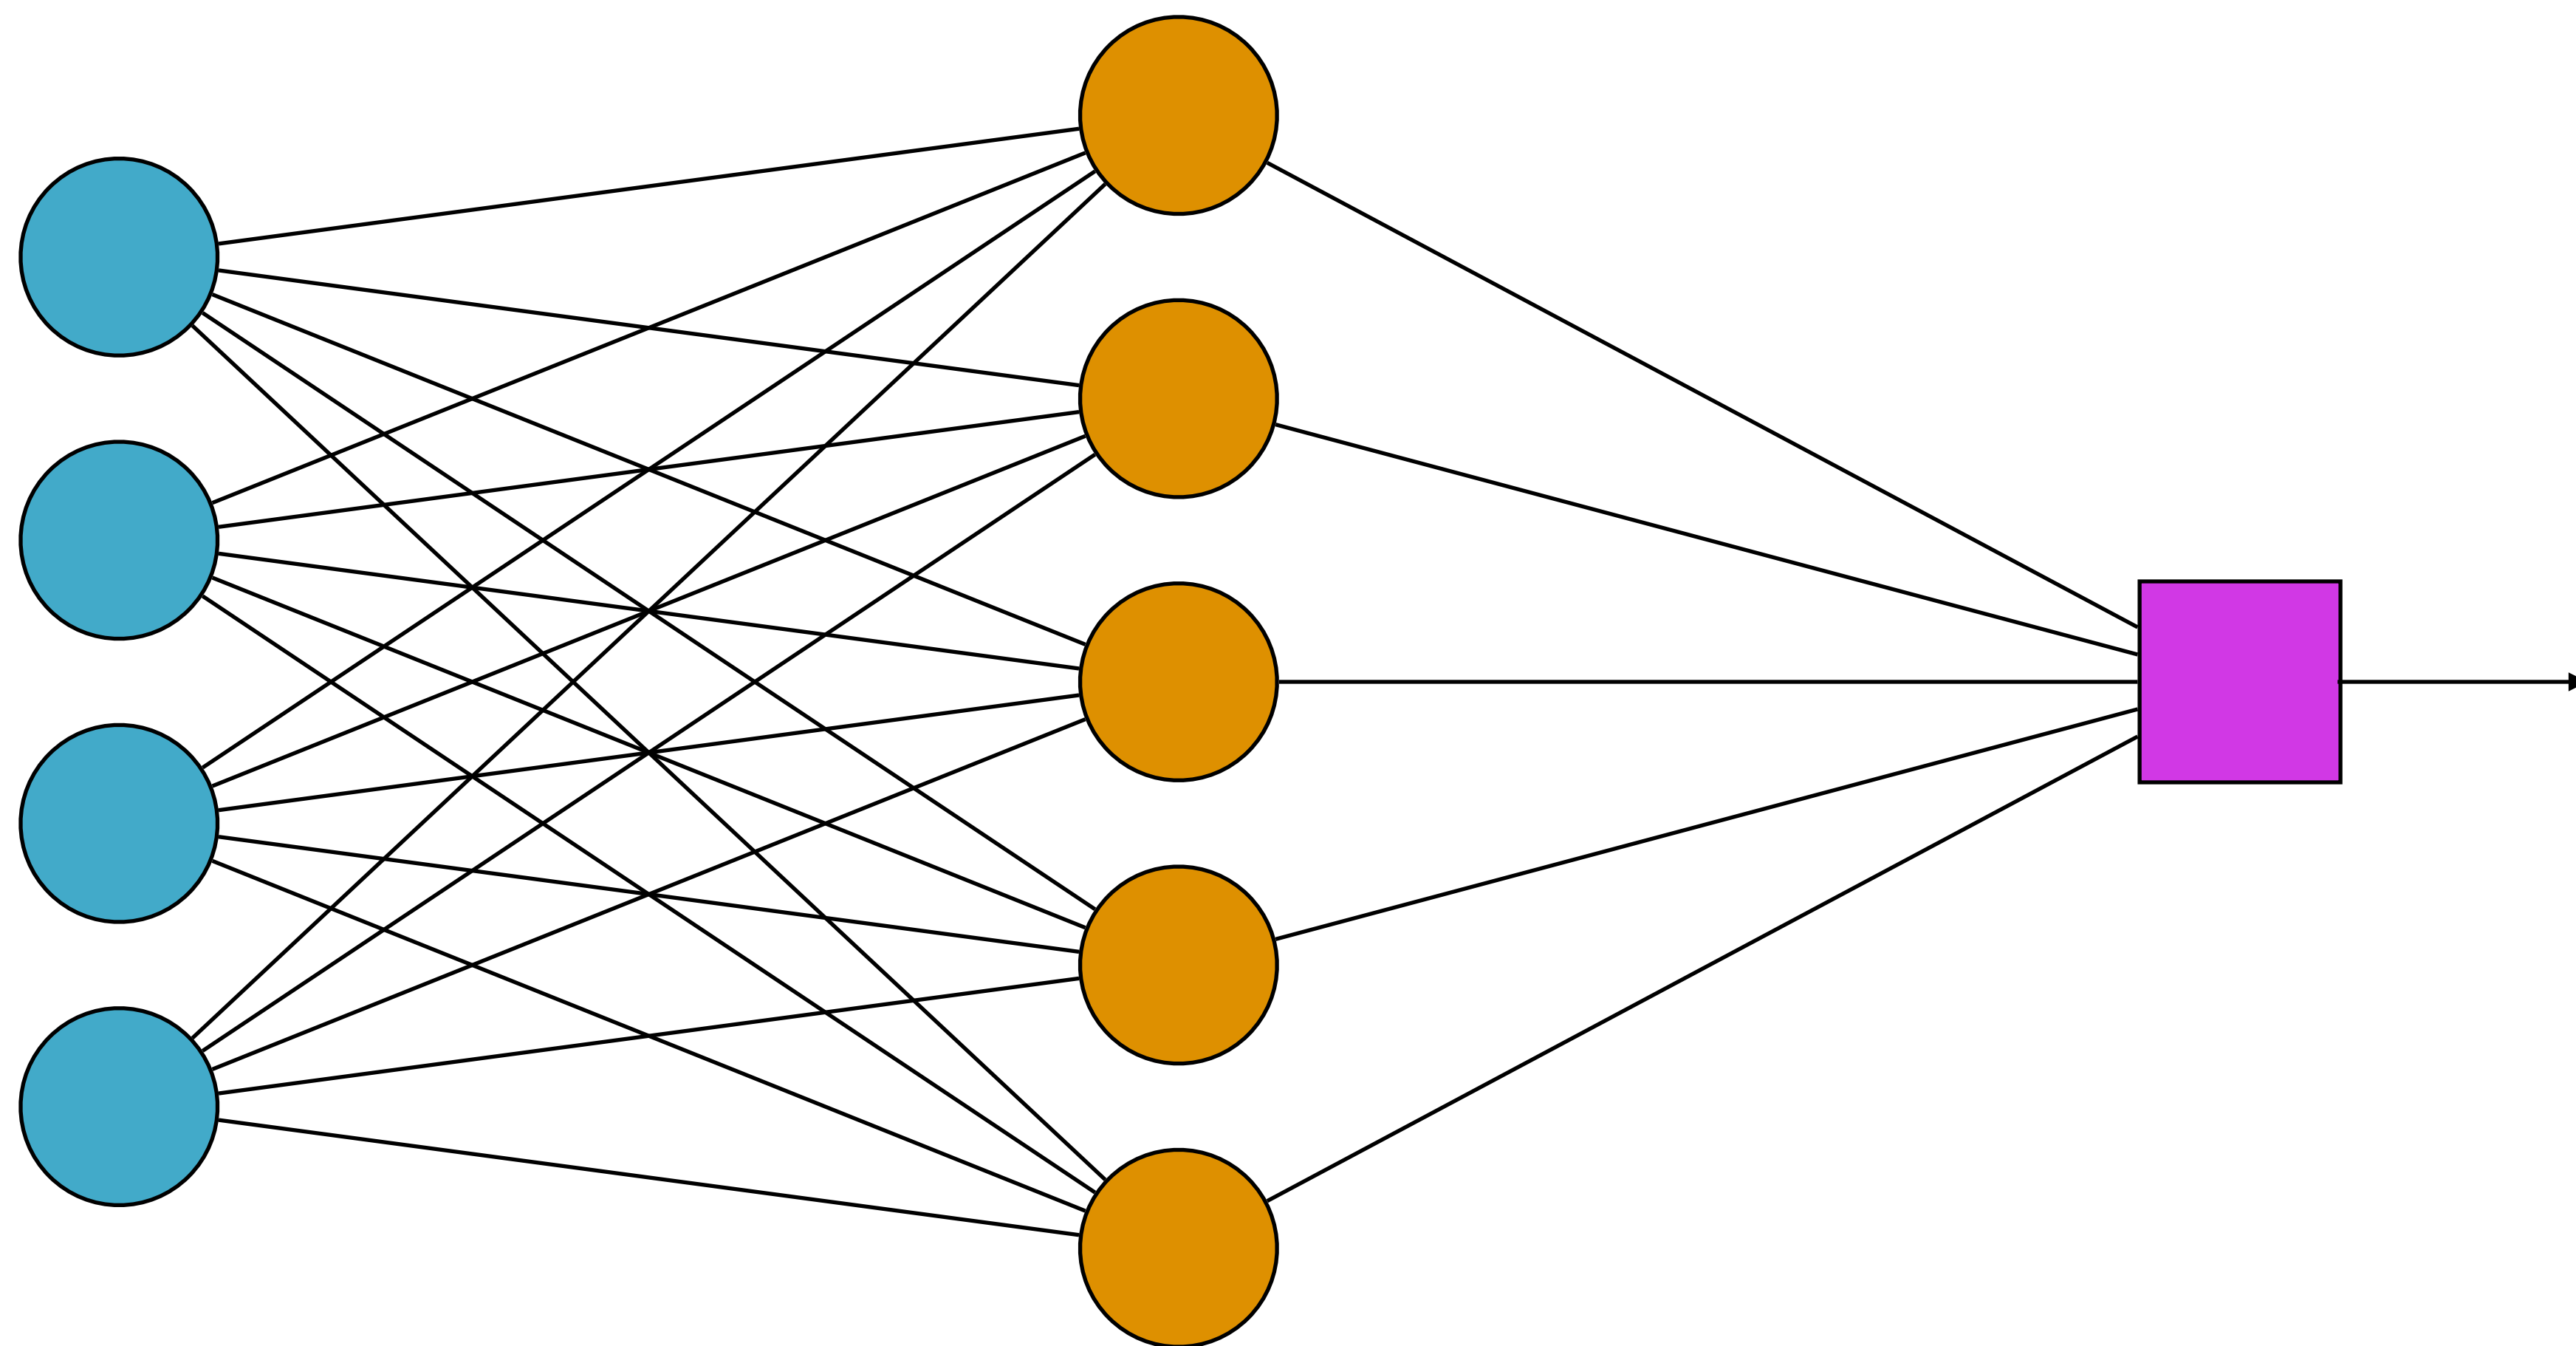


$$f = \frac{1}{1 + e^{-(x_1 w_1 + x_2 w_2)}}$$



CHECK: HOW MANY TRAINABLE PARAMETERS?



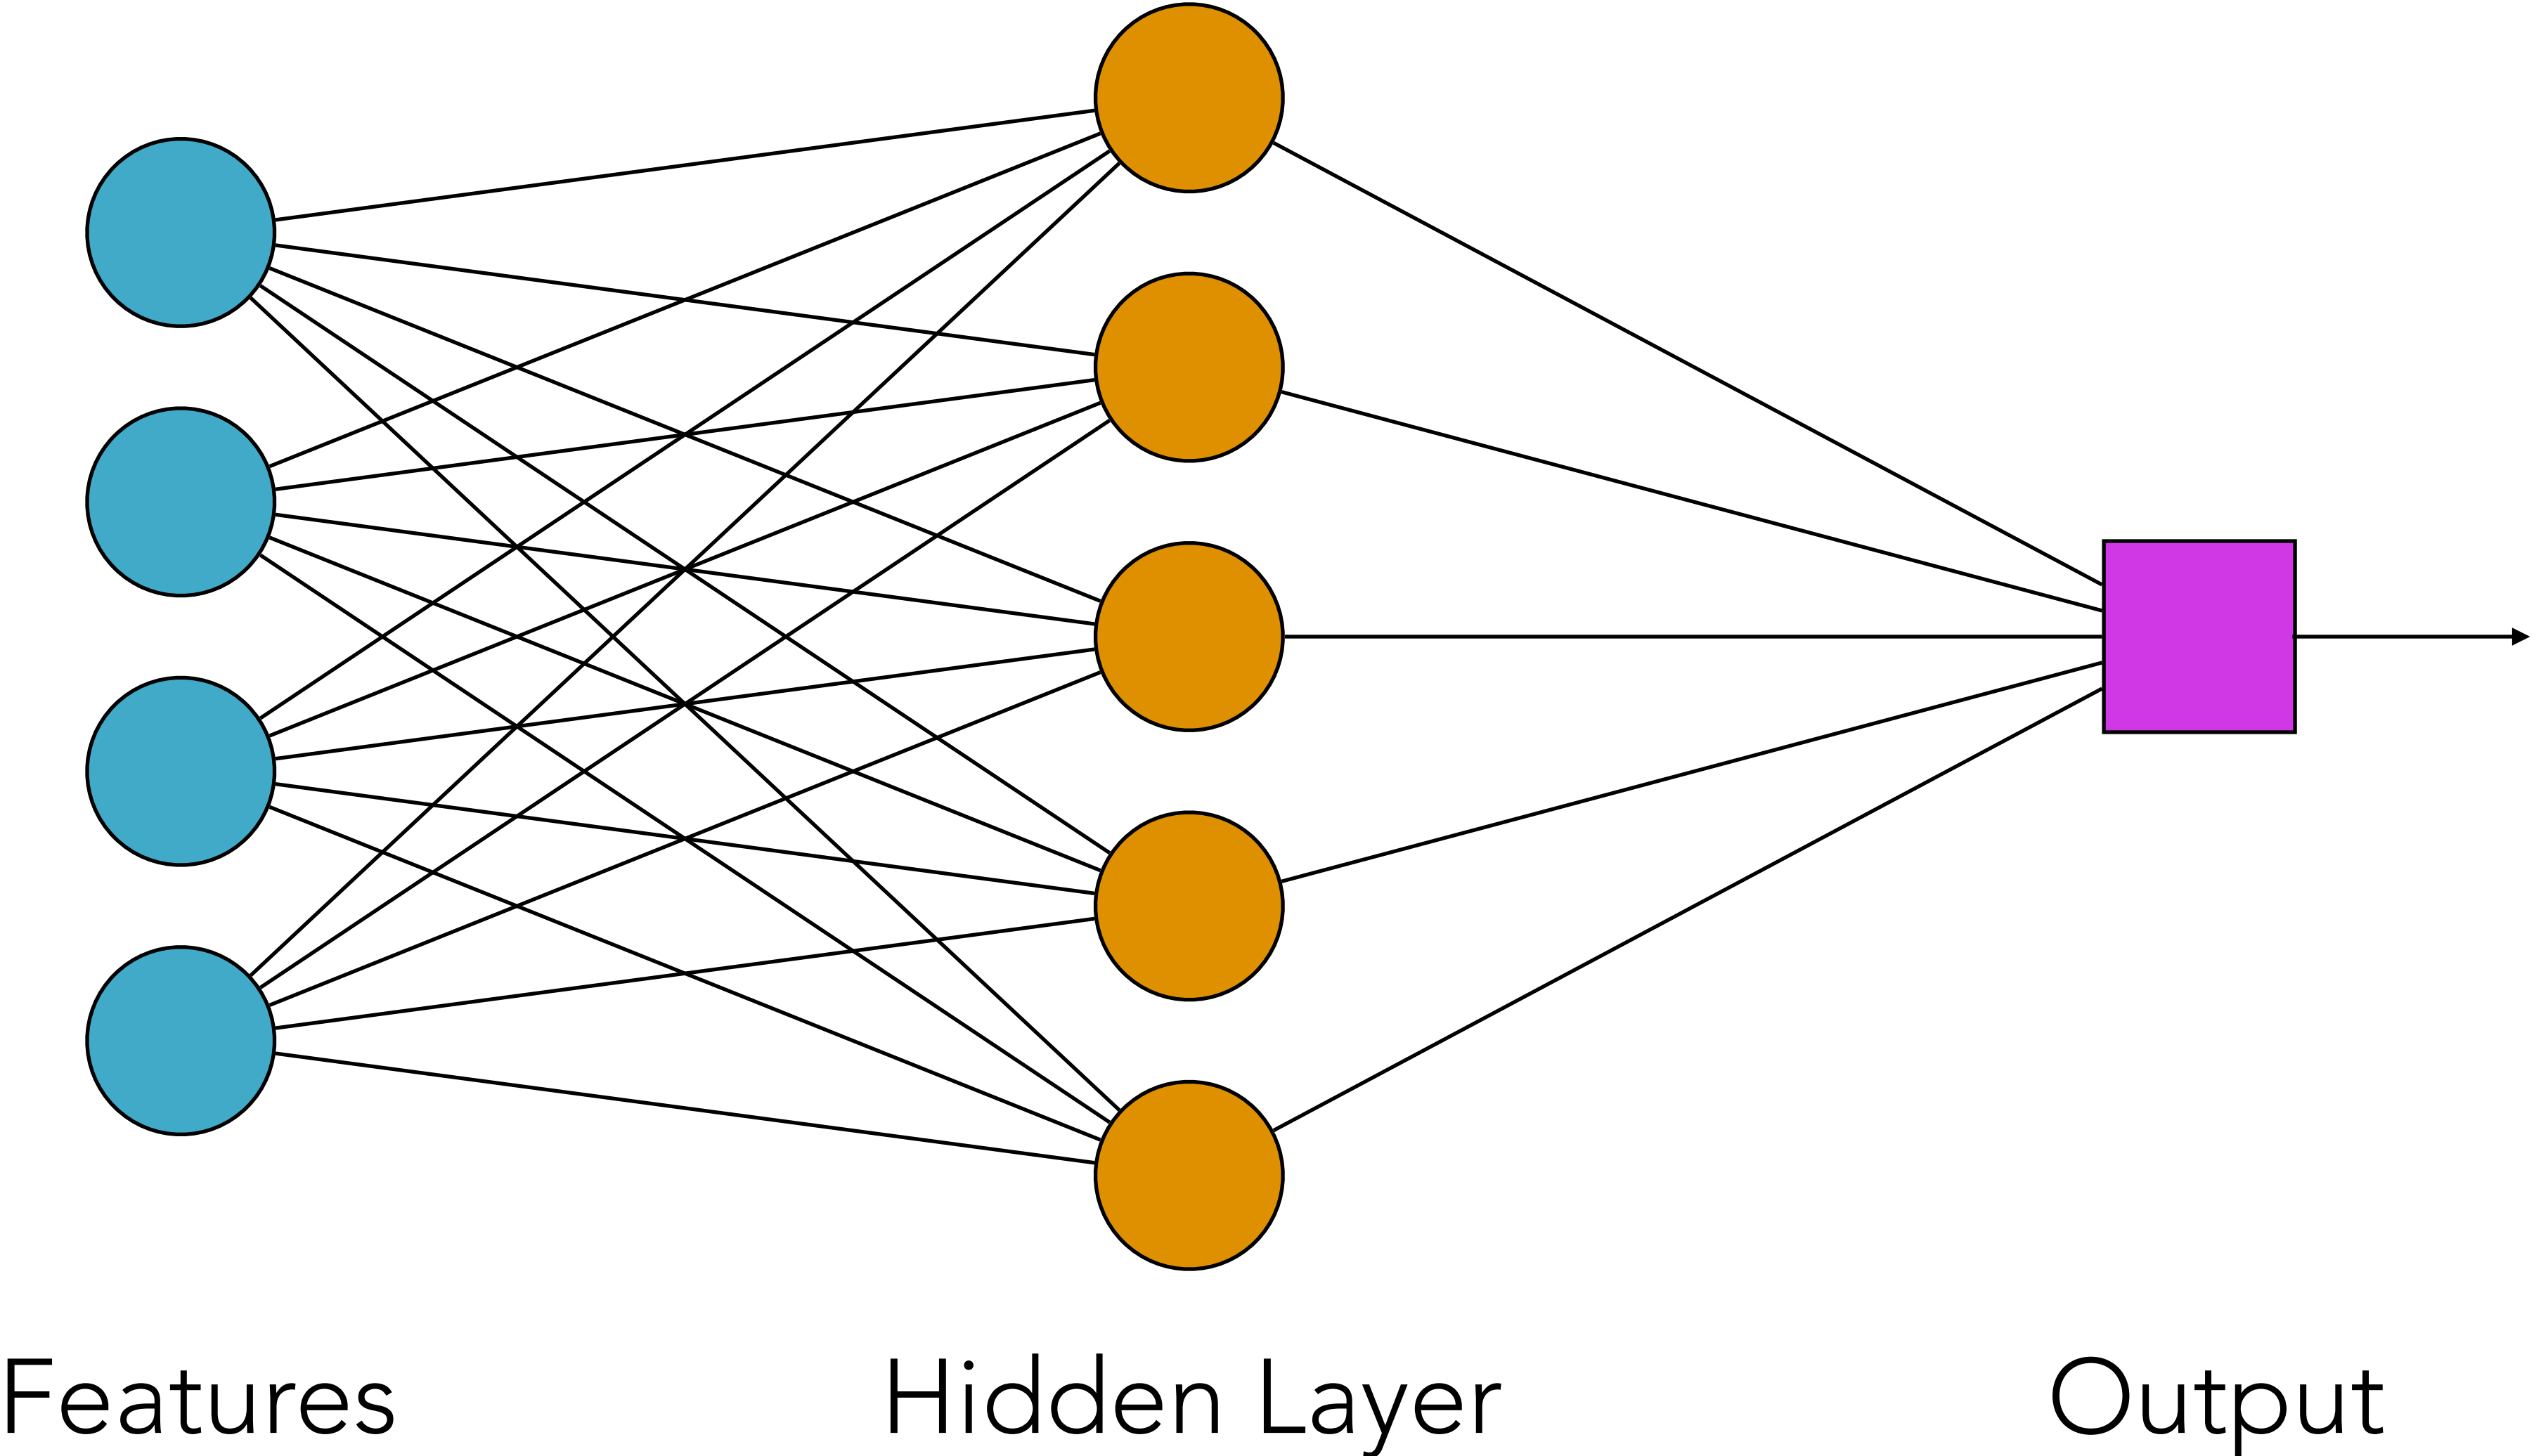


Features

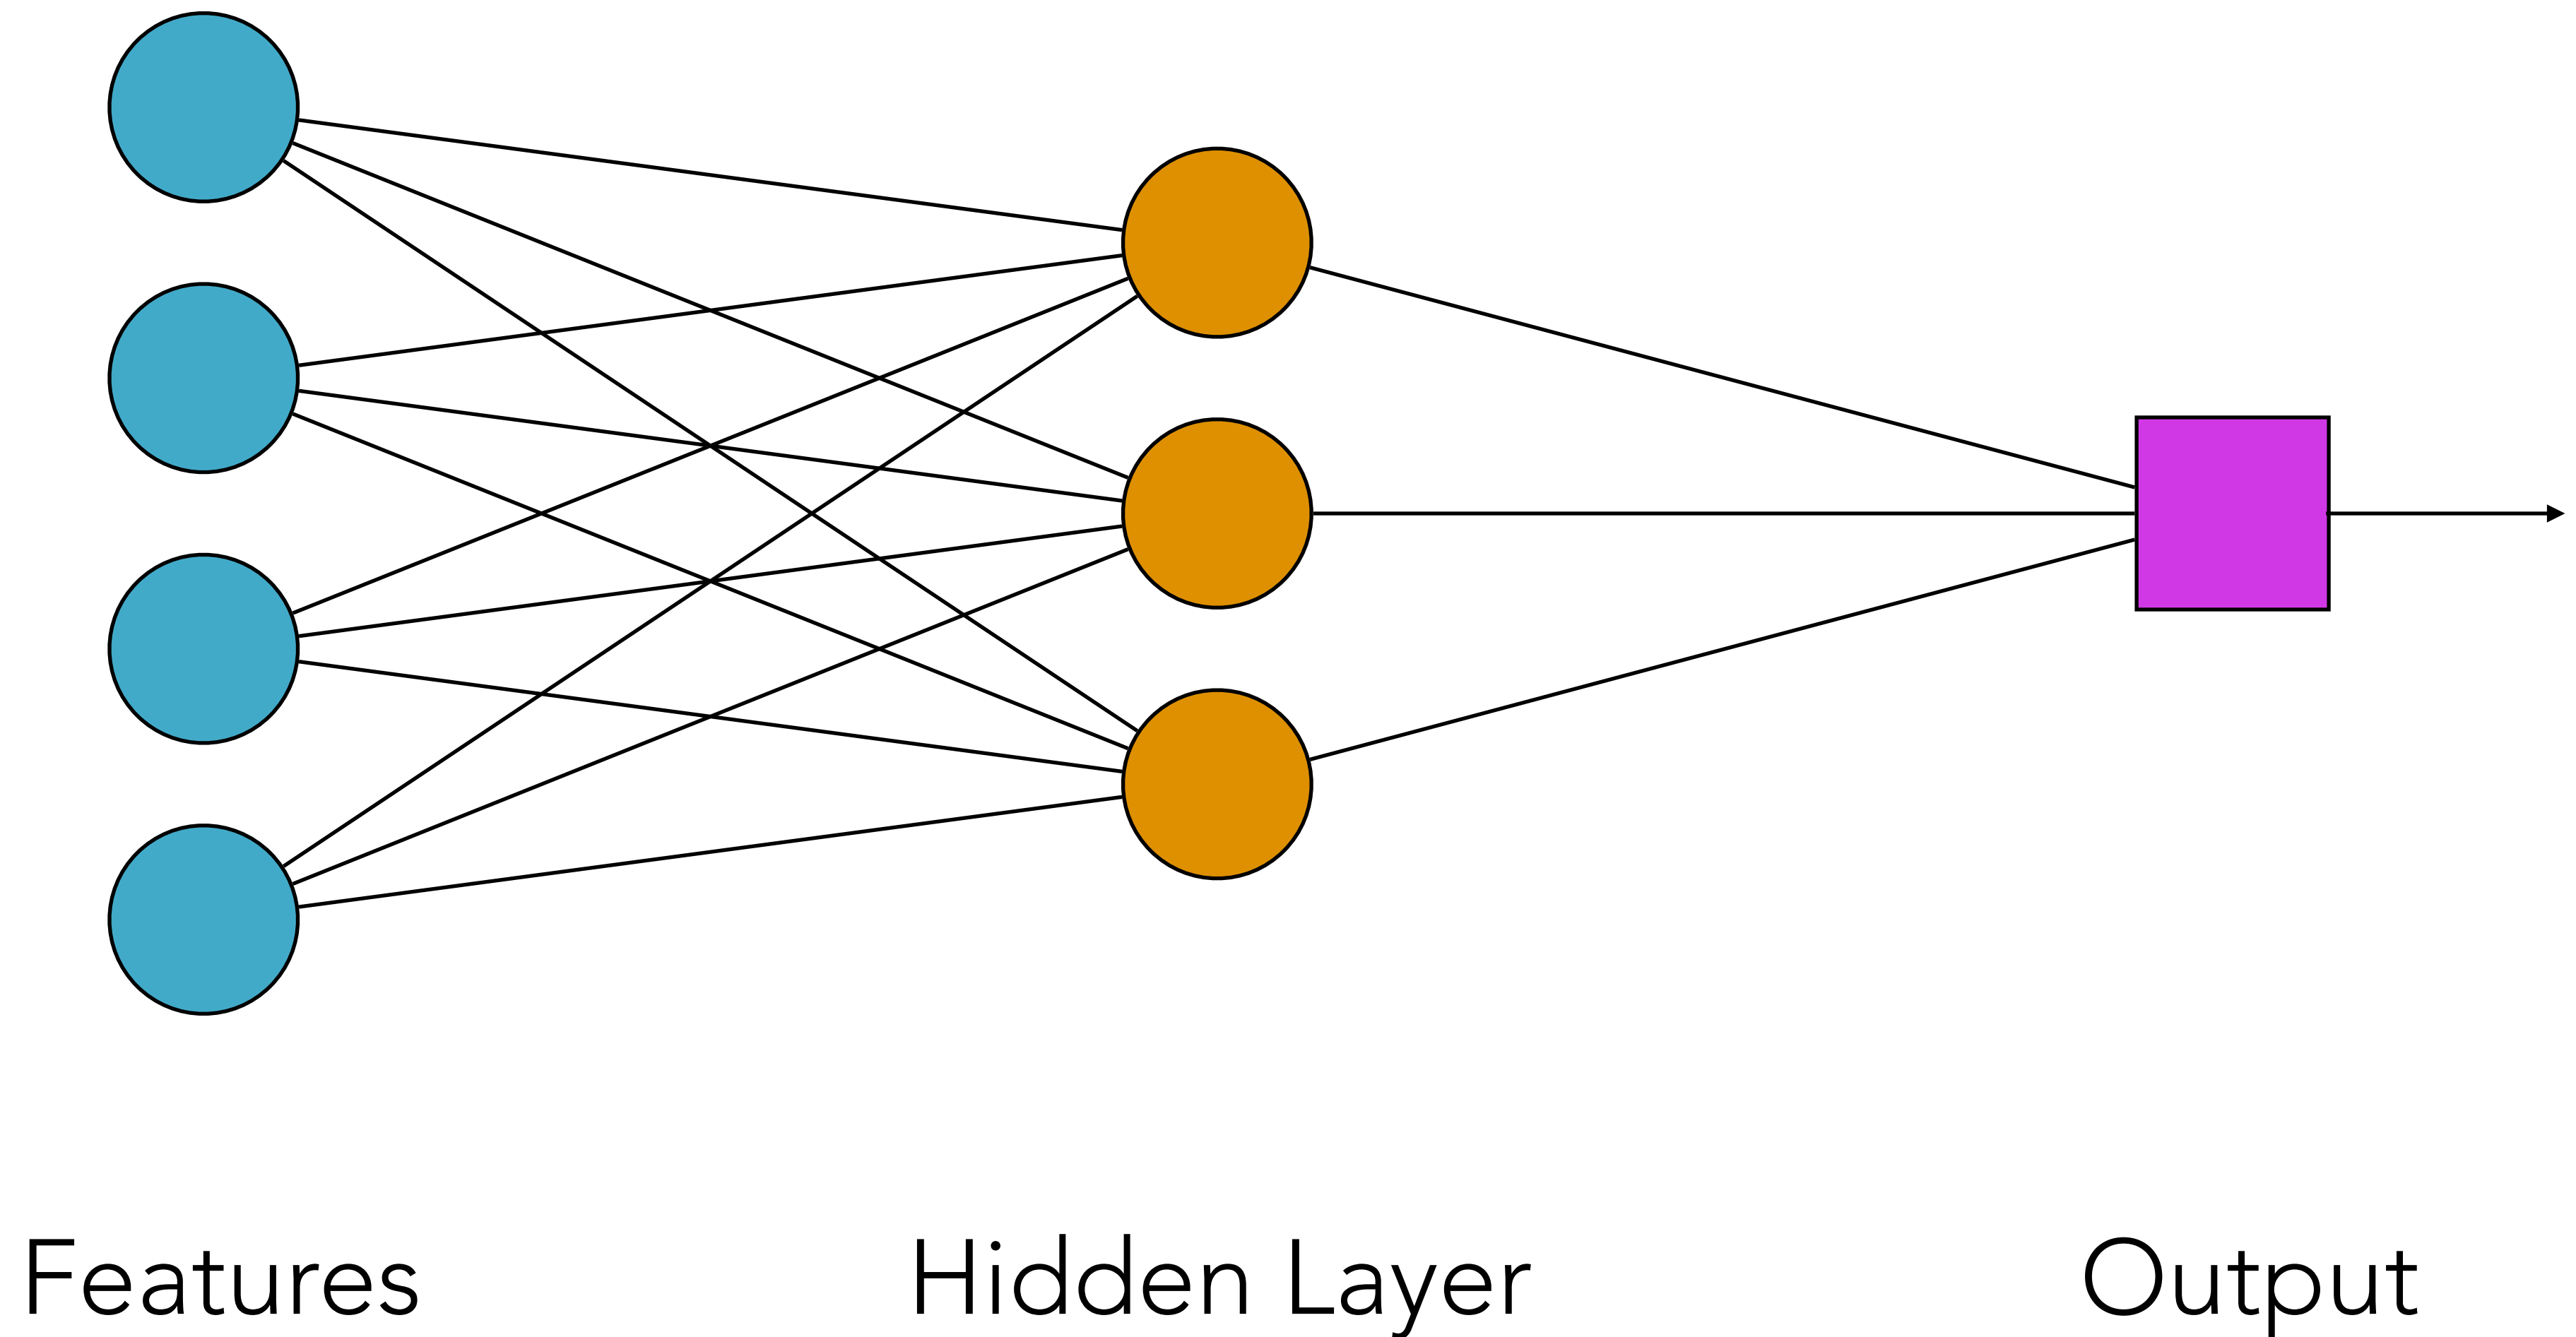
Hidden Layer

Output

CHECK: HOW MANY TRAINABLE PARAMETERS?

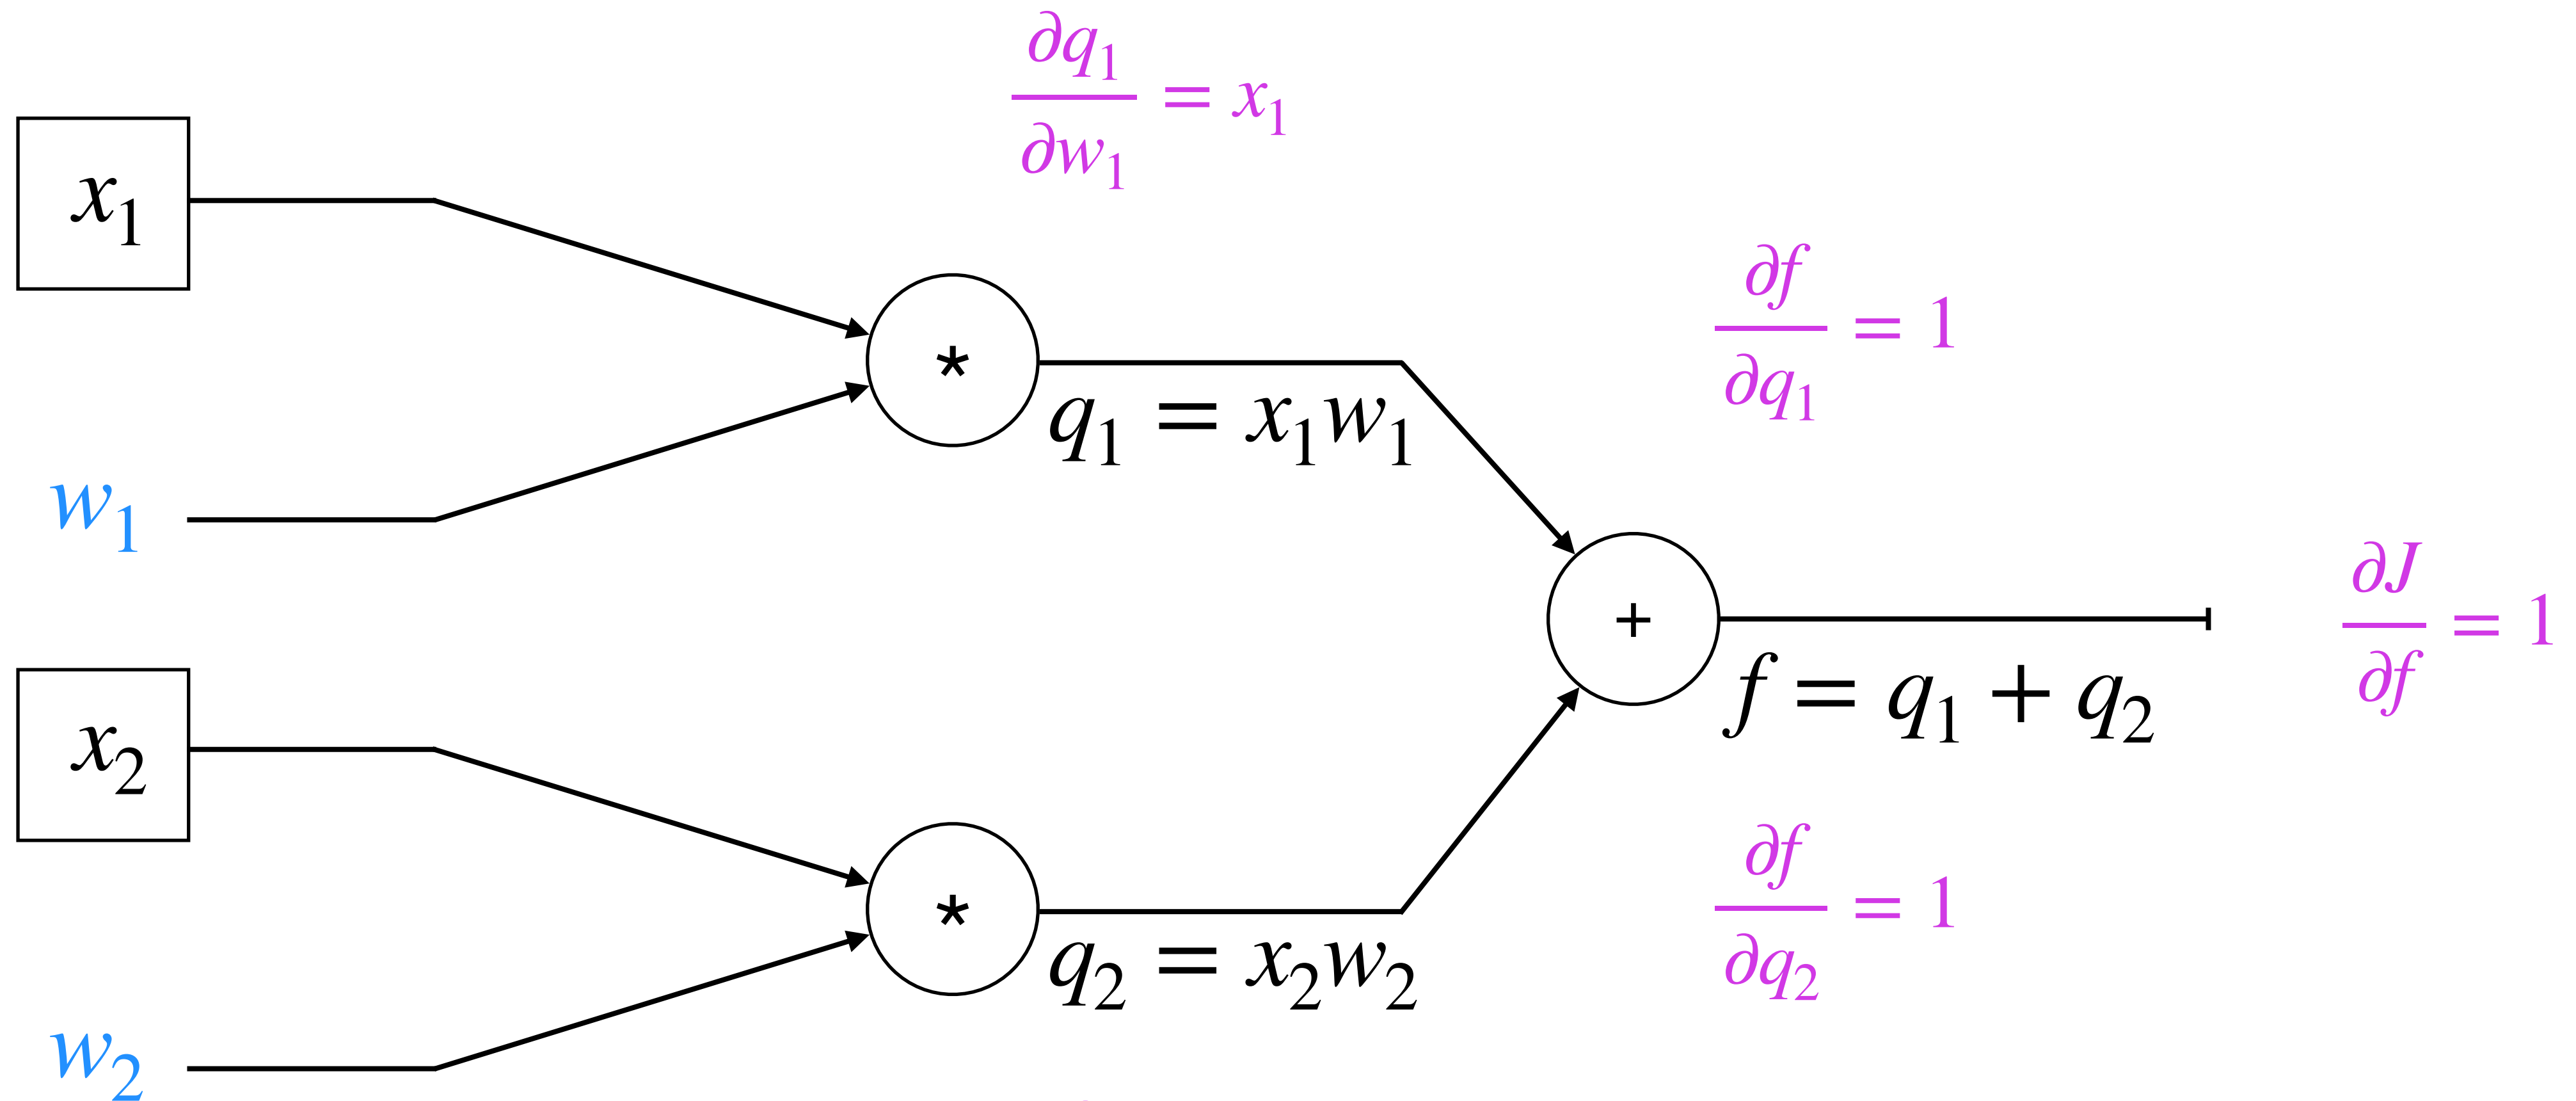


CHECK: HOW MANY TRAINABLE PARAMETERS?



BACKPROPAGATION

$$w_1 = w_1 + \eta * \frac{\partial J}{\partial f} \frac{\partial f}{\partial q_1} \frac{\partial q_1}{\partial w_1}$$

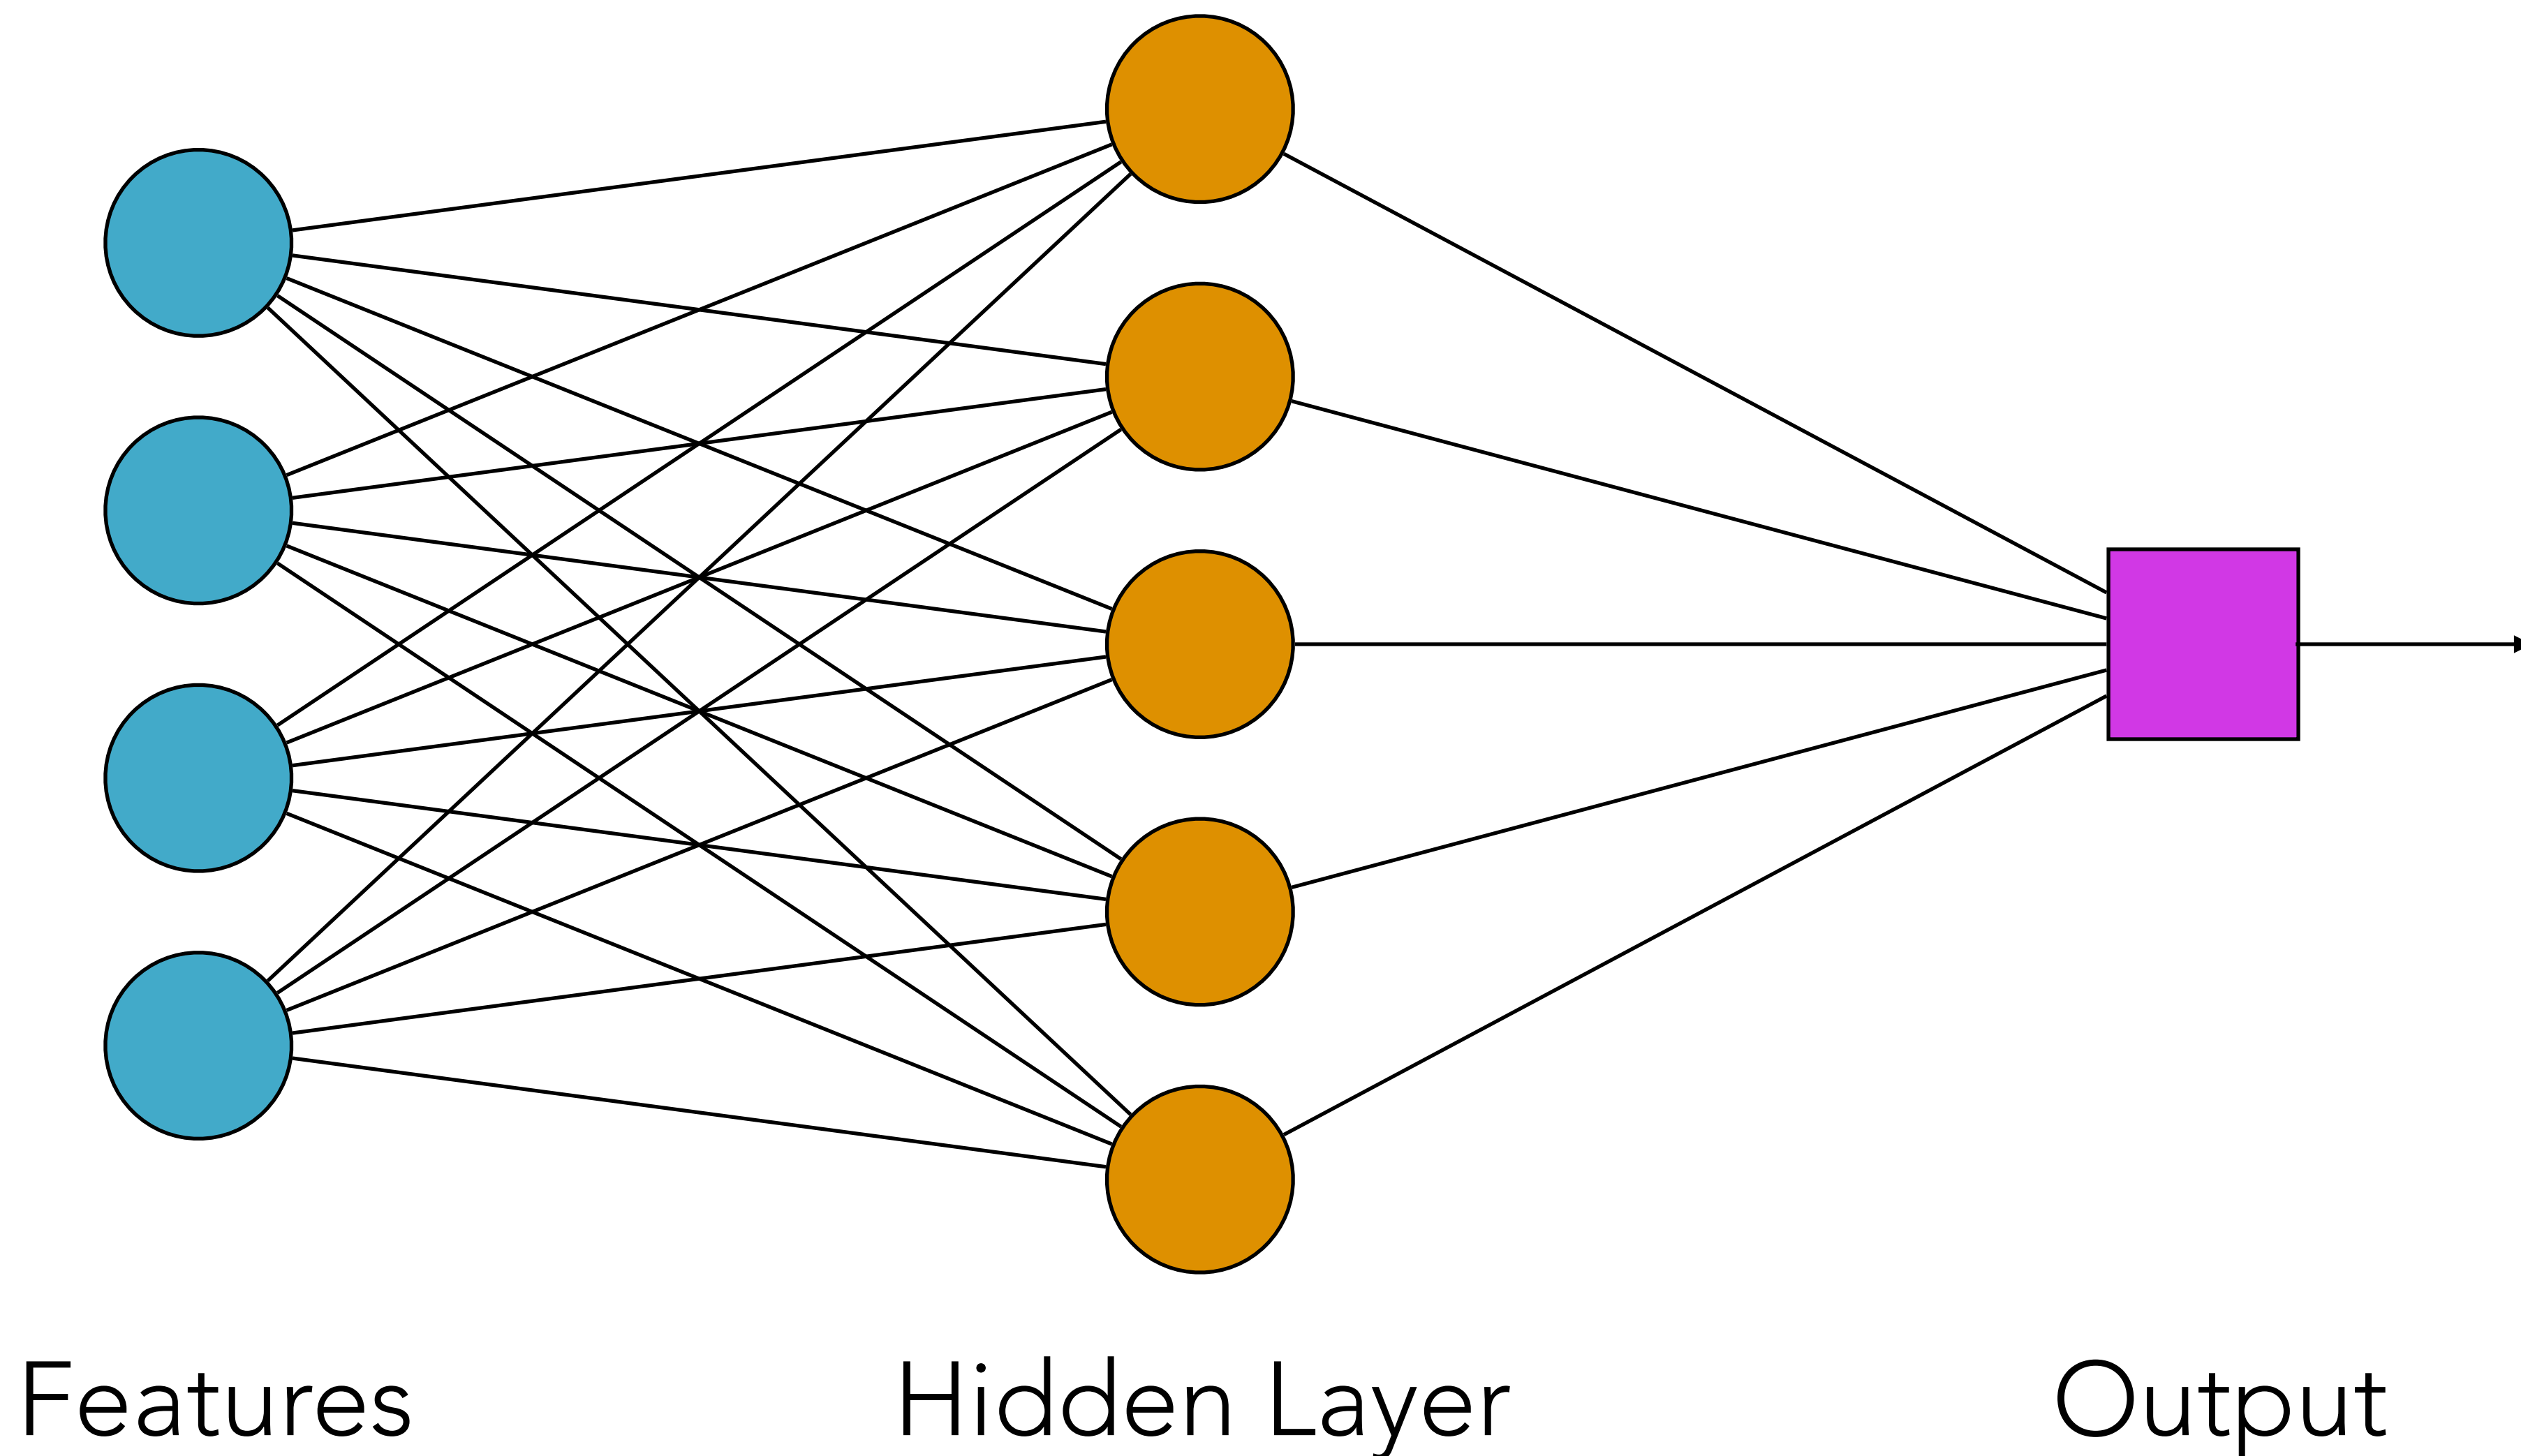


$$w_2 = w_2 + \eta * \frac{\partial J}{\partial f} \frac{\partial f}{\partial q_2} \frac{\partial q_2}{\partial w_2}$$

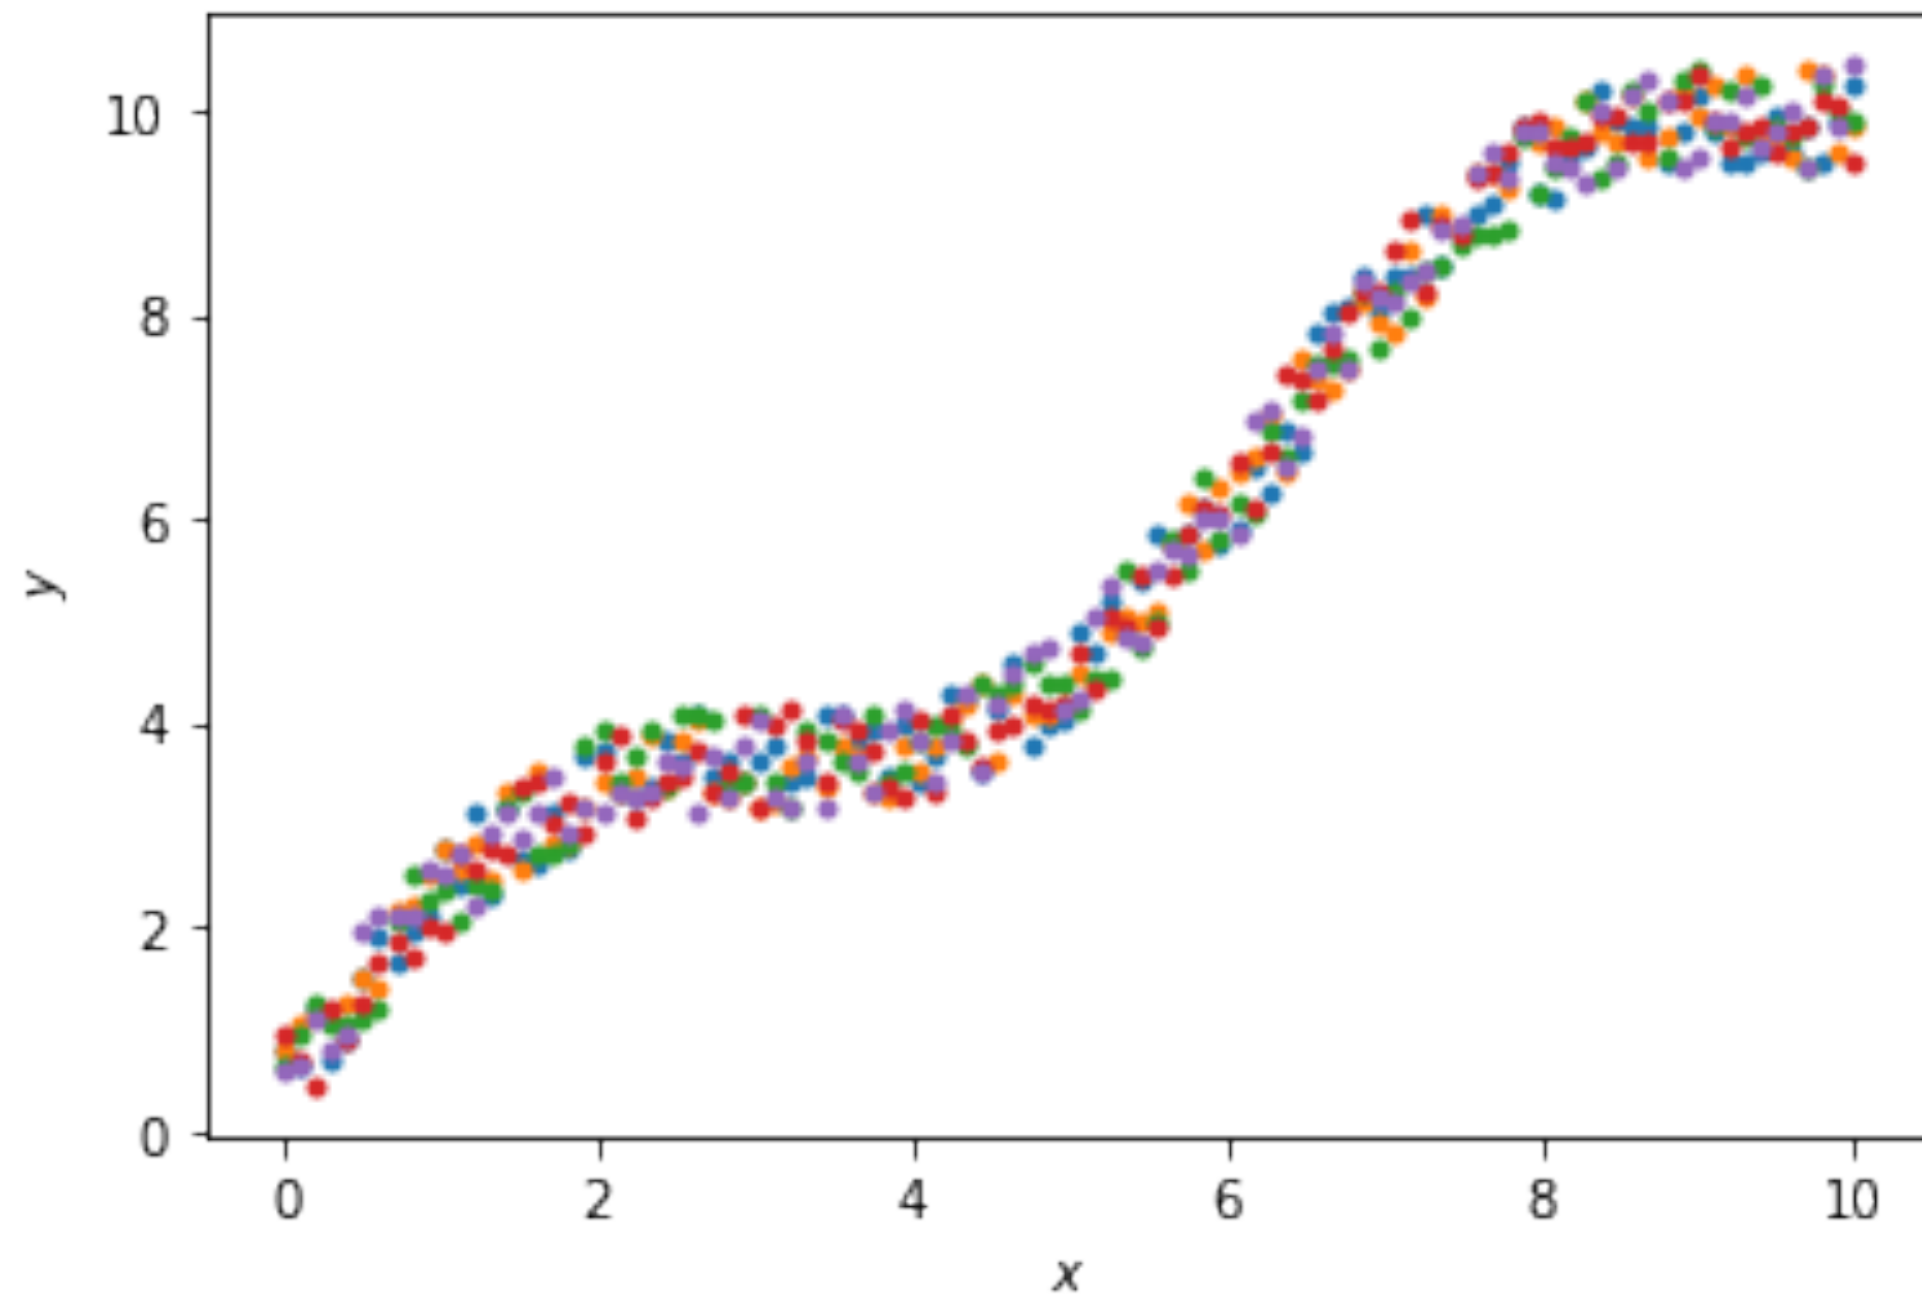
Loss function

$$J(w) = f - \hat{f}$$

Weight initialization: What happens if we initialize all weights to same value?



LOSS FUNCTIONS



Loss function

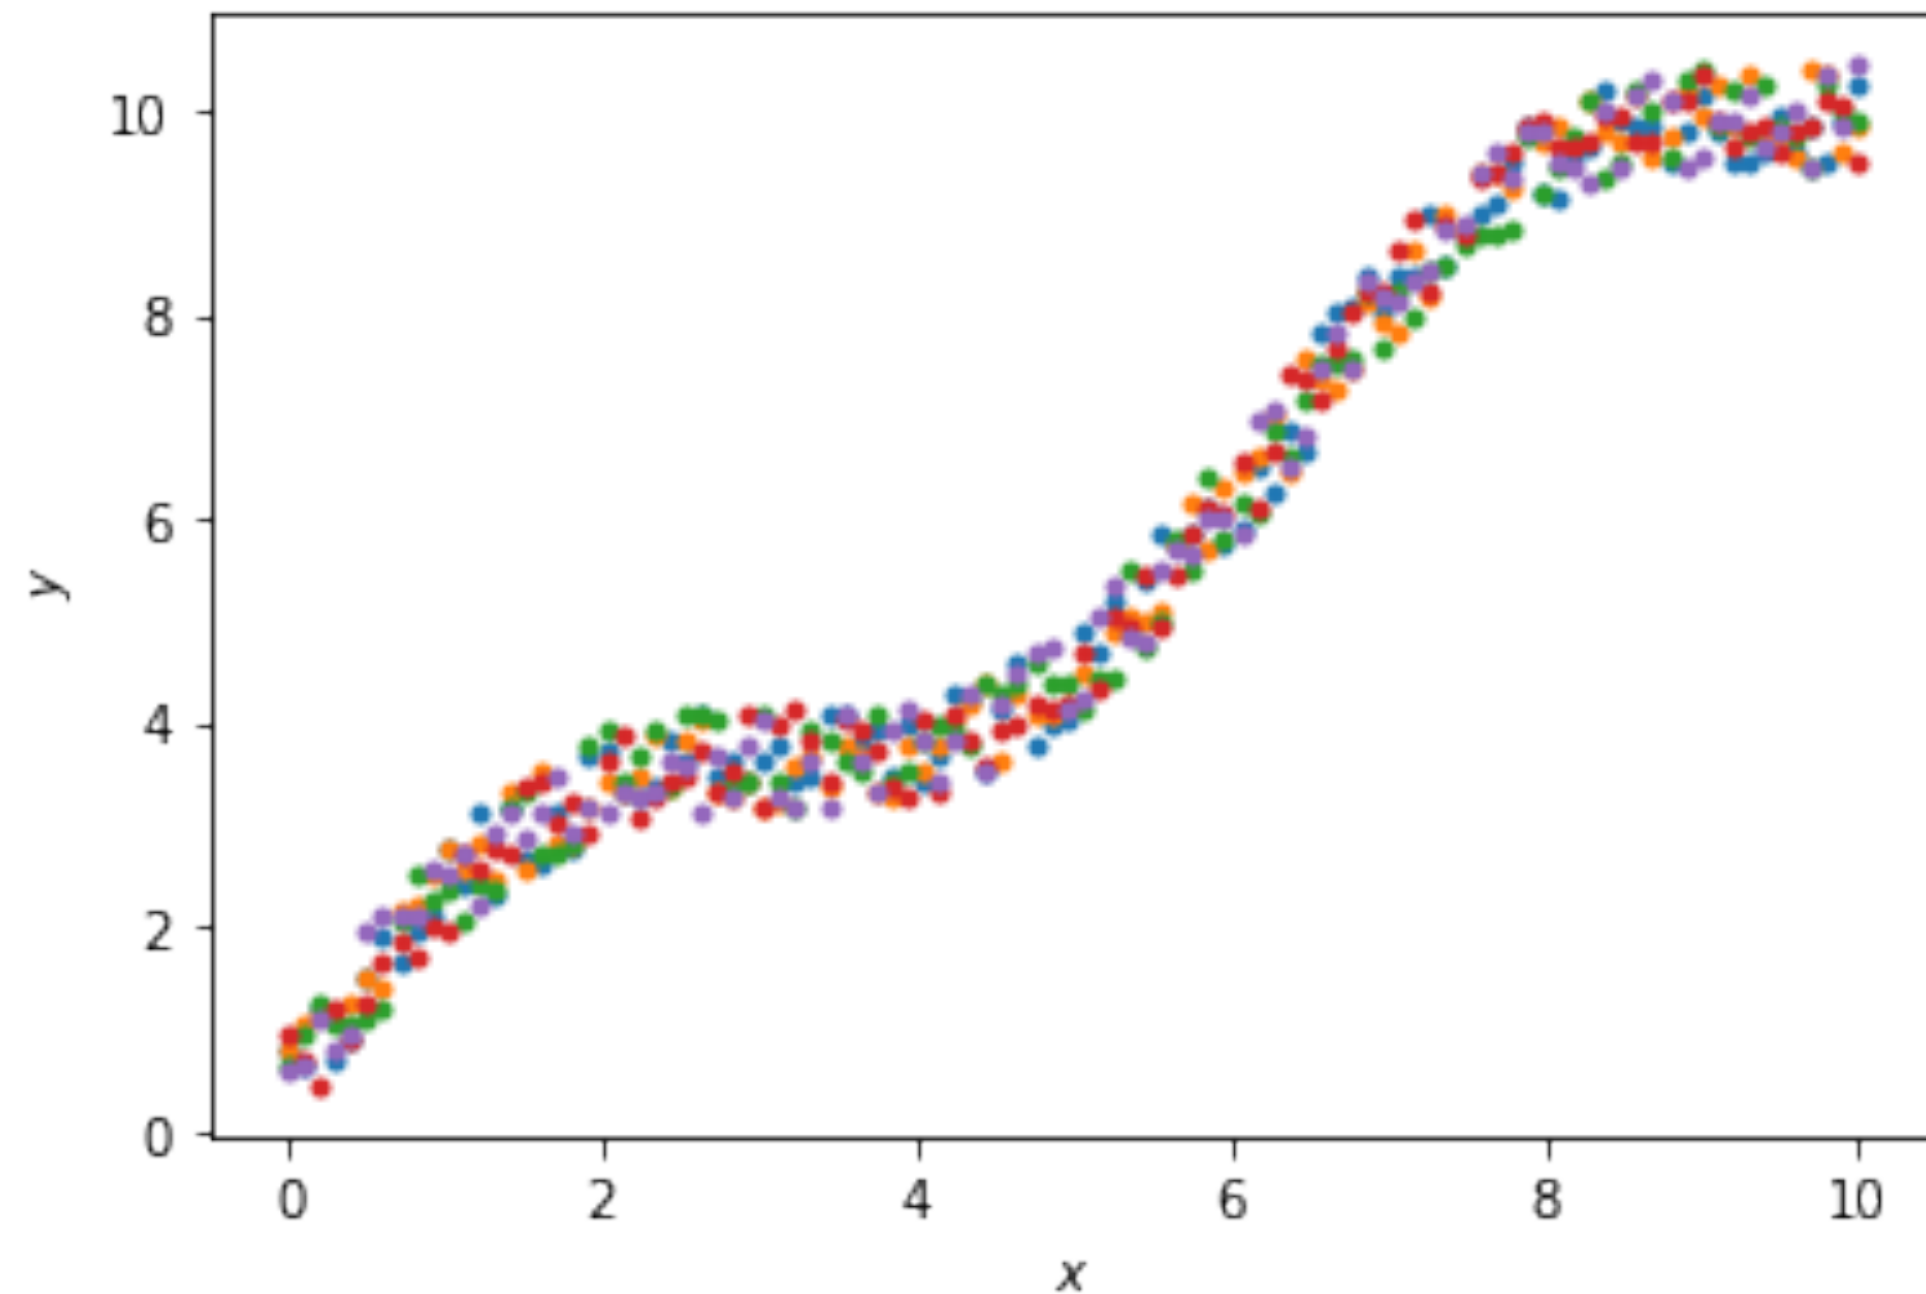
Mean squared error
(MSE)

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

Mean absolute error
(MAE)

$$J(w) = \frac{1}{N} \sum_{i=0}^N |\hat{y}_i - y_i|$$

LOSS FUNCTIONS



Loss function

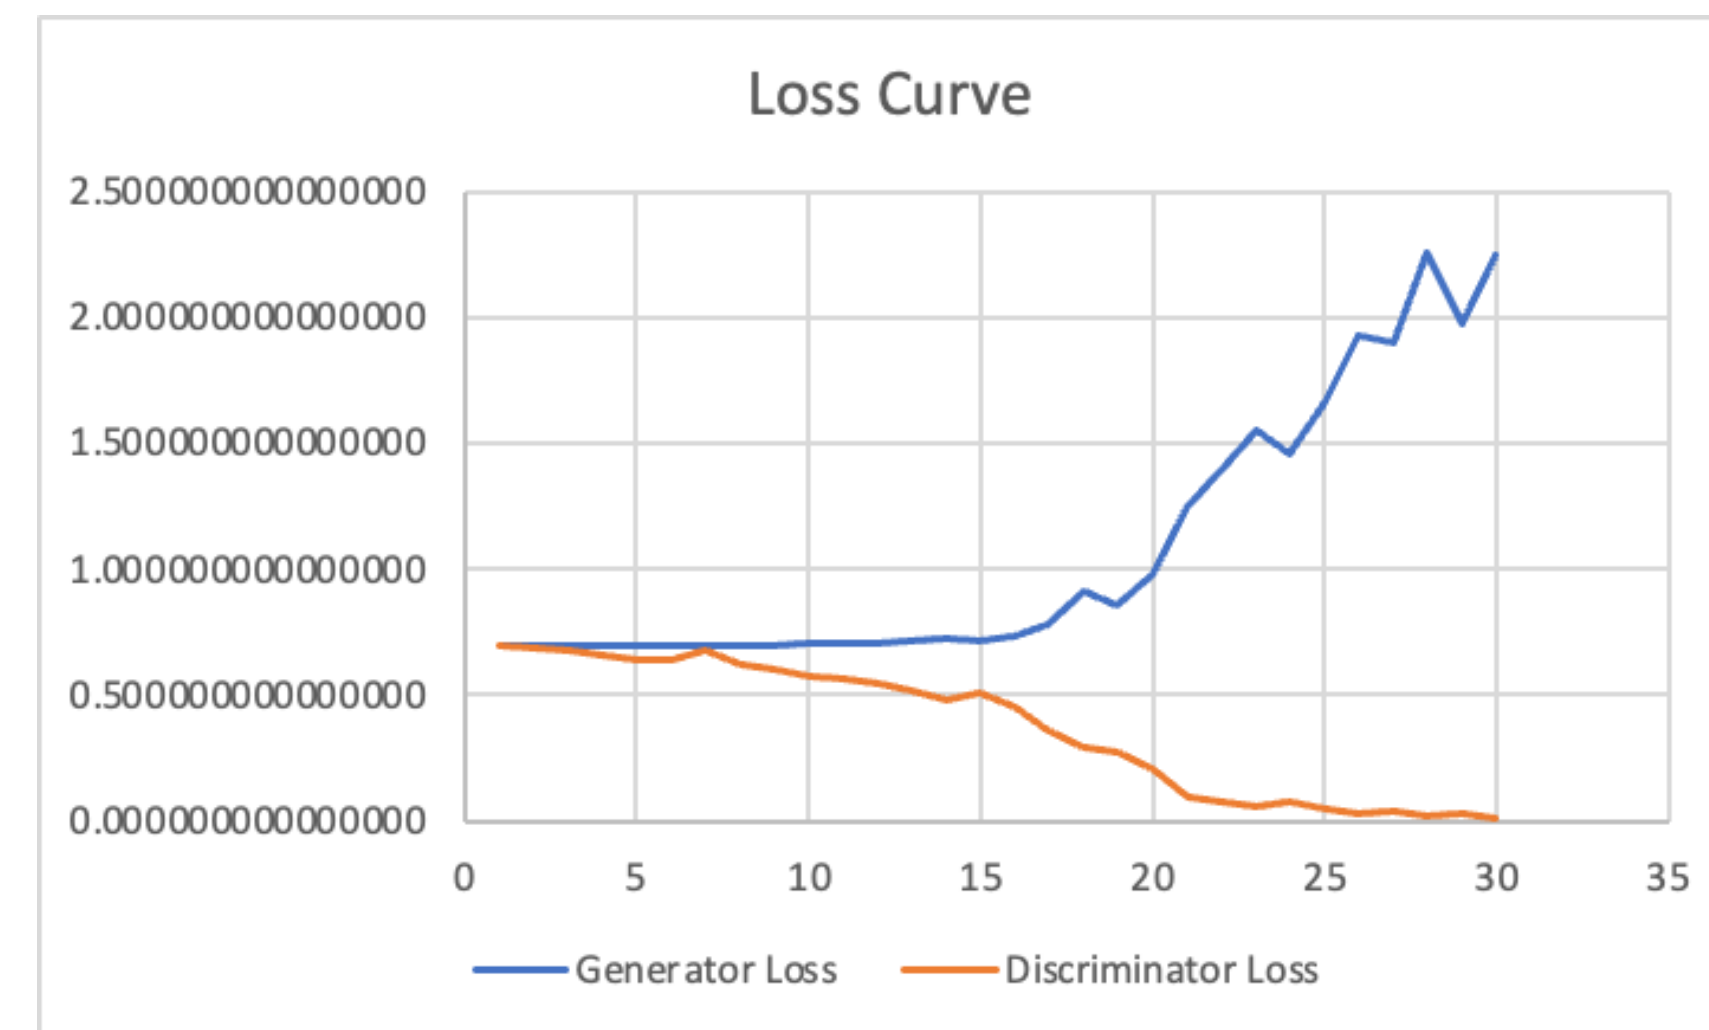
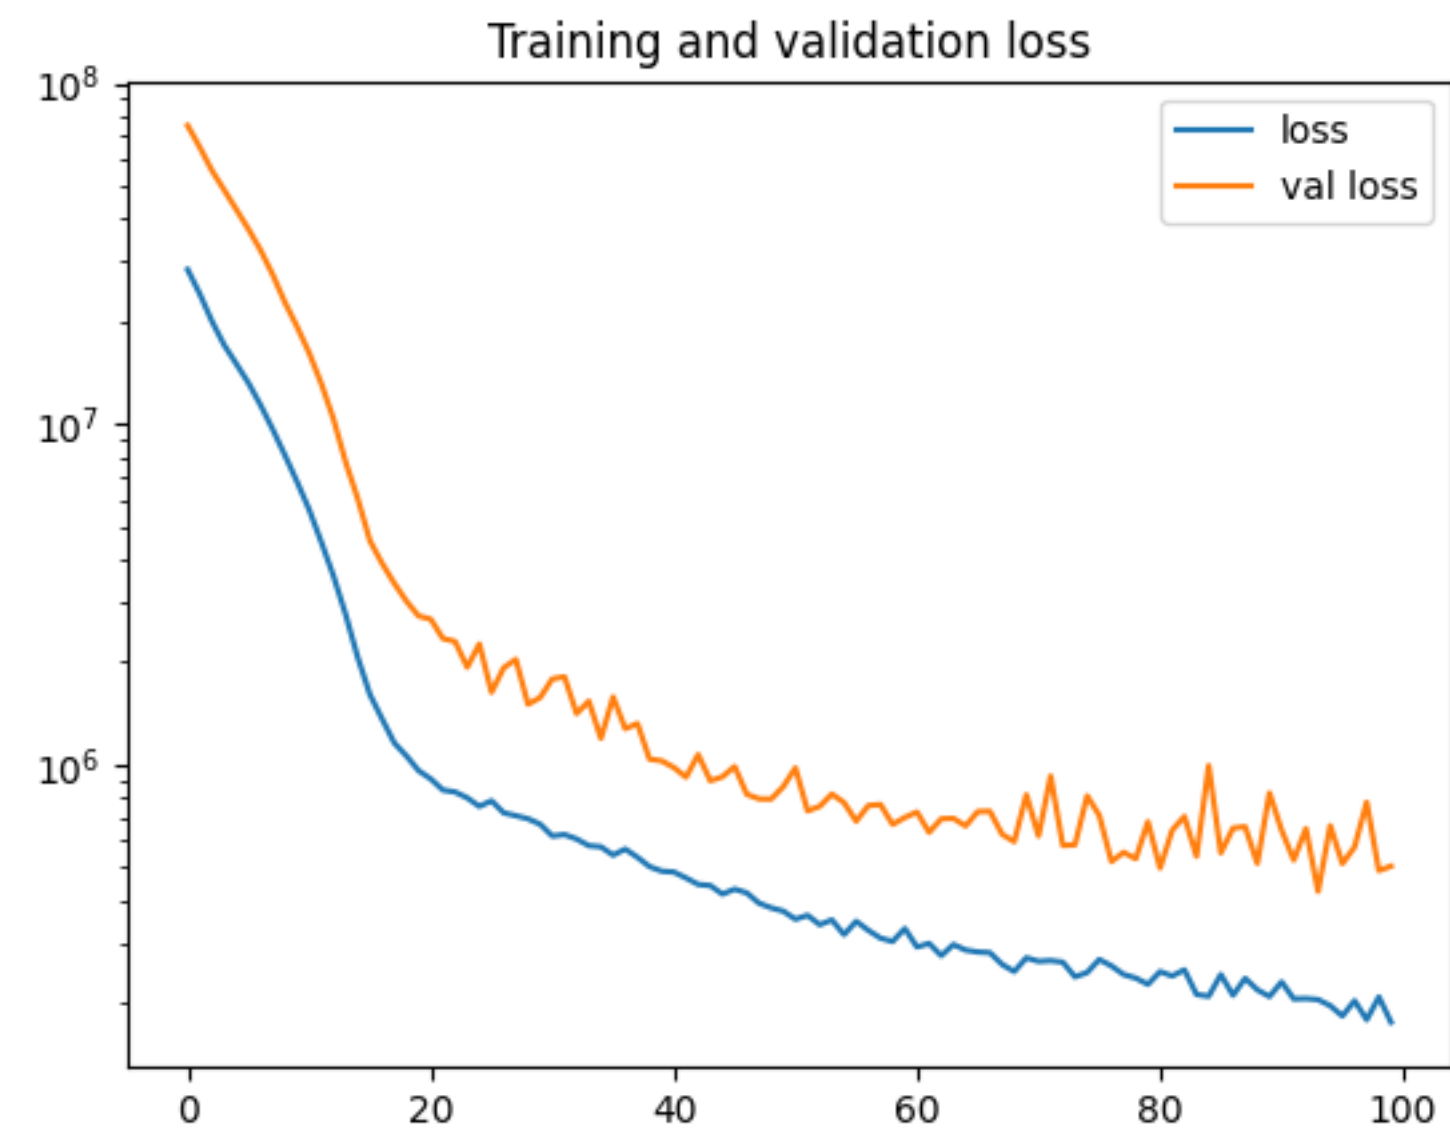
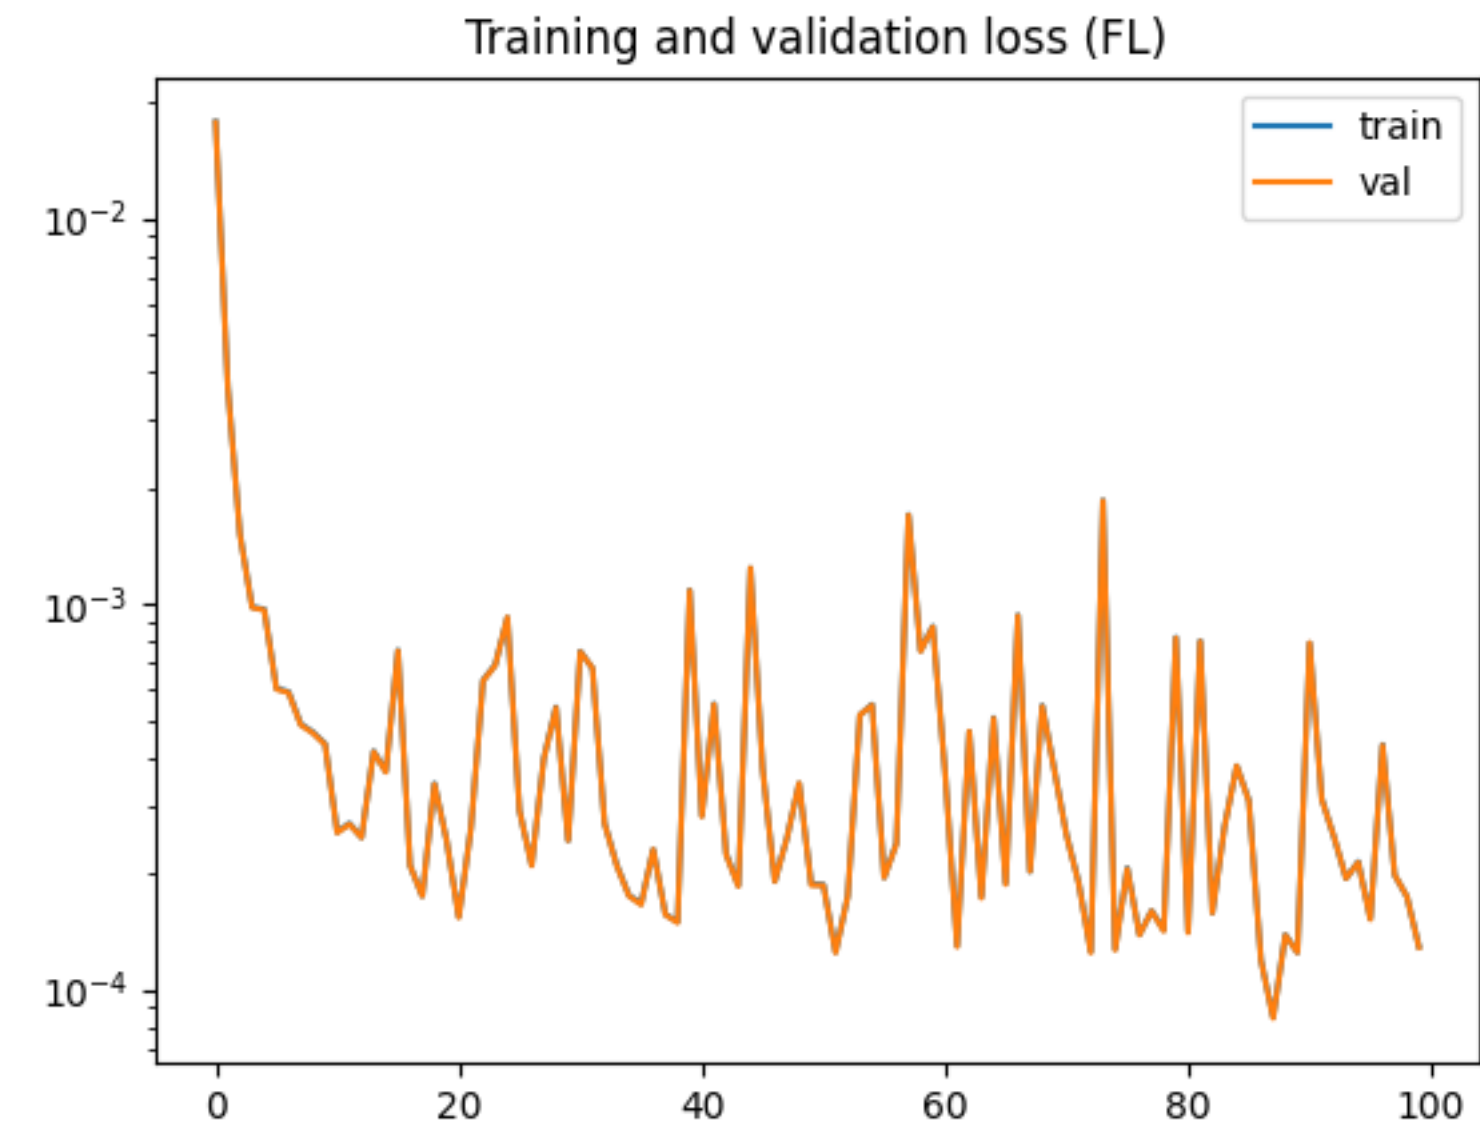
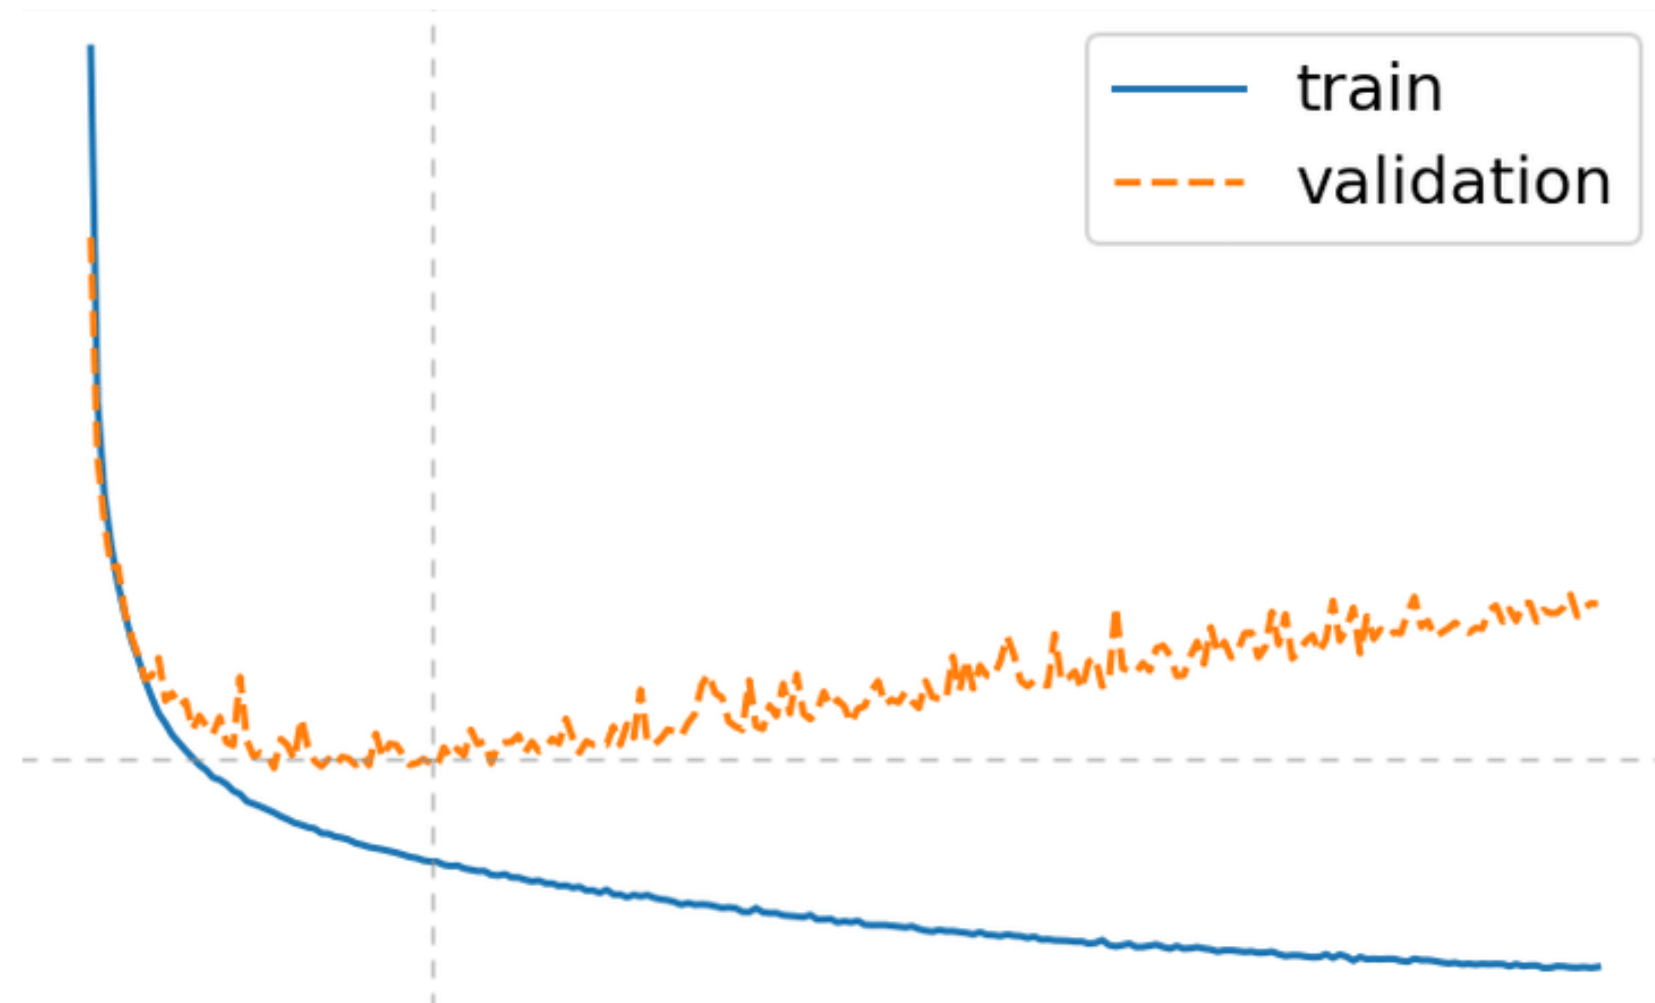
Mean squared error
mean

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

Mean absolute error
median

$$J(w) = \frac{1}{N} \sum_{i=0}^N |\hat{y}_i - y_i|$$

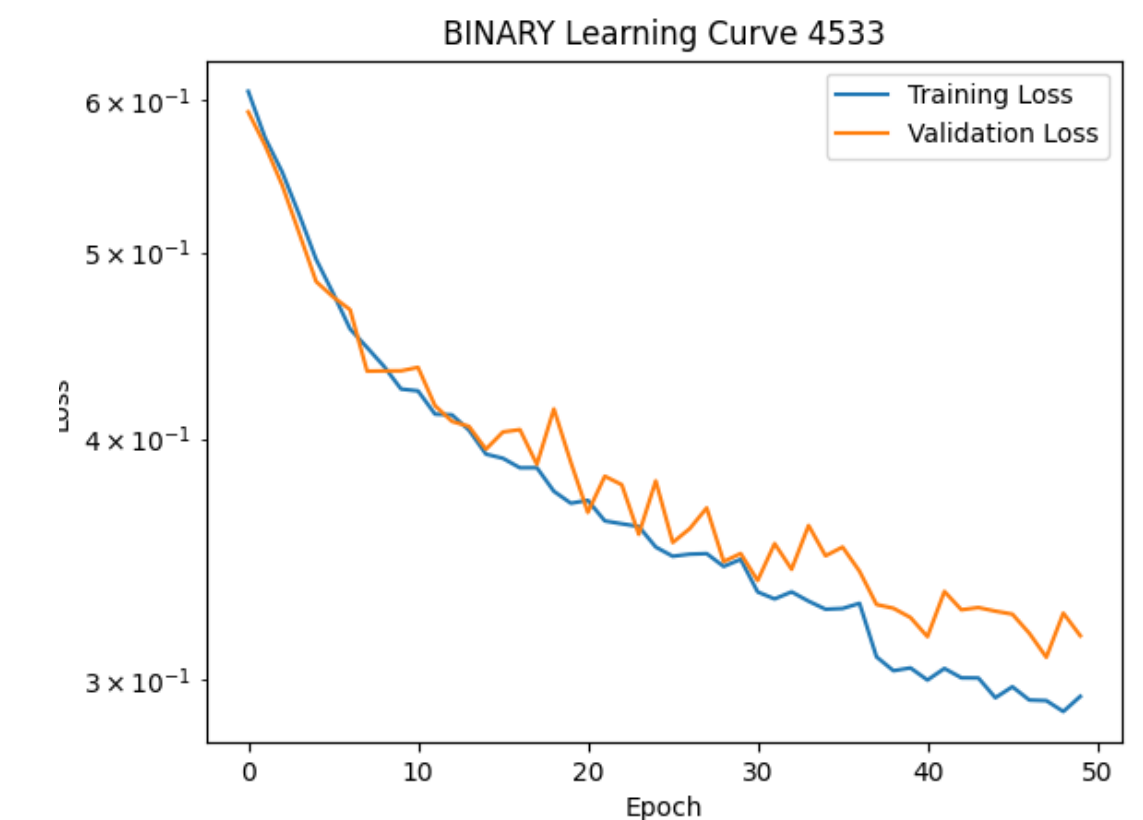
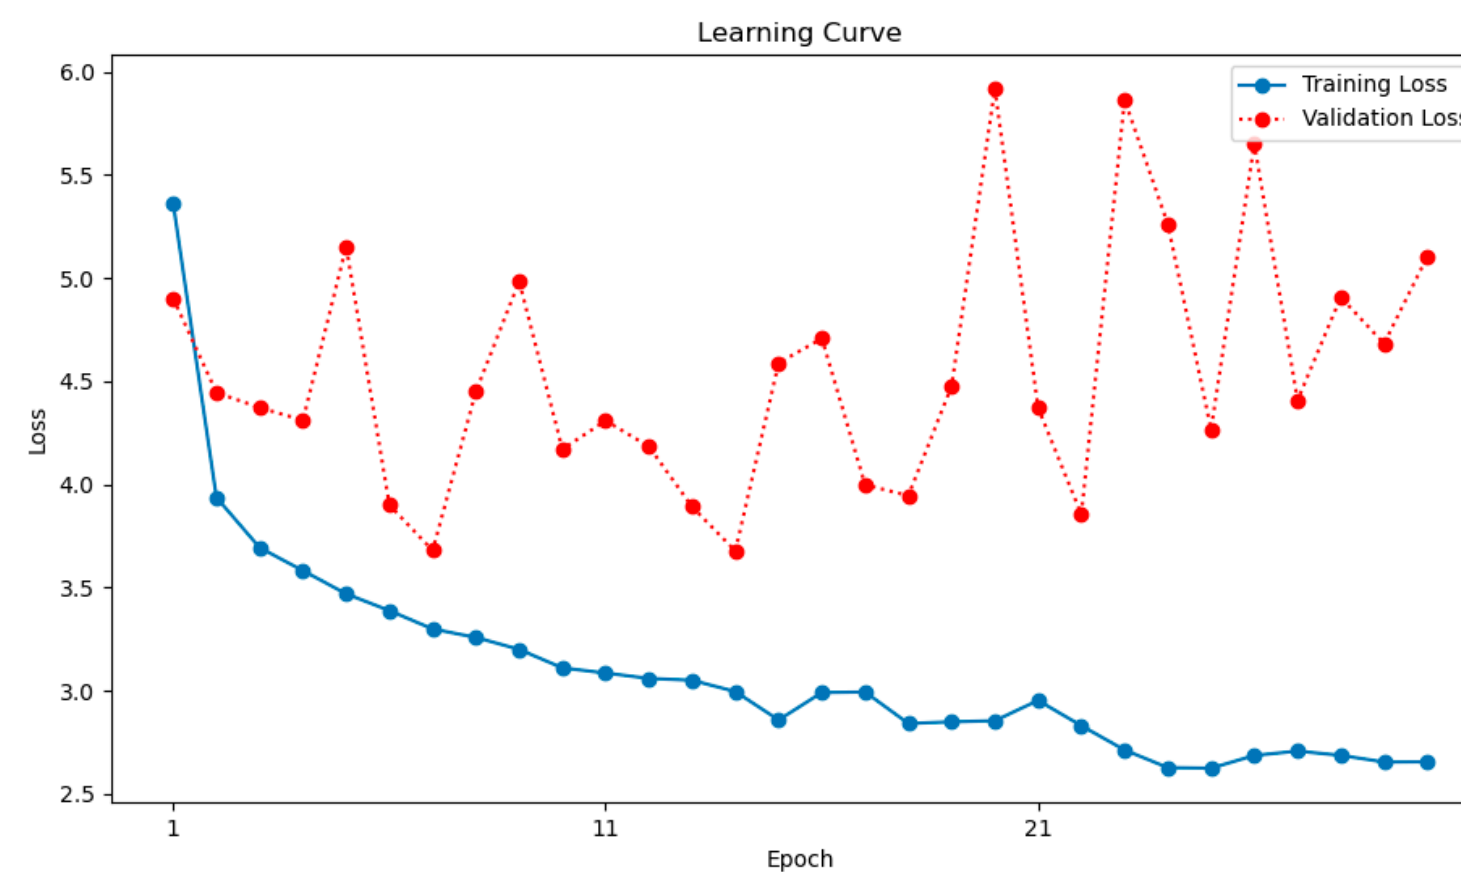
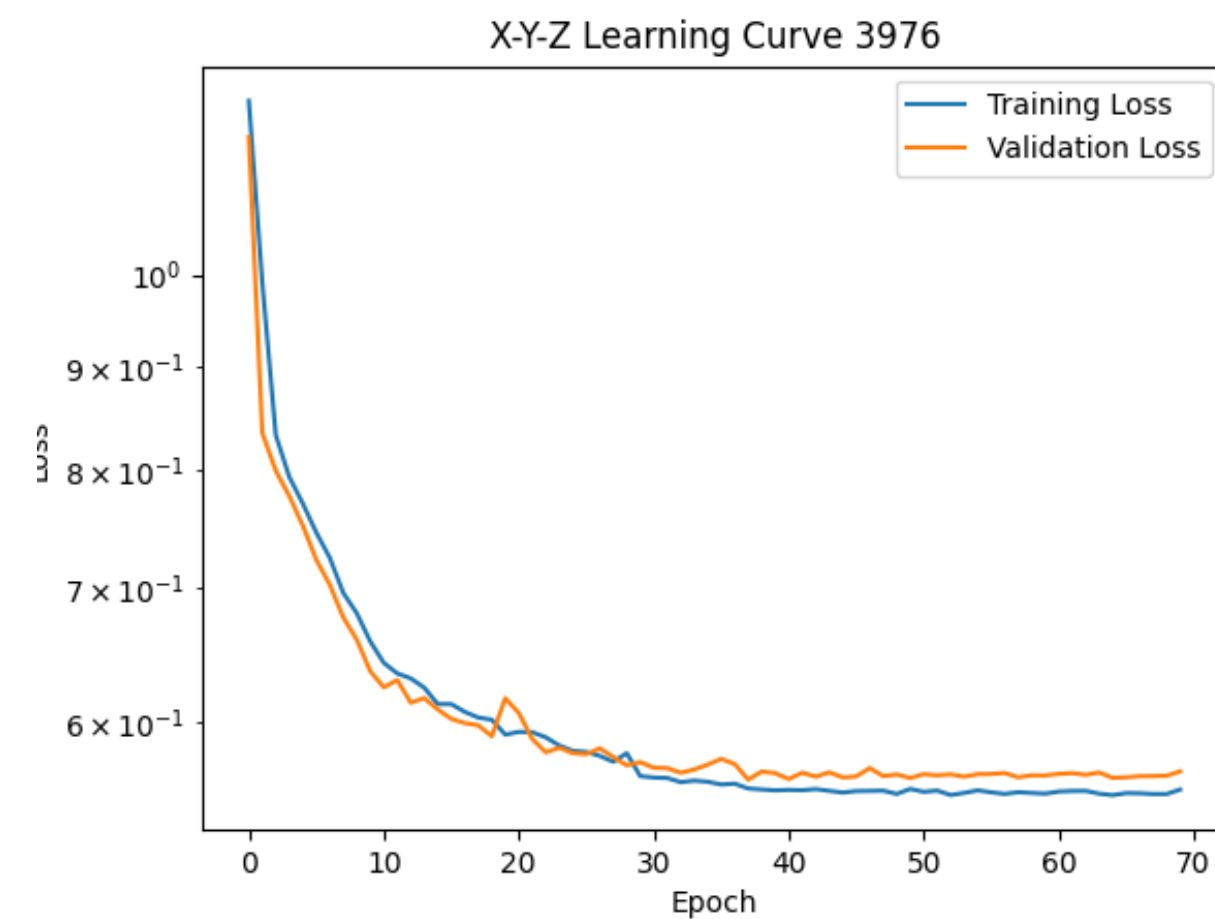
Learning (loss) curves



TRAINING

Remember that our goal is NOT to minimize loss on training data!

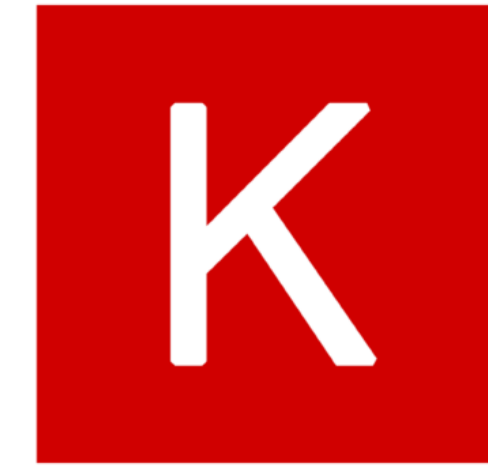
Learning curves



AUTOMATIC DIFFERENTIATION



TensorFlow



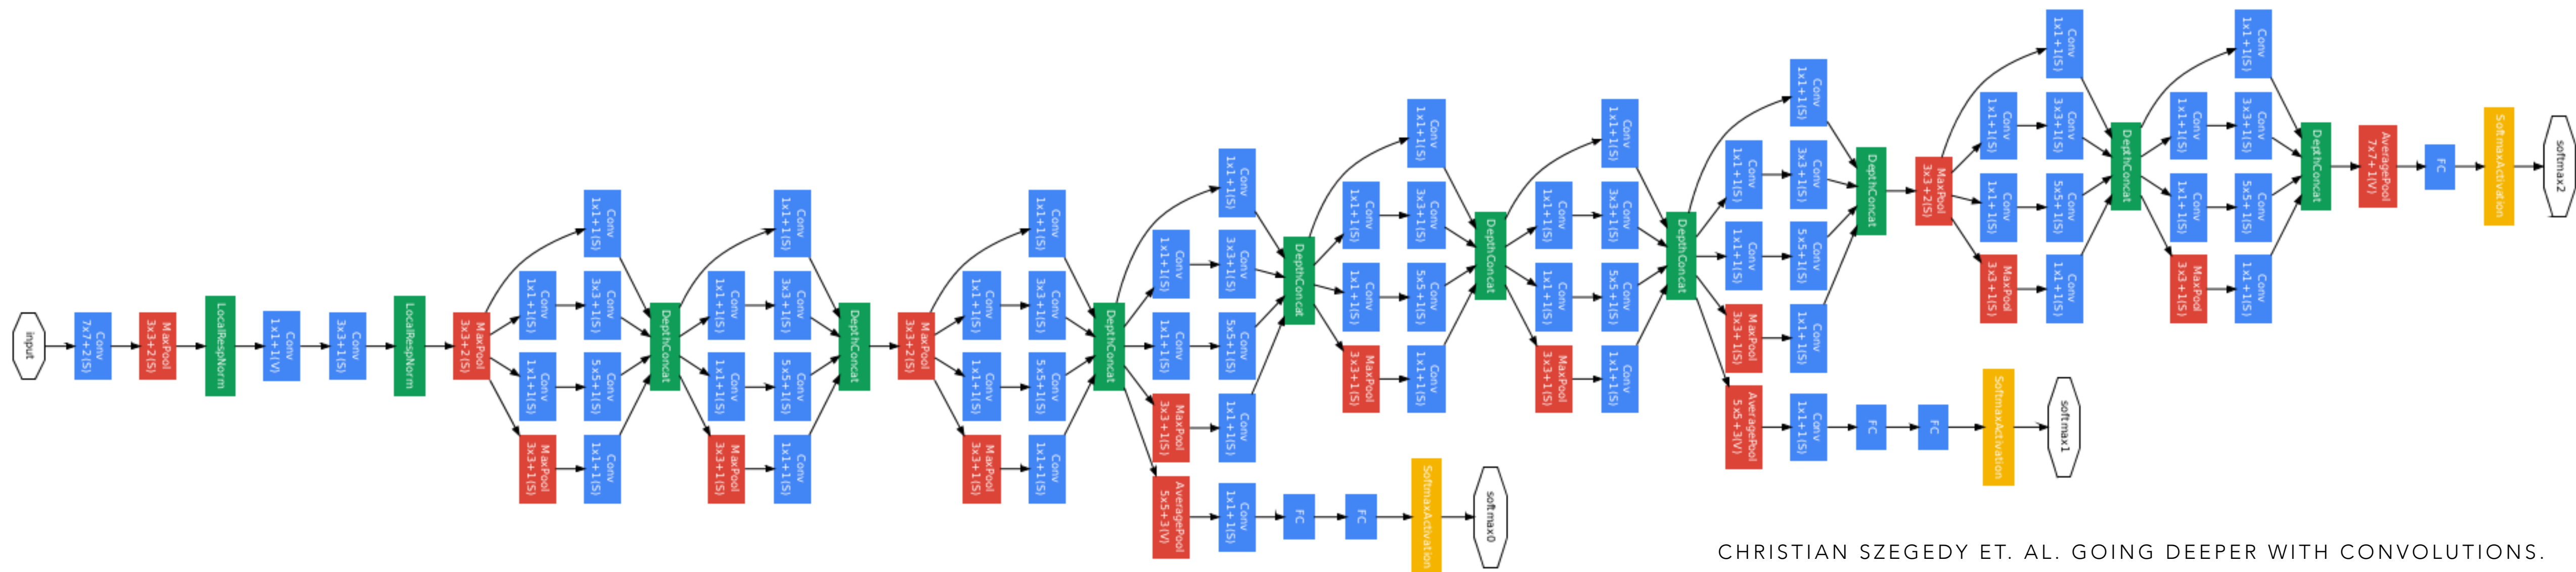
Keras



PyTorch

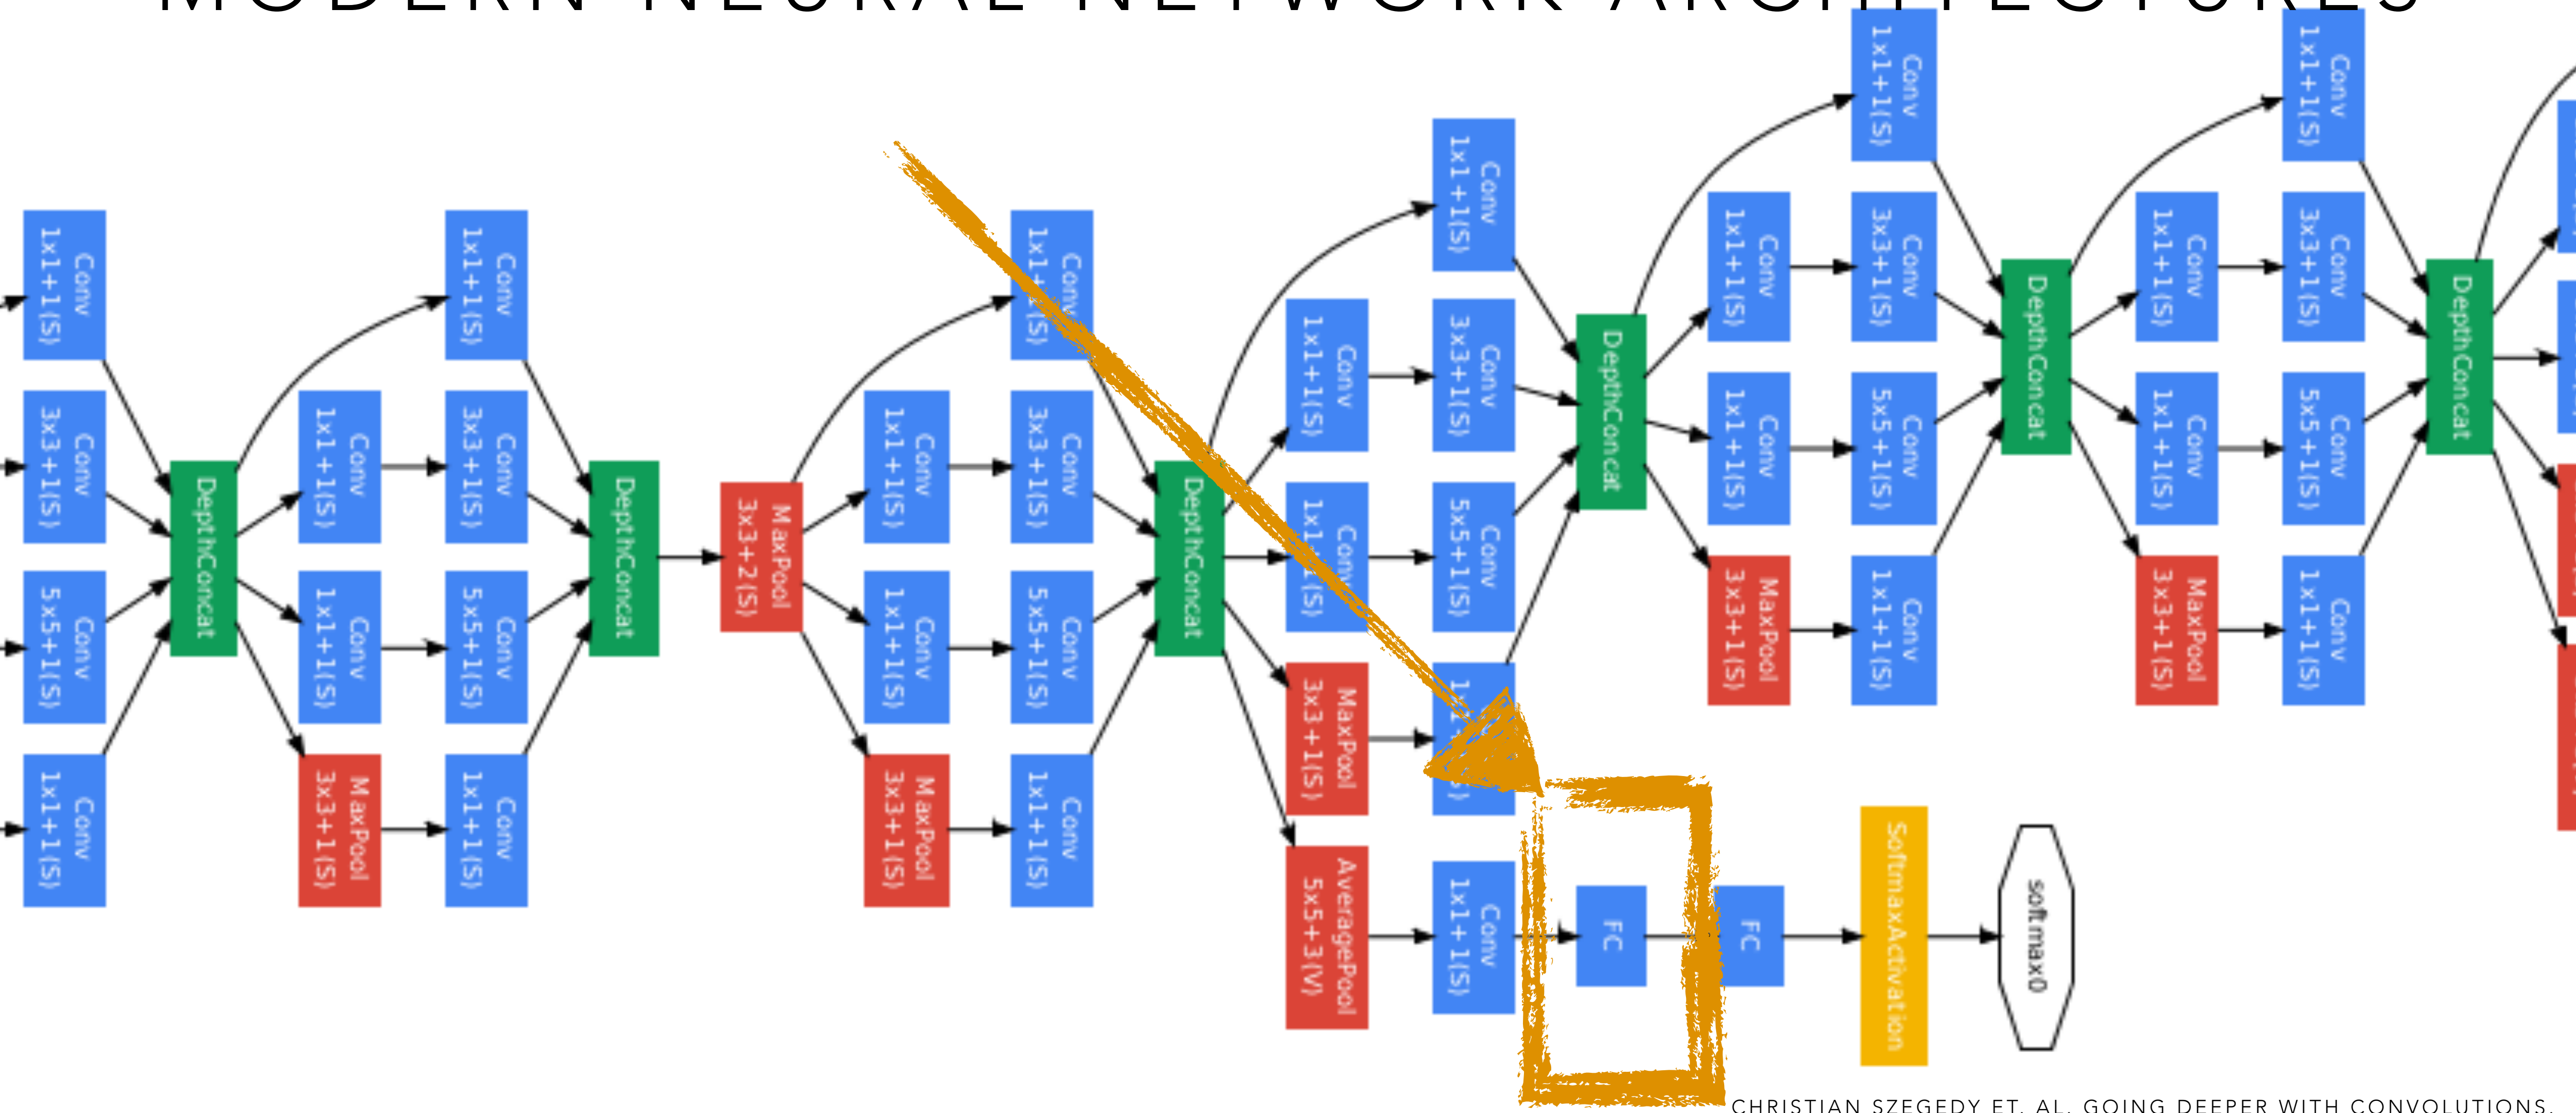


MODERN NEURAL NETWORK ARCHITECTURES



“GoogLeNet network with all the bells and whistles”

MODERN NEURAL NETWORK ARCHITECTURES



“GoogLeNet network with all the bells and whistles”

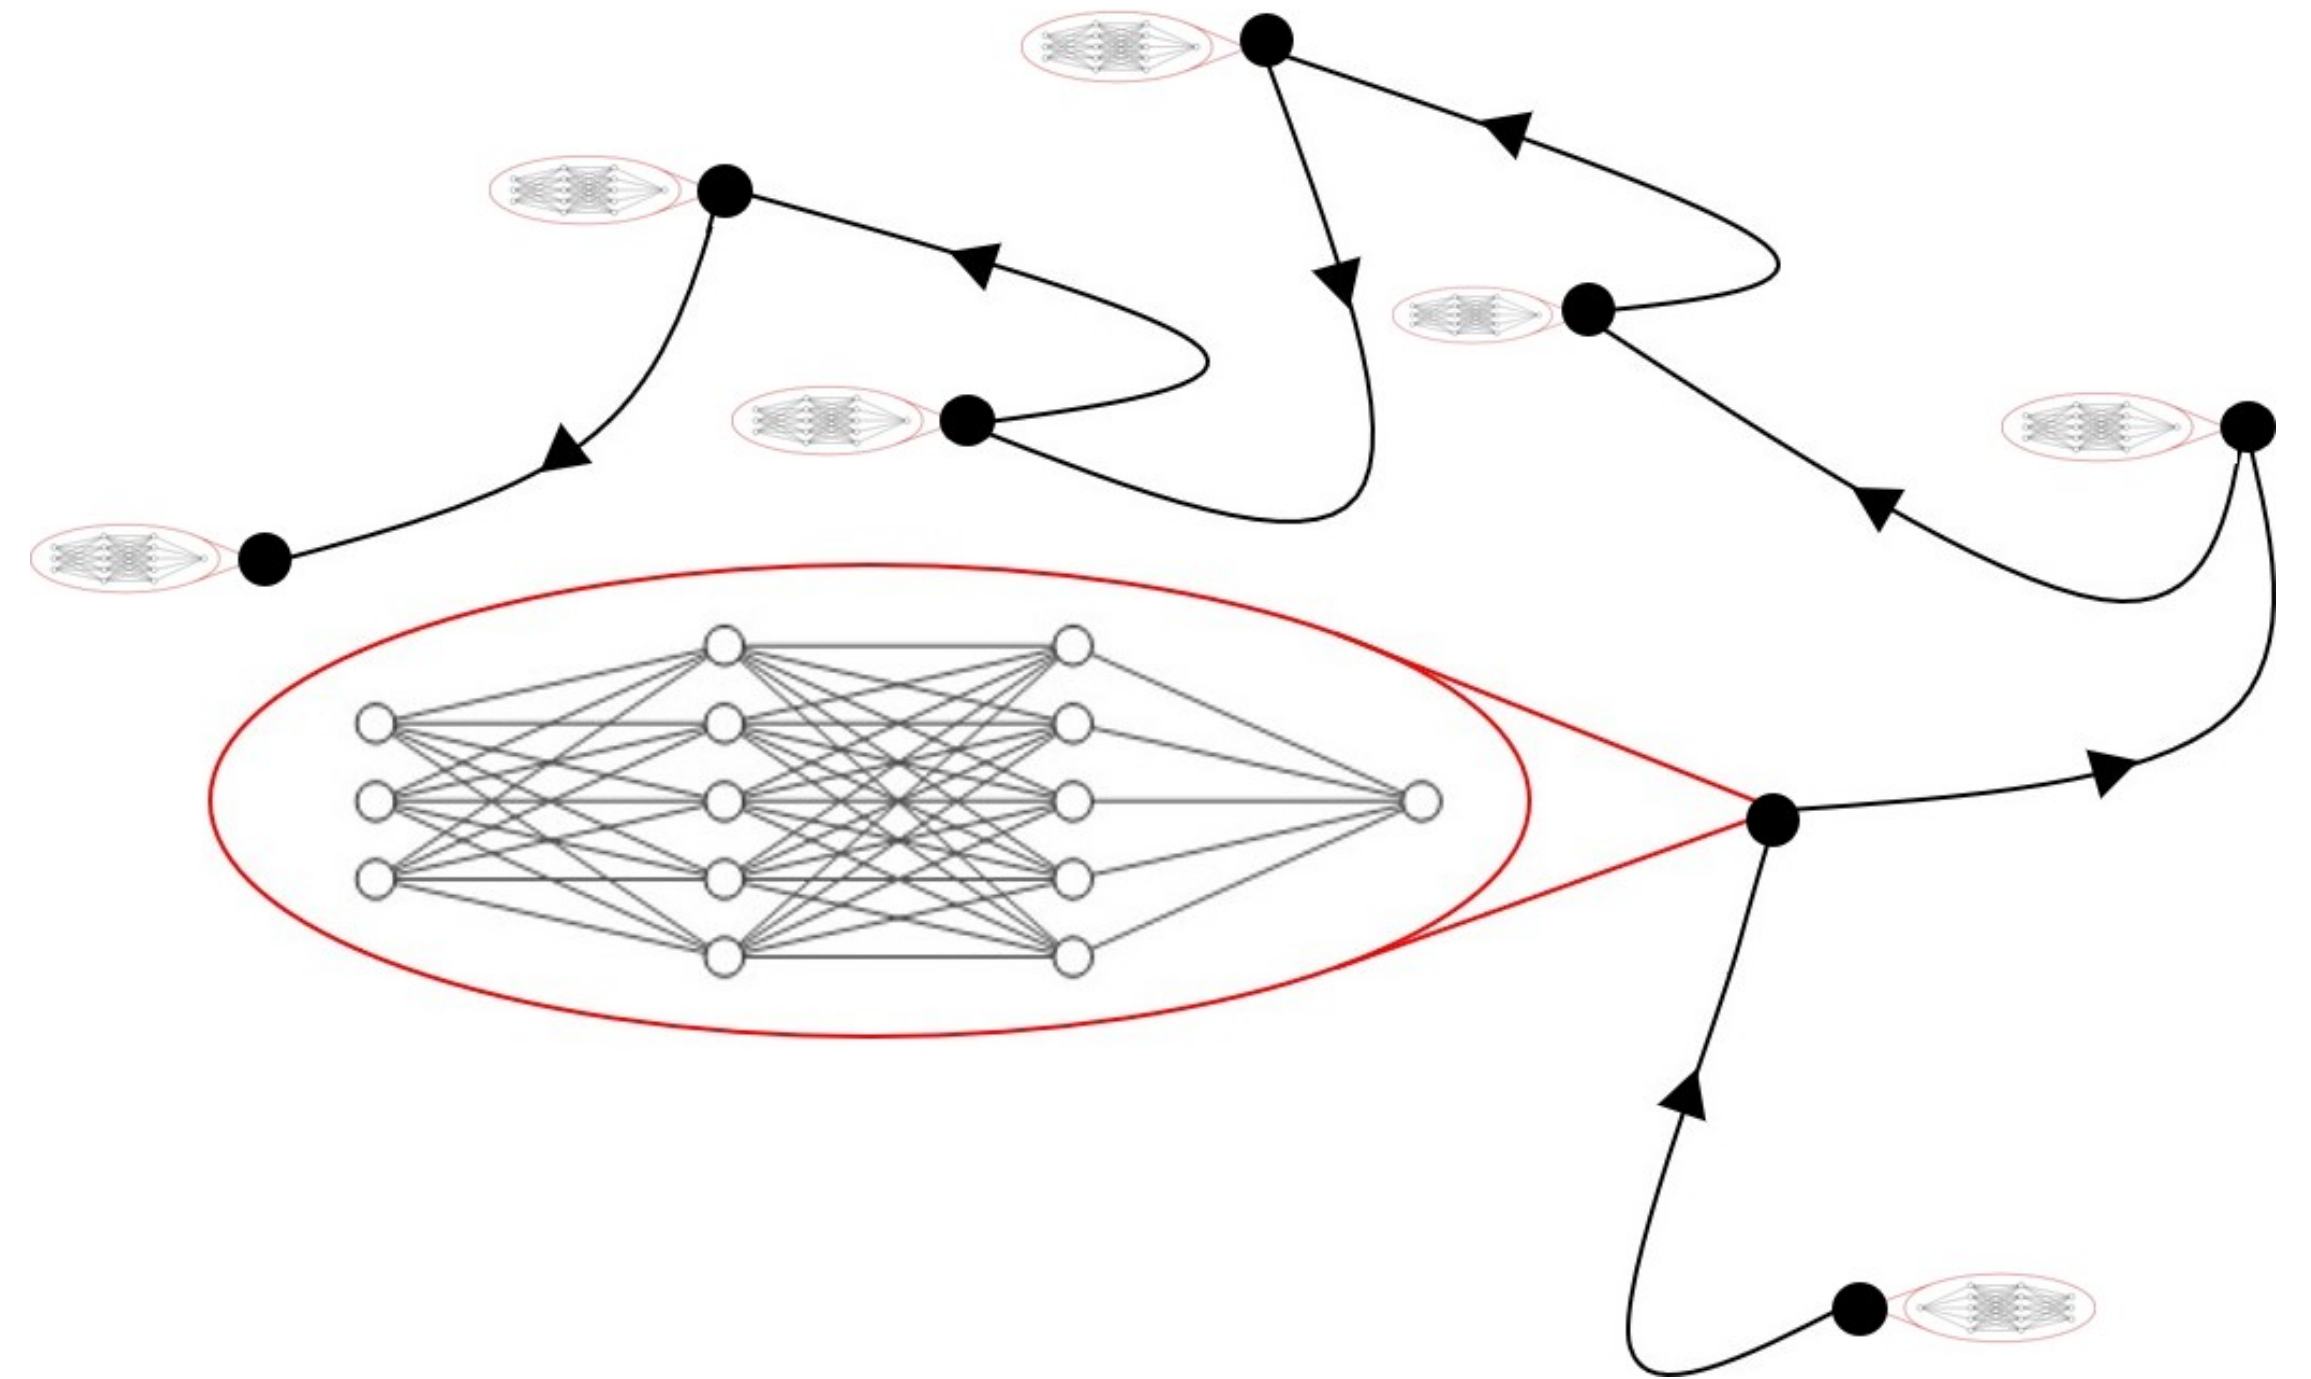
FAST MAPPING TO THEORETICAL PARAMETERS

Bayesian Neural Networks

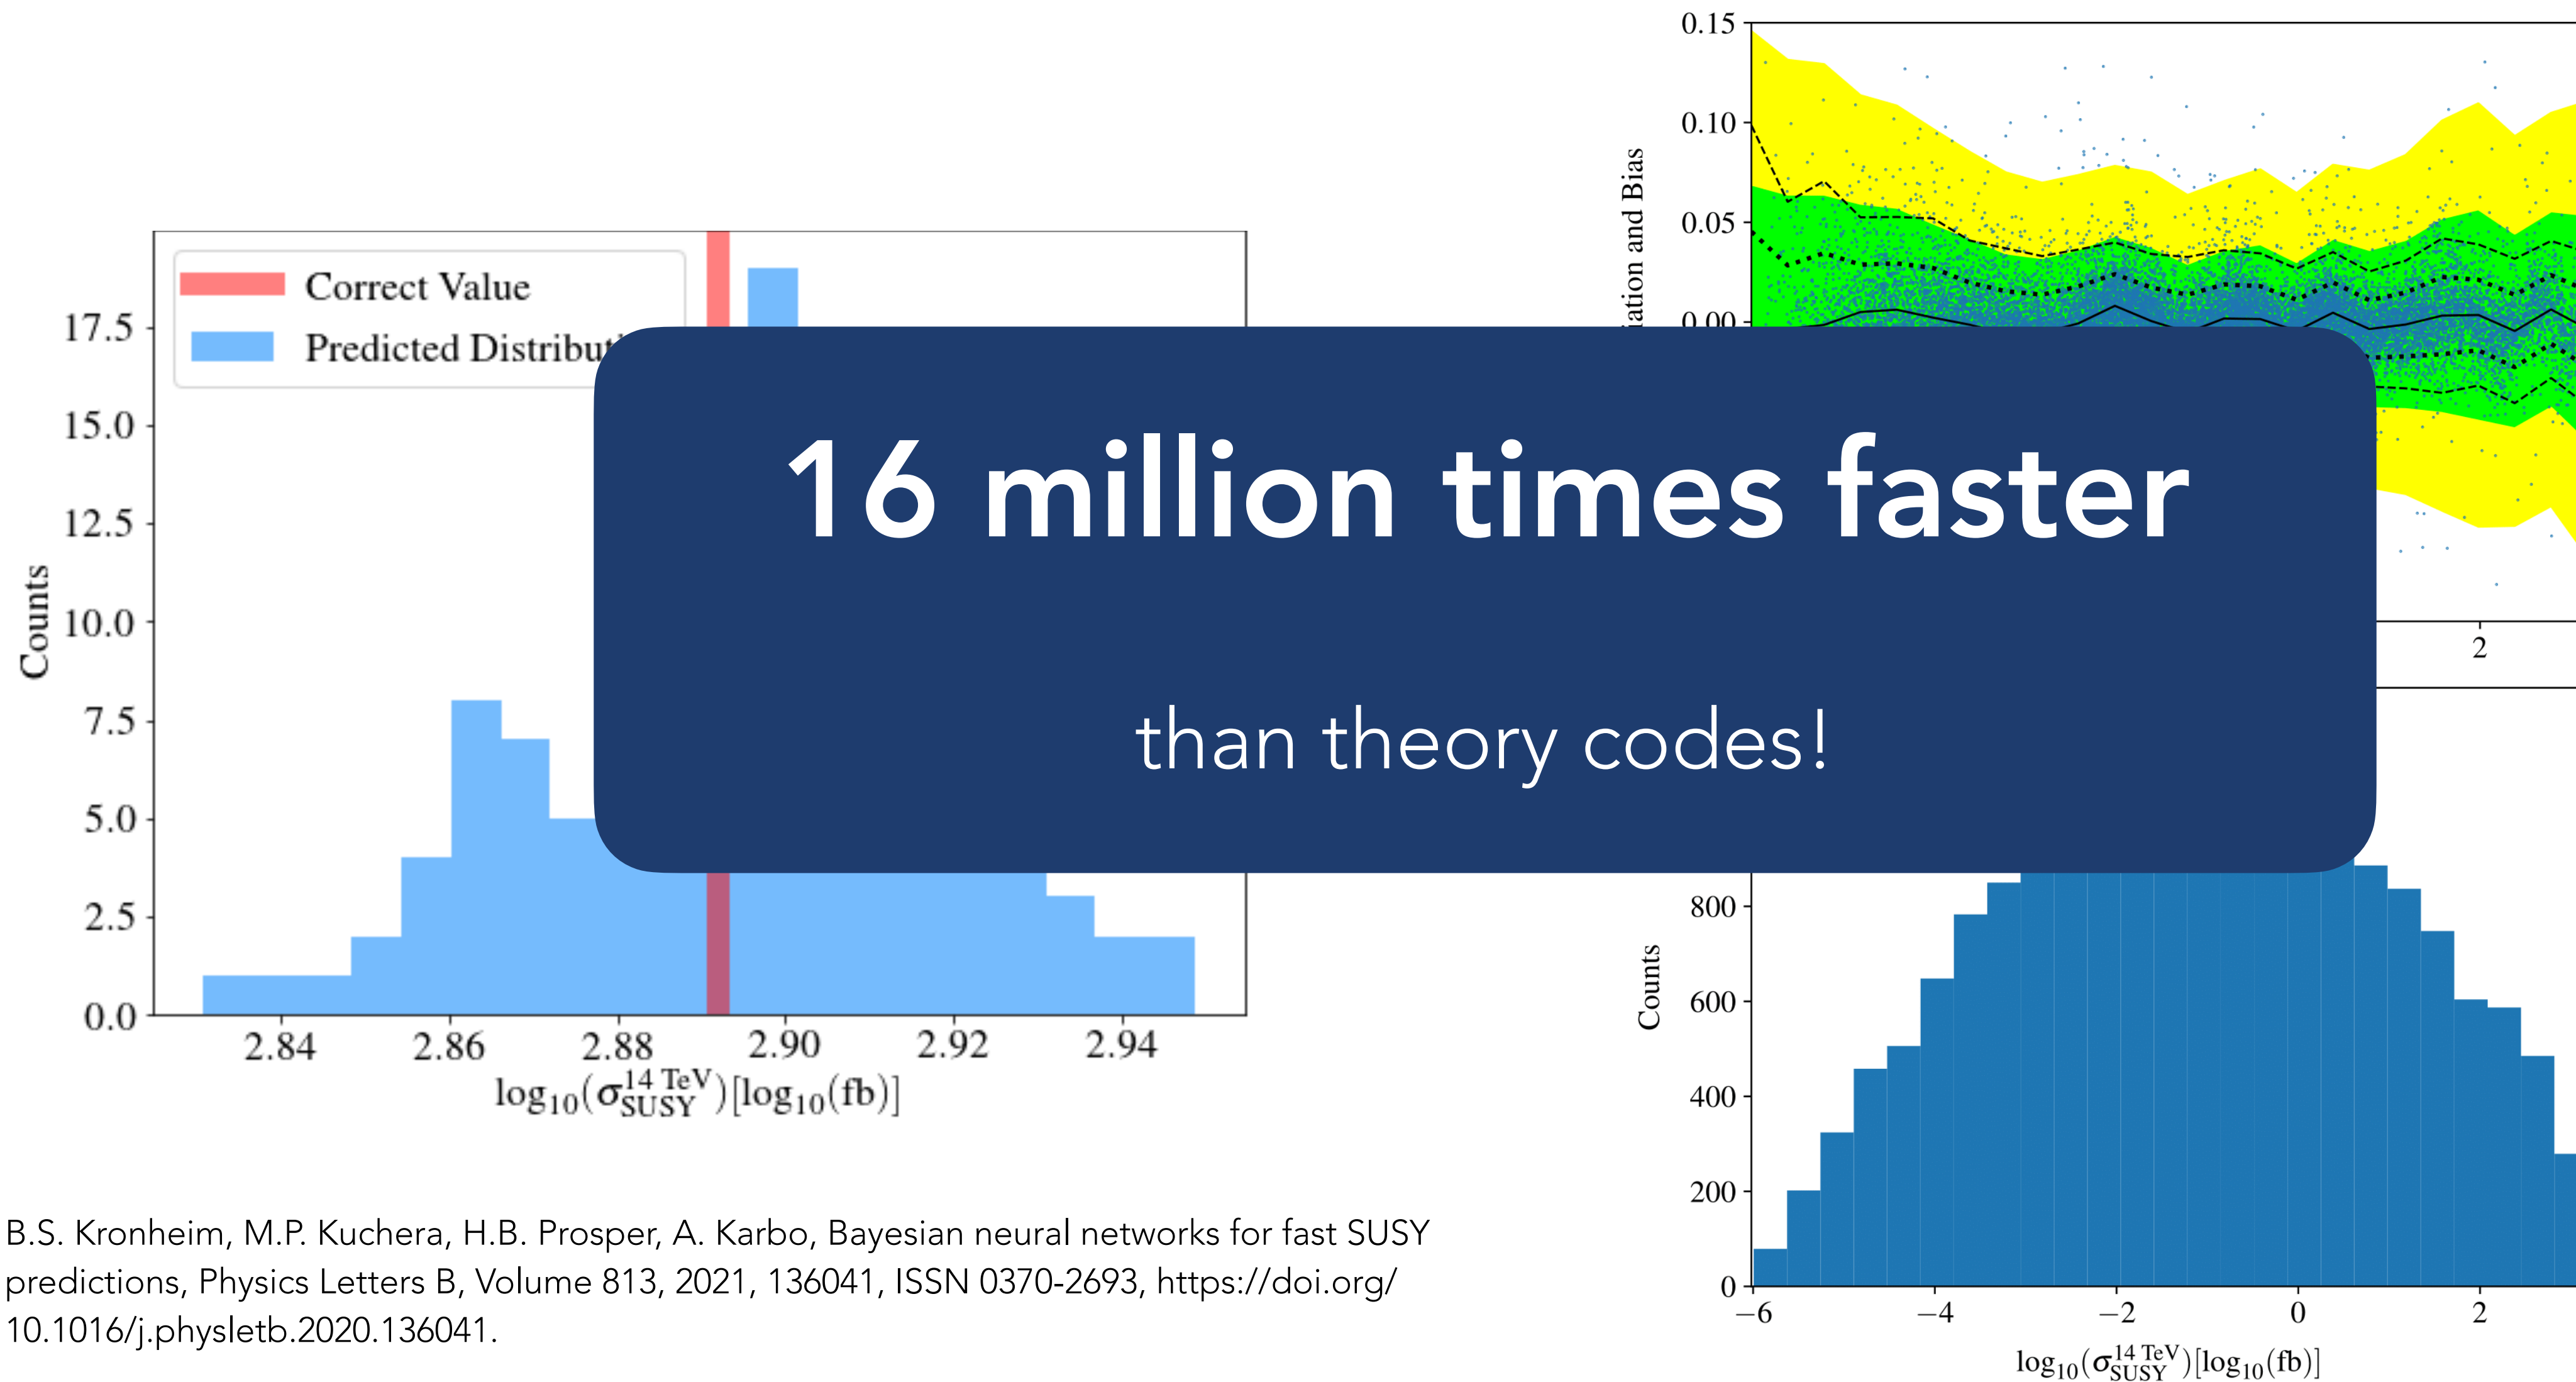
Training — Bayesian inference

Can we make predictions with accurate error estimates?

pMSSM parameters \rightarrow total
SUSY cross section



FAST MAPPING TO THEORETICAL PARAMETERS

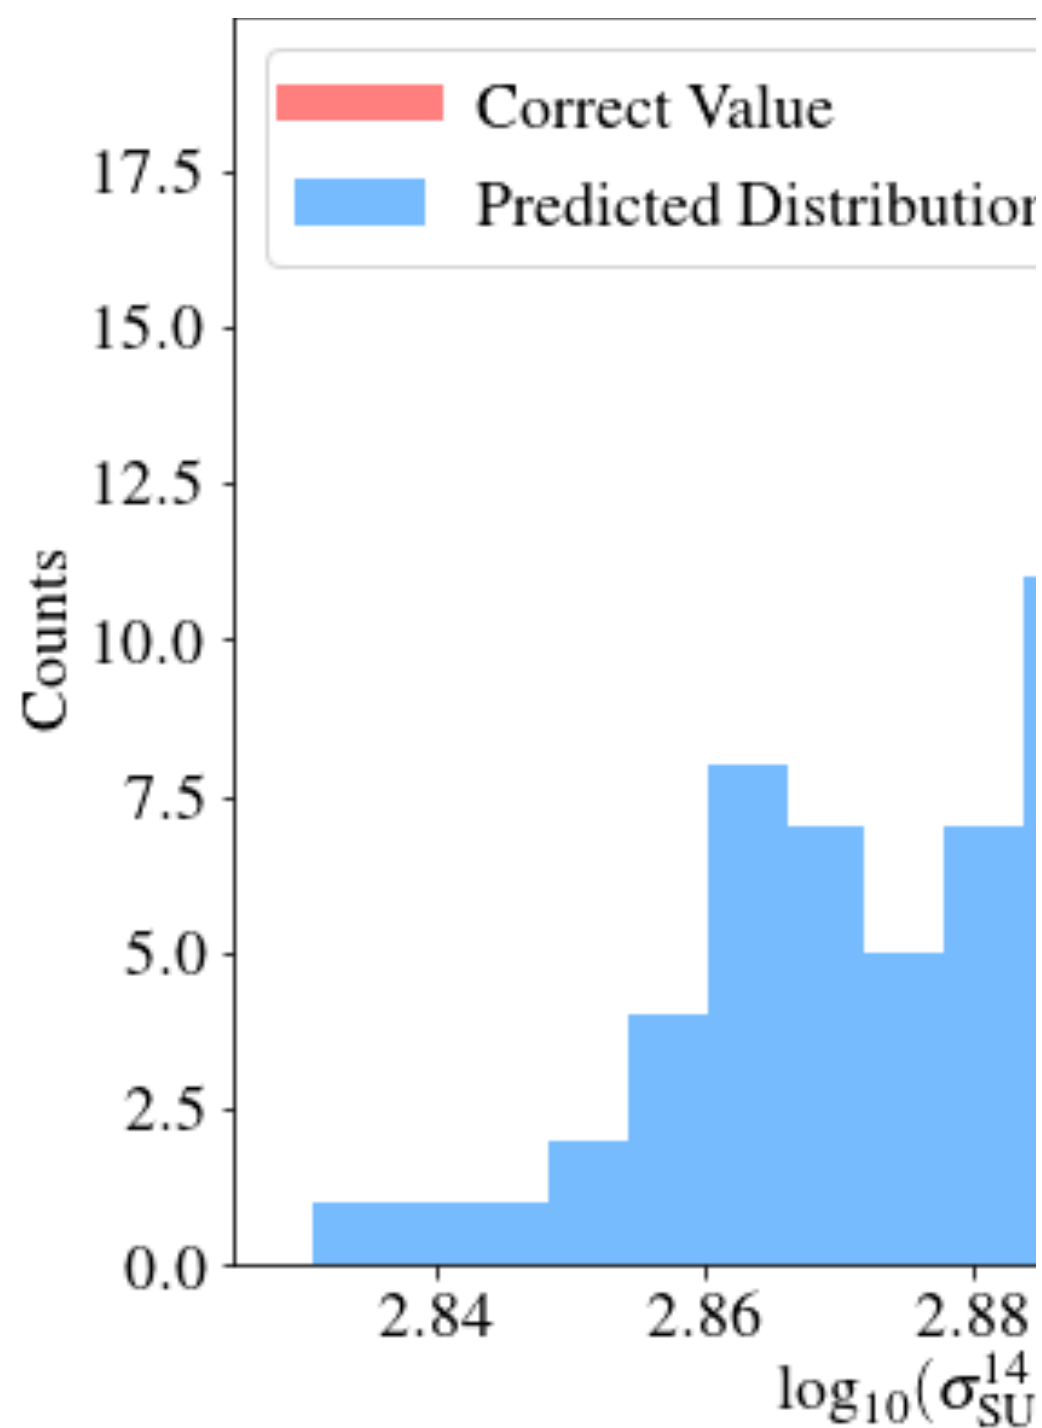


B.S. Kronheim, M.P. Kuchera, H.B. Prosper, A. Karbo, Bayesian neural networks for fast SUSY predictions, Physics Letters B, Volume 813, 2021, 136041, ISSN 0370-2693, <https://doi.org/10.1016/j.physletb.2020.136041>.

<https://arxiv.org/abs/2009.14393>

<https://alpha-davidson.github.io/TensorBNN>

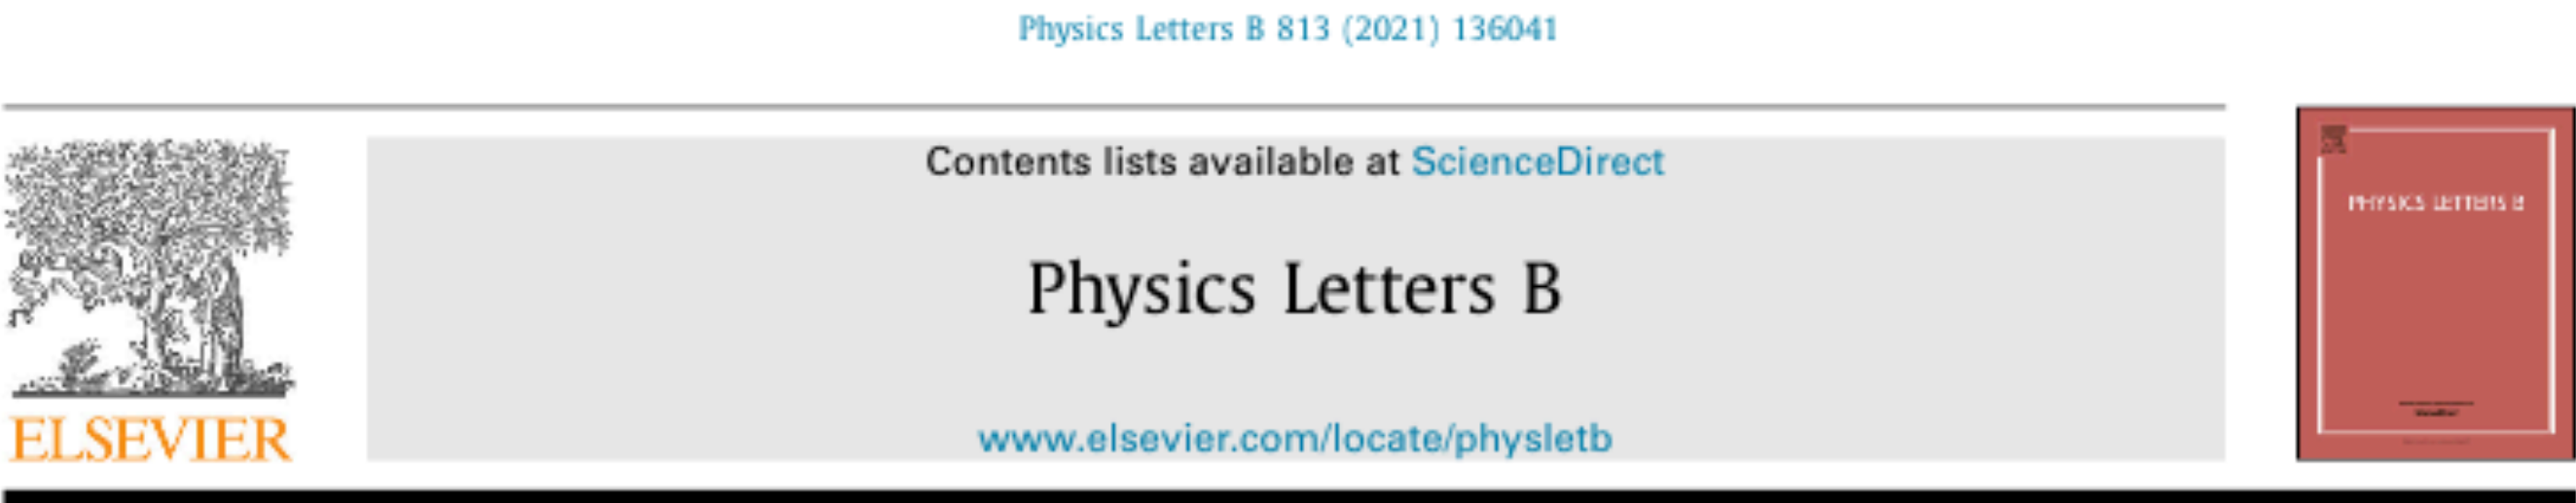
FAST MAPPING TO THEORETICAL PARAMETERS



B.S. Kronheim, M.P. Kuchera, H.B. Prosper, predictions, Physics Letters B, Volume 813 10.1016/j.physletb.2020.136041.

<https://arxiv.org/abs/2009.14393>

<https://alpha-davidson.github.io/TensorBNN>



Bayesian neural networks for fast SUSY predictions

B.S. Kronheim^{a,*}, M.P. Kuchera^a, H.B. Prosper^b, A. Karbo^a

^a Department of Physics, Davidson College, Davidson, NC 28035, USA

^b Department of Physics, Florida State University, Tallahassee, FL 32306, USA



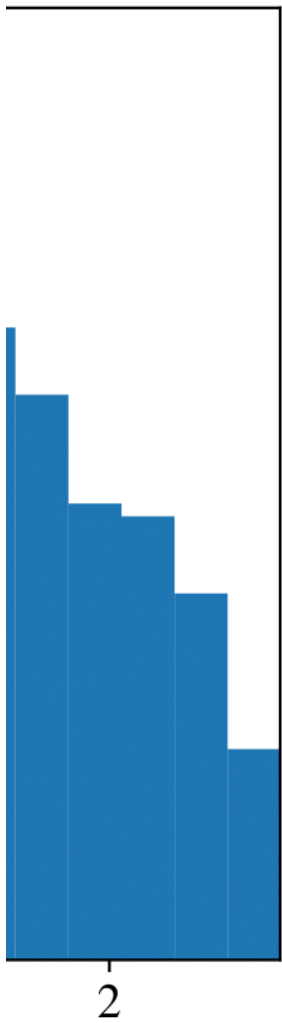
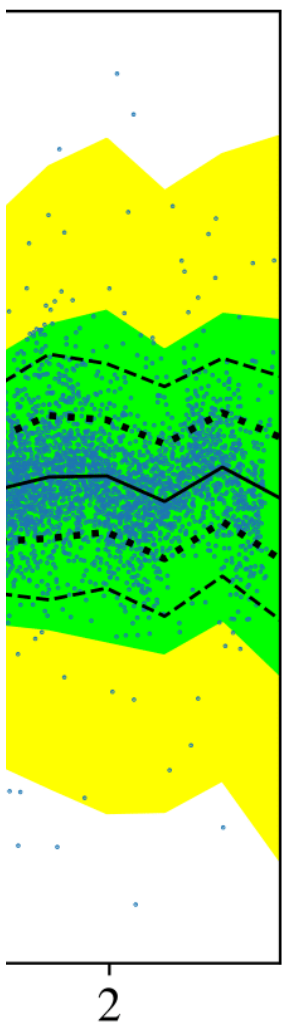
TensorBNN: Bayesian inference for neural networks using TensorFlow[☆]

B.S. Kronheim^{a,*}, M.P. Kuchera^a, H.B. Prosper^b

^a Department of Physics, Davidson College, Davidson, NC 28036, United States of America

^b Department of Physics, Florida State University, Tallahassee, FL 32306, United States of America

$$\log_{10}(\sigma_{SUSY}^{14 \text{ TeV}})[\log_{10}(\text{fb})]$$



PRACTICAL TIPS FOR TRAINING MODELS

DATA

	Feature 1	Feature 2	Feature 3	Target
Example 1				
Example 2				
Example 3				
Example 4				



NORMALIZATION

- Puts each feature on same scale
- Allows default hyperparameters to be a good starting point
 - learning rate, initialization of weights, etc.
- Options depend on data distribution
 - Standardization: mean: 0 stdev: 1
 - Min-max: [0,1]

DATA

	Feature 1	Feature 2	Feature 3	Target
Example 1				
Example 2				
Example 3				
Example 4				

ENCODING

- Non-numeric data
- Class-based features:
 - One-hot encoding: $2 \rightarrow [0 \ 1]$
 - When classes do not have sequential meaning:  cars vs dogs vs plants  months

BUILDING AND TRAINING MODELS

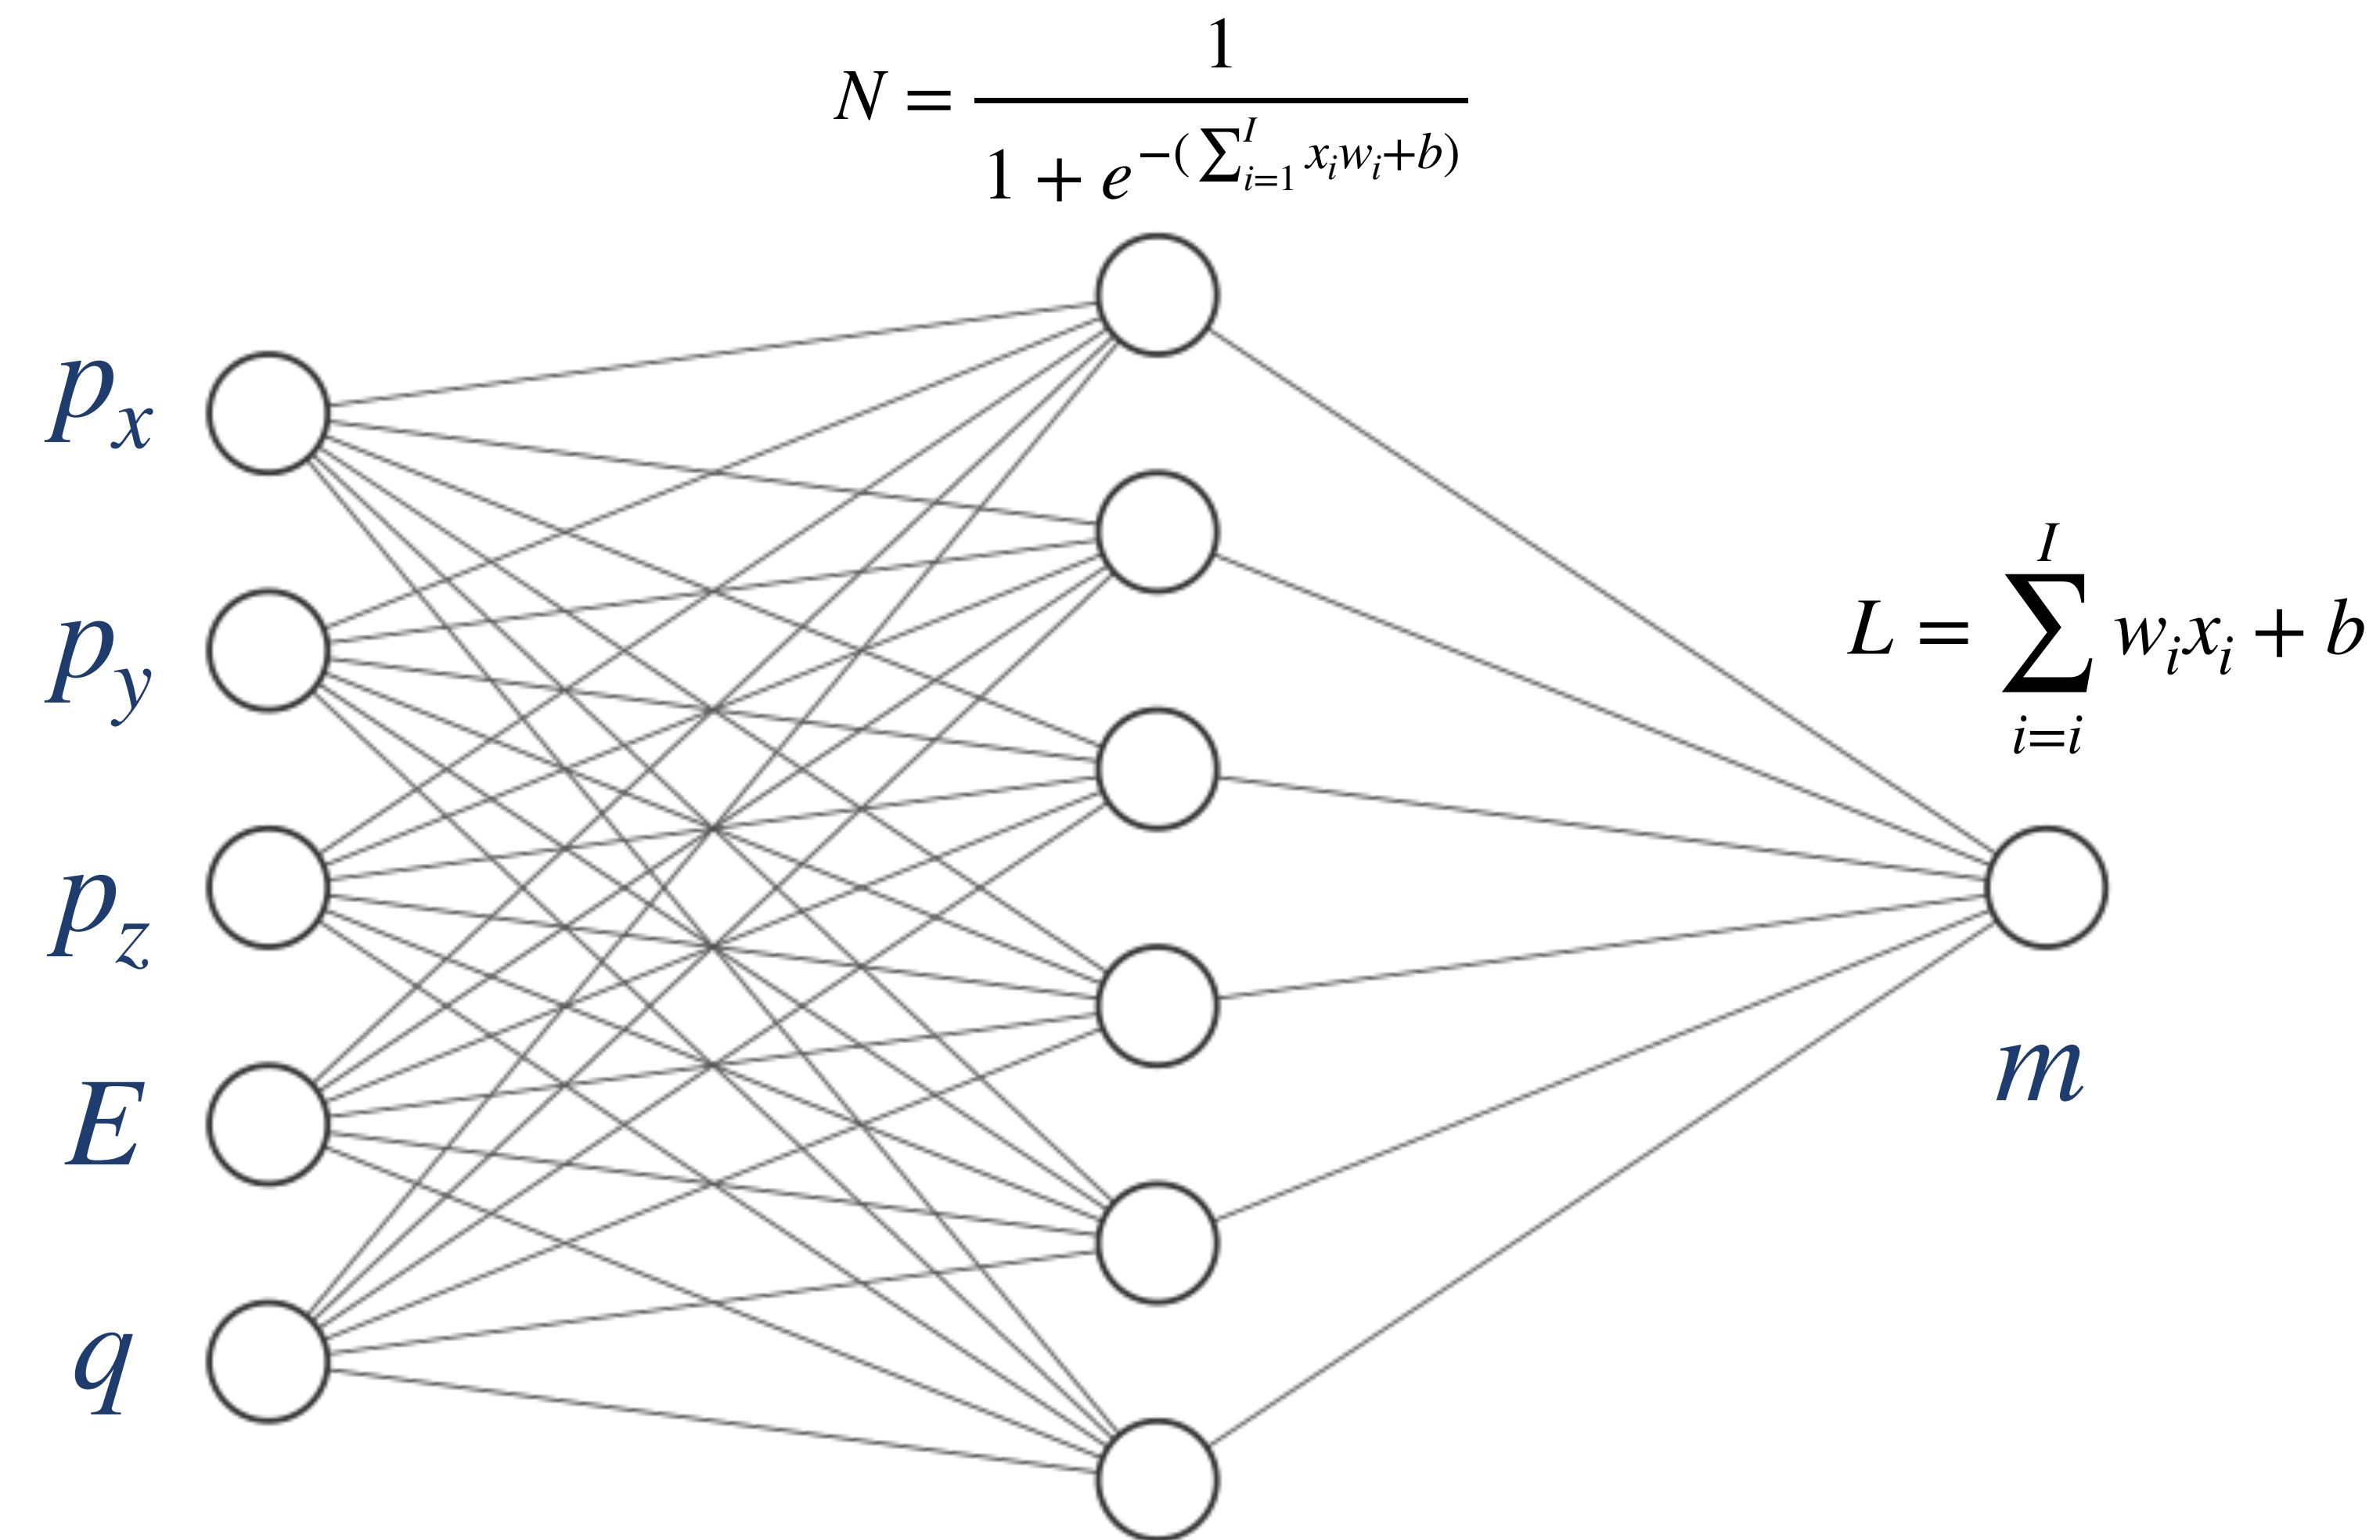
TRAINING

- The most challenging part of machine learning is gaining the experience for tuning models well.
- We will work on this skill!

ACTIVITY DESCRIPTION

- Simulating e+p collisions
- Predicting particle-level invariant mass (regression)
- Advanced: try a generative model (e.g. autoencoders)

ACTIVITY DESCRIPTION



$$m^2 = E^2 - \|p\|^2$$

- Sigmoid activation for hidden layer and linear for output (regression model)
- How many "trainable parameters" in our model?

COMMUNITY

- Each of you arrived here with your own backgrounds, specialty, and path in life
- Your experience and expertise are valuable here, no matter what it is
- If the activity is within your background, help others!
- If you are totally (or a little) lost, ask for help!
- It is our shared goal to have **each** of us leave with some new skill/knowledge/understanding

GETTING STARTED

- Click the link under this tutorial on the workshop page
- If you have access to a google login, click “open in colab”
- Otherwise, download and open in Deepnote or download onto your personal computer (with appropriate dependencies)